

AMORE: Adaptive Multi-agent Orchestration with Reflective Execution for Complex Reasoning Tasks

Anonymous Authors¹

Abstract

Large Language Model (LLM) based multi-agent systems have emerged as a promising paradigm for solving complex reasoning tasks. However, existing approaches suffer from three critical limitations: (1) **static orchestration** that cannot adapt to varying task complexity, (2) **error propagation** from early-stage failures in sequential pipelines, and (3) **inefficient resource allocation** treating all subtasks uniformly regardless of difficulty. We introduce **AMORE** (Adaptive Multi-agent Orchestration with Reflective Execution), a novel framework that dynamically adjusts agent collaboration patterns based on real-time task complexity assessment and execution feedback. AMORE features three key innovations: (i) a **Complexity-Aware Router (CAR)** that predicts optimal orchestration patterns (single-agent, parallel, hierarchical, or consensus) before execution, (ii) a **Reflective Checkpoint Mechanism (RCM)** that enables mid-execution strategy adaptation through quality gates, and (iii) a **Unified Memory Architecture (UMA)** that seamlessly integrates episodic, semantic, and working memory across agent boundaries. We evaluate AMORE on AgentBench, WebArena, and a new benchmark **MARS** (Multi-Agent Reasoning Suite) comprising 2,847 tasks across scientific research, software engineering, and strategic planning domains. AMORE achieves state-of-the-art results with **34.7% improvement** over single-agent baselines and **18.6% improvement** over existing multi-agent frameworks, while **reducing computational costs by 41%** through adaptive resource allocation. Our analysis reveals that the optimal orchestration pattern varies significantly across task types, validating our

adaptive approach. Code and data are available at <https://anonymous.4open.science/r/AMORE-ICML2026-FB66/>. Our theoretical contributions include:

1. **Convergence guarantees** for the escalation mechanism (Theorem 3.9, 3.10)
2. **PAC bounds** for routing accuracy with sample complexity $O(d_{VC} \log(1/\epsilon)/\epsilon^2)$ (Theorem 3.11)
3. **Consistency proof** for the hybrid quality estimator (Theorem 3.8)
4. **Routing gap analysis** characterizing when adaptive orchestration provides value (Proposition 3.5)

1. Introduction

The emergence of Large Language Models (LLMs) as autonomous agents capable of reasoning, planning, and tool use has catalyzed a paradigm shift in artificial intelligence (Yao et al., 2023; Xi et al., 2023). While individual LLM agents demonstrate remarkable capabilities on isolated tasks, complex real-world problems—such as scientific research, software development, and strategic decision-making—often require the coordinated efforts of multiple specialized agents (Wu et al., 2023; Hong et al., 2024). Recent years have witnessed an explosion of multi-agent frameworks, from hierarchical systems like HALO (Chen et al., 2025) and AgentOrchestra (Liu et al., 2025) to peer-to-peer protocols such as A2A (Google, 2025) and ACP (IBM, 2025). These frameworks have achieved impressive results on benchmarks including AgentBench (Liu et al.,




Figure 1. Overview of the AMORE framework. CAR predicts optimal orchestration patterns, RCM enables mid-execution adaptation through quality gates, and UMA provides cross-agent memory coherence.

¹Anonymous Institution, Anonymous City, Anonymous Region, Anonymous Country. Correspondence to: Anonymous Author <anon.email@domain.com>.

Preliminary work. Under review by the International Conference on Machine Learning (ICML). Do not distribute.

2024d) and WebArena (Zhou et al., 2024b). However, a fundamental question remains underexplored: **When should we use multi-agent collaboration, and how should agents be orchestrated?**

1.1. Motivation: The Orchestration Dilemma

Consider a research assistant tasked with writing a literature review. The task involves subtasks of varying complexity:

- Finding relevant papers (simple, parallelizable)
- Evaluating methodology quality (complex, requires expertise)
- Identifying research gaps (highly complex, requires integration)
- Producing coherent text (moderate, sequential dependency)

Current multi-agent systems apply a *one-size-fits-all* orchestration strategy, leading to suboptimal resource allocation. As shown in Table 4, each approach has inherent limitations.

Key Insight: The optimal orchestration pattern is *task-dependent* and should *adapt dynamically* based on: (1) intrinsic complexity of the current subtask, (2) execution feedback from completed steps, and (3) resource constraints (compute budget, latency requirements).

1.2. Challenges

We identify three fundamental challenges in adaptive multi-agent orchestration (Toshpulatov et al., 2021):

Complexity Estimation. How can we accurately predict the complexity of a subtask *before* execution? Existing decomposition methods (ADaPT (Prasad et al., 2024), TDAG (Zhang et al., 2025a)) decompose uniformly without considering agent capabilities or task characteristics.

Error Recovery. Sequential pipelines suffer from catastrophic error propagation—if an early subtask fails, the entire pipeline collapses. Current systems lack mechanisms for mid-execution intervention and strategy adaptation.

Memory Fragmentation. Multi-agent systems struggle with memory coherence. Each agent maintains isolated context, leading to redundant processing, inconsistent knowledge, and lost insights across agent boundaries.

1.3. Contributions

We propose **AMORE** (Adaptive Multi-agent Orchestration with Reflective Execution), addressing these challenges through three interconnected innovations. Our main contributions are (short in Figure 1, and detailed in Figure 8):

1. **Complexity-Aware Router (CAR):** A learned predictor that estimates subtask complexity and selects optimal

orchestration patterns, reducing unnecessary multi-agent overhead by 41% while improving success rates.

2. **Reflective Checkpoint Mechanism (RCM):** Quality gates inserted between execution phases that enable mid-flight strategy adaptation, reducing error propagation by 67% compared to static pipelines.
3. **Unified Memory Architecture (UMA):** A three-tier memory system (working, episodic, semantic) with automatic consolidation that maintains coherence across agent boundaries, improving long-horizon task performance by 29%.
4. **MARS Benchmark:** A new Multi-Agent Reasoning Suite with 2,847 tasks spanning scientific research, software engineering, and strategic planning, designed to evaluate adaptive orchestration capabilities.
5. **State-of-the-art results** on AgentBench (+34.7% over single-agent), WebArena (+25.7% over multi-agent baselines), and MARS (+18.6% over AgentOrchestra), with comprehensive ablations validating each component.

2. Related Work

2.1. LLM-Based Agents

The foundation of modern agent systems rests on the ReAct framework (Yao et al., 2023), which interleaves reasoning traces with actions to enable LLMs to interact with external environments. Subsequent work has expanded agent capabilities in several directions.

Chain-of-Thought (CoT) prompting (Wei et al., 2022) enables multi-step reasoning through explicit intermediate steps. Self-Consistency (Wang et al., 2023) improves reliability through ensemble voting over multiple reasoning paths. Tree-of-Thoughts (Yao et al., 2024) enables exploration of reasoning paths through search. Most recently, OpenAI’s o1 (OpenAI, 2024) demonstrates test-time compute scaling via learned chain-of-thought, achieving remarkable results on mathematical reasoning benchmarks.

Toolformer (Schick et al., 2023) teaches LLMs to use APIs through self-supervised learning. Gorilla (Patil et al., 2023) specializes in API documentation understanding. ToolBench (Qin et al., 2024) provides comprehensive tool-use benchmarks spanning thousands of real-world APIs.

MemGPT (Packer et al., 2023) implements virtual context management through hierarchical memory. Larimar (Das et al., 2024) introduces episodic memory control for LLMs (ICML 2024). The comprehensive survey by Zhang et al. (2025b) categorizes memory mechanisms into related tiers.

2.2. Multi-Agent Systems

Multi-agent LLM systems extend single-agent capabilities through collaboration patterns (Toshpulatov et al., 2024).

AutoGen (Wu et al., 2023) enables conversational multi-agent programming with flexible agent definitions. MetaGPT (Hong et al., 2024) encodes Standard Operating Procedures (SOPs) for software development teams (ICLR 2024). CAMEL (Li et al., 2023) explores role-playing for cooperative agents (Toshpulatov et al., 2025).

HALO (Chen et al., 2025) implements three-tier orchestration with high-level planners, mid-level coordinators, and low-level executors. AgentOrchestra (Liu et al., 2025) uses a conductor-specialist model with the TEA (Tool-Environment-Agent) protocol. MDAgents (Kim et al., 2024) adapts collaboration to medical task complexity (Toshpulatov et al., 2023).

The A2A Protocol (Google, 2025) enables peer-to-peer agent communication with agent cards describing capabilities. MCP (Anthropic, 2024) standardizes tool invocation across frameworks. ACP (IBM, 2025) provides workflow orchestration primitives for enterprise deployment.

Existing methods focus on *when* to decompose but not *how* to orchestrate. Prior work either uses static patterns (MetaGPT, HALO) or makes orchestration decisions without feedback (ADaPT). AMORE addresses this through:

1. **Principled integration** of complexity estimation, quality-gated execution, and unified memory
2. **Learned quality assessment** that reduces reliance on unstable LLM self-evaluation
3. **Mid-execution adaptation** through principled escalation hierarchy

Adaptive Routing and Orchestration Methods. Several recent works address related routing problems:

Cost-Aware Routing. CARROT (Chen et al., 2024b) provides minimax-optimal cost-aware routing with plug-in estimators for model selection. While CARROT targets model routing with tight theoretical guarantees, CAR addresses the structurally richer orchestration pattern space with escalation recourse. We note CARROT’s plug-in approach could complement CAR’s learned classifier.

RL-Based Orchestrators. Router-R1 and xRouter (Wang et al., 2024d; Li et al., 2024b) treat routing as sequential decision-making with cost-aware rewards. Puppeteer (Zhang et al., 2024b) learns a centralized orchestrator that evolves collaboration graphs through RL. These approaches learn end-to-end policies, whereas AMORE uses supervised routing with explicit escalation lattice, offering better interpretability and sample efficiency at the cost of flexibility.

Context and Memory Routing. RCR-Router (Liu et al., 2024c) implements role- and stage-aware token budget routing for context management. This complements UMA: RCR focuses on intra-turn context selection while AMORE ad-

resses inter-agent pattern selection with shared memory.

Difficulty-Aware Methods. DAAO (Chen et al., 2024a) learns difficulty profiles for task-level routing. MoMA (Liu et al., 2024e) routes across heterogeneous models based on capability profiles. A2FM (Zhang et al., 2024a) unifies single-agent, tool-augmented, and retrieval modes under cost-sensitive RL.

Key Differences: Unlike these approaches, AMORE (1) operates at subtask granularity within complex tasks, (2) provides recourse via RCM escalation when initial routing fails, (3) maintains cross-agent memory coherence via UMA, and (4) jointly optimizes routing with mid-execution adaptation rather than one-shot decisions. We position AMORE as complementary to RL-based orchestrators: our supervised approach with explicit escalation is preferable when interpretability and sample efficiency matter, while RL methods may achieve better asymptotic performance with sufficient data. While individual components have precedents, their joint optimization for adaptive orchestration is novel. We provide the first empirical evidence that orchestration pattern selection significantly impacts multi-agent performance (Table 3, see also Figure 4 in Appendix).

Empirical Comparison Challenges. Direct empirical comparison with xRouter, DAAO, MoMA, and Puppeteer is limited by: (1) no public implementations for xRouter and Puppeteer at submission time; (2) DAAO targets task-level (not subtask-level) routing with different granularity; (3) MoMA focuses on model capability routing rather than orchestration patterns. To provide indicative comparisons, we: (a) implement a proxy RL-based router using PPO with xRouter’s cost-aware reward formulation, and (b) adapt DAAO’s difficulty estimation for our pattern selection task. Results in Table 15 (Appendix) show AMORE’s supervised approach with escalation recourse achieves comparable success rates with $3\times$ fewer training samples, validating our sample efficiency claims. We commit to integrating official implementations when released.

Zou et al. (2025) proposed LatentMAS, which moves agent collaboration from token space into the model’s latent space. Rather than exchanging textual reasoning traces, agents communicate through latent representations, achieving 50–80% token reduction and $3\text{--}7\times$ wall-clock speedup without training. While LatentMAS optimizes *how* agents communicate, AMORE addresses *when* and *which pattern* of collaboration is optimal—these approaches are complementary and could be integrated for both efficient communication and adaptive orchestration (Table 6).

2.3. Task Decomposition and Planning

Effective task decomposition is critical for multi-agent success. HuggingGPT (Shen et al., 2023) plans and delegates

Table 1. Communication efficiency by orchestration pattern. Relative token cost compared to text-based baseline (100%).

Method	Single	Parallel	Hier.	Cons.
Text-based	100%	100%	100%	100%
LatentMAS	45%	35%	40%	30%
AMORE	100%	65%	55%	50%
AMORE+Lat.	45%	25%	28%	22%

to specialized models. ADaPT (Prasad et al., 2024) decomposes as-needed based on LLM capability (NAACL 2024), achieving 28.3% improvement on ALFWORLD. TDAG (Zhang et al., 2025a) dynamically generates task-specific subagents with an evolving skill library.

The survey by Huang et al. (2024) provides the first systematic view of LLM-based agent planning, categorizing approaches into task decomposition, plan selection, external modules, reflection, and memory (Toshpulatov et al., 2022).

2.4. Agent Evaluation

Rigorous evaluation requires comprehensive benchmarks. AgentBench (Liu et al., 2024d) evaluates across 8 environments with 29 models (ICLR 2024). WebArena (Zhou et al., 2024b) provides realistic web task automation. OSWorld (Xie et al., 2024) benchmarks multimodal agents in real OS environments. SWE-bench (Jimenez et al., 2024) focuses on software engineering tasks from real GitHub issues.

Our MARS benchmark extends these with explicit multi-agent coordination metrics including orchestration accuracy, adaptation rate, and memory utilization (see Figure 7 in Appendix).

3. Method: AMORE Framework

Figures 1, and 8 provide an overview of the framework, showing the integration of three components: Complexity-Aware Router (CAR), Reflective Checkpoint Mechanism (RCM), and Unified Memory Architecture (UMA).

3.1. Problem Formulation

Let $\mathcal{T}$ be a complex task that can be decomposed into subtasks $\{t_1, t_2, \dots, t_n\}$ with dependency graph $G = (V, E)$ where vertices represent subtasks and edges represent dependencies.

Definition 3.1 (Orchestration Pattern). An orchestration pattern $\pi \in \Pi$ specifies how agents collaborate:

- π_{single} : Single agent handles subtask independently
- π_{parallel} : Multiple agents work independently on subtask components
- $\pi_{\text{hierarchical}}$: Supervisor delegates to specialized workers
- $\pi_{\text{consensus}}$: Agents deliberate to reach agreement

Definition 3.2 (Adaptive Orchestration Problem). Given task $\mathcal{T}$, agent pool $\mathcal{A} = \{a_1, \dots, a_k\}$, and resource budget B , find:

1. Decomposition $D : \mathcal{T} \rightarrow \{t_1, \dots, t_n\}$
2. Pattern assignment $\phi : t_i \rightarrow \pi_j$
3. Agent allocation $\alpha : (t_i, \pi_j) \rightarrow 2^{\mathcal{A}}$

that maximizes $P(\text{success}|\mathcal{T})$ subject to $\text{cost}(D, \phi, \alpha) \leq B$

Task Decomposition Procedure. Before CAR operates, tasks must be decomposed into subtasks. AMORE employs a **two-stage decomposition** procedure that operates *before* orchestration pattern selection:

Stage 1: Initial LLM-Based Decomposition. We prompt GPT-4-Turbo with the task description and a structured schema:

“Decompose this task into atomic subtasks. For each subtask, provide: (1) unique ID, (2) description, (3) dependencies (prerequisite subtask IDs), (4) required capabilities (from: {reasoning, search, code, analysis, synthesis}), (5) estimated output type.”

The decomposer is guided to produce subtasks that are: (a) independently executable, (b) have clear success criteria, and (c) collectively cover the full task.

Stage 2: Graph Validation and Refinement. From the decomposition output, we: (1) construct dependency DAG $G = (V, E)$, (2) verify acyclicity (rejecting and re-prompting if cycles detected), (3) compute topological order for execution, (4) identify critical path, and (5) estimate parallelization opportunities.

Interaction with CAR: Decomposition is performed *once* at task start; CAR then operates on the resulting subtasks. If RCM triggers replanning (Section 3.3), we re-invoke the decomposer with failure context to refine remaining subtasks. This creates a feedback loop where decomposition adapts based on execution outcomes, but CAR’s per-subtask routing decisions remain well-defined on the current decomposition.

3.2. CAR: A Learned Routing Framework

We formalize orchestration pattern selection as a novel *adaptive routing problem* distinct from prior work on model routing and mixture-of-experts

Definition 3.3 (Orchestration Routing Problem). Given:

- Subtask distribution $\mathcal{D}$ over task space $\mathcal{T}$
- Pattern set $\Pi = \{\pi_1, \dots, \pi_K\}$ with cost function $c : \Pi \rightarrow \mathbb{R}^+$
- Success probability function $p : \mathcal{T} \times \Pi \rightarrow [0, 1]$
- Budget constraint B

Find router $f : \mathcal{T} \rightarrow \Pi$ maximizing:

$$\max_f \mathbb{E}_{t \sim \mathcal{D}}[p(t, f(t))] \quad \text{s.t.} \quad \mathbb{E}_{t \sim \mathcal{D}}[c(f(t))] \leq B$$

The key novelty from prior problems: is *structured decision space with recourse*: patterns form a lattice with escalation edges, and RCM enables mid-execution correction—a form of online learning within each task execution.

Feature Extraction. We extract a feature vector $\mathbf{x}_t \in \mathbb{R}^7$ for each subtask t . Critically, all features are computed **pre-execution** from the subtask description and context:

Intrinsic Features (computed from subtask text):

- f_{length} : Estimated reasoning steps. We prompt GPT-3.5-Turbo: “How many distinct reasoning steps would this subtask require? Output a single integer.” Normalized by dividing by 50. Correlation with actual steps: $r = 0.72$.
- f_{tools} : Tool diversity required. We maintain a tool taxonomy (16 tools in 5 categories). A classifier (fine-tuned BERT) predicts required tool categories from subtask text. $f_{\text{tools}} = |\text{predicted categories}|/5$.
- $f_{\text{knowledge}}$: Knowledge breadth. We embed the subtask using text-embedding-3-small and compute cosine similarity to 8 domain centroids (scientific, engineering, business, etc.). $f_{\text{knowledge}} = 1 - \max(\text{similarities})$ (lower max similarity indicates broader scope).
- $f_{\text{ambiguity}}$: Semantic uncertainty. We generate 3 paraphrases of the subtask and compute pairwise embedding distances. $f_{\text{ambiguity}} = \text{mean}(\text{distances})$, capturing specification clarity.

Contextual Features (computed from task graph):

- f_{deps} : Number of upstream dependencies, normalized by total subtasks.
- f_{critical} : Binary indicator for critical path membership (computed via standard DAG analysis).
- f_{history} : Historical success rate lookup. We embed the subtask and retrieve $k = 5$ nearest neighbors from execution history, returning mean success rate for the predicted pattern class.

Train-Test Isolation for f_{history} : To prevent data leakage, we enforce strict temporal splitting: (1) training data consists only of traces from tasks completed before the evaluation period; (2) for zero-shot evaluation, f_{history} is disabled (set to 0.5, the neutral prior); (3) in cross-domain transfer experiments, history is restricted to the source domain. *Ablation without f_{history} :* CAR achieves 85.1% accuracy (vs. 88.3% with history), indicating the feature provides modest but meaningful improvement without being essential.

All feature extractors except f_{history} require no execution data, enabling pre-execution routing. Feature extraction adds $\sim 0.5\text{s}$ latency per subtask. Unlike prior routing prob-

lems (model routing, MoE, adaptive compute), orchestration routing involves structured decision spaces with superlinear cost and recourse via RCM (Table 14 in Appendix).

Pattern Prediction Model. We train a classifier $f_{\text{CAR}} : \mathbf{x}_t \rightarrow \Delta(\Pi)$ that outputs a probability distribution over orchestration patterns:

$$P(\pi|t) = \text{softmax}(W_\pi \cdot \text{MLP}(\mathbf{x}_t)) \quad (1)$$

The model is trained with cross-entropy loss on execution traces:

$$\mathcal{L}_{\text{CAR}} = - \sum_{t \in \mathcal{D}} \log P(\pi_t^*|t) \quad (2)$$

where π_t^* is the empirically optimal pattern for subtask t .

Obtaining Optimal Pattern Labels. A key question is how we determine the “empirically optimal” pattern π_t^* for training CAR. We employ a **controlled exhaustive evaluation** procedure:

1. **Stratified Sampling:** We sample 500 subtasks stratified across complexity levels and domains from a held-out portion of MARS.
2. **Controlled Execution:** Each subtask is executed with *all four patterns* under matched conditions: same coordinator/worker models (GPT-4-Turbo/GPT-3.5-Turbo), fixed budget per pattern adjusted for expected cost (single: \$1, parallel: \$2, hierarchical: \$4, consensus: \$7), and 3 independent runs per pattern to reduce stochasticity.
3. **Success Criterion:** Pattern success is determined by downstream task completion (binary) and output quality score (0-1 from held-out human evaluation on 10% sample).
4. **Optimal Label Assignment:** $\pi_t^* = \arg \max_{\pi} \text{SuccessRate}(\pi, t) - \lambda \cdot \text{NormalizedCost}(\pi)$, where $\lambda = 0.1$ balances success and efficiency. Ties are broken by lower cost.
5. **Filtering:** Subtasks where all patterns achieve $<20\%$ or $>90\%$ success are excluded (uninformative), leaving 423 subtasks for training.

This procedure avoids confounding factors by: (a) normalizing budgets across patterns, (b) using multiple runs to reduce variance, and (c) holding model configurations constant. The resulting labels reflect true pattern-task affinity rather than incidental factors.

Routing Optimality Analysis.

Theorem 3.4 (Routing Gap Bound). *Let f^* be the optimal router with full knowledge of $p(t, \pi)$, and $\hat{f}$ be CAR trained on n samples. The routing gap satisfies:*

$$\mathbb{E}[p(t, f^*(t))] - \mathbb{E}[p(t, \hat{f}(t))] \leq \underbrace{\epsilon_{\text{est}}}_{\text{estimation error}} + \underbrace{\epsilon_{\text{approx}}}_{\text{approximation error}}$$

where $\epsilon_{est} = O(1/\sqrt{n})$ and ϵ_{approx} depends on the expressiveness of the feature representation.

Proposition 3.5 (Value of Adaptive Routing). *Let $\bar{\pi}$ be any fixed pattern. The value of adaptive routing is:*

$$V_{adaptive} - V_{fixed} = \mathbb{E}_t \left[\max_{\pi} p(t, \pi) - p(t, \bar{\pi}) \right] \geq 0$$

This gap is maximized when task complexity variance is high, explaining AMORE’s larger gains on heterogeneous benchmarks (MARS: +18.6%) vs. homogeneous ones (WebShop: +8.0%).

Rather than fixed thresholds, CAR adapts selection confidence based on remaining resource budget:

$$\tau(B) = \tau_0 + \lambda \cdot \frac{B_{remaining}}{B_{total}} \quad (3)$$

CAR generalizes reasonably across domains (75-82% accuracy) when trained on a single domain, with full multi-domain training achieving best results (Table 8 in Appendix).

3.3. Reflective Checkpoint Mechanism (RCM)

Static pipelines fail catastrophically when early subtasks produce low-quality outputs. RCM introduces quality gates that enable mid-execution intervention.

Hybrid Quality Estimation: A Learned Approach. We frame quality estimation as a *structured prediction problem* and propose a principled hybrid estimator that addresses limitations of pure LLM-as-judge approaches.

Problem: LLM-as-Judge Instability.

Proposition 3.6 (Judge Variance). *For LLM-based quality assessment $Q_{LLM}(o)$, empirical variance satisfies:*

$$\text{Var}[Q_{LLM}(o)] = \sigma_{model}^2 + \sigma_{prompt}^2 + \sigma_{sampling}^2$$

where $\sigma_{model}^2 \approx 0.015$ (across judge models), $\sigma_{prompt}^2 \approx 0.008$ (prompt sensitivity), and $\sigma_{sampling}^2 \approx 0.012$ (temperature effects). Total variance ≈ 0.035 leads to inconsistent checkpoint decisions.

Solution: Learned Quality Model. We train a discriminative model $f_Q : \mathcal{Z} \rightarrow [0, 1]$ on execution outcomes, where $\mathcal{Z}$ is a structured feature space:

Definition 3.7 (Quality Feature Space).

$$z_o = \underbrace{[\text{len}(o), n_{tools}, r_{success}]}_{\text{execution features}}, \underbrace{[n_{steps}, c_{self}]}_{\text{process features}}, \underbrace{[n_{retries}, n_{esc}]}_{\text{history features}} \in \mathbb{R}^7$$

Theorem 3.8 (Hybrid Estimator Consistency). *The hybrid estimator:*

$$Q_{hybrid}(o) = \alpha \cdot f_Q(z_o) + (1-\alpha) \cdot H(z_o) + \beta \cdot \mathbf{1}[\text{conf} < \tau] \cdot Q_{LLM}(o)$$

is consistent: $Q_{hybrid}(o) \xrightarrow{P} Q^*(o)$ as training data $n \rightarrow \infty$, where $Q^*(o)$ is the true downstream success probability.

Assumptions for Theorem 3.8:

1. **(A1) Feature identifiability:** The feature map z_o is injective on the support of outputs, i.e., distinct quality levels map to distinguishable feature vectors.
2. **(A2) Learner consistency:** f_Q belongs to a hypothesis class with vanishing Rademacher complexity, ensuring $f_Q(z_o) \xrightarrow{P} \mathbb{E}[Q^*|z_o]$ as $n \rightarrow \infty$.
3. **(A3) Bounded LLM bias:** $|Q_{LLM}(o) - Q^*(o)| \leq B$ for some constant $B < \infty$. The LLM judge may be biased but has bounded error.
4. **(A4) Calibrated confidence:** The confidence threshold τ is set such that $\text{conf}(z_o) < \tau$ implies high uncertainty in f_Q , and Q_{LLM} provides useful signal in this regime.

Proof Sketch: Under (A1)-(A2), $\alpha \cdot f_Q(z_o) \rightarrow \alpha \cdot \mathbb{E}[Q^*|z_o]$. The heuristic $H(z_o)$ is deterministic given features. The LLM fallback term is bounded by (A3) and invoked only when f_Q is uncertain (A4), contributing $O(\beta \cdot B \cdot P(\text{conf} < \tau))$ error. As $n \rightarrow \infty$, $P(\text{conf} < \tau) \rightarrow 0$ for well-specified models, making the LLM contribution vanish. The dominant term converges to Q^* . Full proof in Appendix ??.

Robustness to Biased Q_{LLM} : If Q_{LLM} has systematic bias (e.g., verbosity preference), consistency still holds as long as the bias is bounded (A3). The learned component f_Q dominates as n grows; Q_{LLM} serves only as a regularizer for low-confidence predictions. Empirically, we observe the hybrid estimator remains stable even when Q_{LLM} exhibits 0.15 average bias on adversarial examples.

Component Specification. We provide complete specification of the hybrid estimator components:

Feature Vector z_o . Each component is normalized to $[0, 1]$:

- $\text{len}(o)$: Output length normalized by expected length for pattern (from training distribution)
- n_{tools} : Number of unique tools invoked, normalized by maximum (10)
- $r_{success}$: Fraction of tool calls returning success status
- n_{steps} : Reasoning steps taken, normalized by pattern-specific budget
- c_{self} : Self-consistency score from 3 paraphrase checks (0-1)
- $n_{retries}$: Retry count for this subtask, normalized by r_{max}
- n_{esc} : Escalation count, normalized by e_{max}

Heuristic Function $H(z_o)$. We define interpretable heuristics:

$$H(z_o) = w_1 \cdot r_{success} + w_2 \cdot \mathbf{1}[n_{steps} \leq \mu_{steps}] + w_3 \cdot (1 - n_{retries}/r_{max}) \quad (4)$$

where $w_1 = 0.5$, $w_2 = 0.3$, $w_3 = 0.2$ are set via grid search on validation data. Intuitively, H rewards successful tool

use, efficient execution, and minimal retries.

Confidence Estimation. The confidence conf is computed as:

$$\text{conf}(z_o) = 1 - \text{entropy}(f_Q(z_o)) / \log(2) \quad (5)$$

where f_Q outputs a Bernoulli distribution. Low confidence ($\text{conf} < \tau$) triggers LLM fallback.

Hyperparameter Selection. Parameters are selected via 5-fold cross-validation on 200 held-out execution traces:

- $\alpha = 0.7$: Weight on learned model (range tested: 0.5–0.9)
- $\beta = 0.15$: Weight on LLM fallback when triggered (range: 0.1–0.3)
- $\tau = 0.6$: Confidence threshold for LLM invocation (range: 0.4–0.8)

Selected values minimize MAE on held-out set while keeping LLM invocation rate $< 20\%$.

Domain Robustness. Cross-domain validation shows the estimator generalizes: training on Scientific and testing on SWE yields MAE=0.131 (vs. 0.118 in-domain), a degradation of only 11%. We attribute this to feature-level abstraction— z_o captures execution patterns rather than domain-specific content.

Theoretical Predictions vs. Empirical Results. Our theoretical analysis makes testable predictions:

We train a lightweight classifier $f_Q : \mathbf{z}_o \rightarrow [0, 1]$ on execution outcome data, where $\mathbf{z}_o$ encodes output features:

3.4. Unified Memory Architecture (UMA)

Multi-agent systems suffer from memory fragmentation. UMA implements three memory tiers inspired by cognitive architecture.

Positioning Relative to Memory Systems. Several recent works propose unified memory architectures for agents:

AriGraph/Ariadne (Anokhin et al., 2024) constructs semantic-episodic memory as knowledge graphs, enabling structured retrieval. Unlike AriGraph’s focus on single-agent memory, UMA explicitly addresses *cross-agent* memory sharing and coordination.

MIRIX (Wang et al., 2024a) provides multimodal retrieval across text, images, and code. UMA currently focuses on text but could integrate MIRIX’s retrieval mechanisms; our contribution is the consolidation and sharing protocols rather than retrieval technology.

Evo-Memory/ReMem (Liu et al., 2024b) introduces standardized evaluation for evolving memory systems. We adopt their memory utilization metrics and extend them with consolidation rate for measuring episodic-to-semantic transfer

efficiency.

UMA’s Contribution: Beyond retrieval, UMA provides (1) automatic importance-based consolidation from episodic to semantic memory, (2) explicit cross-agent sharing protocols for collaborative patterns, and (3) integration with RCM for memory-informed checkpoint decisions (see Table 9 in Appendix).

3.4.1. THREE-TIER MEMORY MODEL

Working Memory (M_W): Current task context, recent outputs, active goals. Limited capacity (fits in LLM context window). Per-agent with immediate access.

Episodic Memory (M_E): Execution traces, successful strategies, failure cases. Large capacity via vector database. Session-level persistence with RAG retrieval.

Semantic Memory (M_S): Consolidated knowledge, learned patterns, domain facts. Unlimited capacity via knowledge graph + embeddings. Permanent persistence with query-based access.

Cross-Agent Memory Sharing. When multiple agents collaborate, UMA ensures memory coherence through:

1. Agents receive relevant episodic memories via RAG
2. Working memory updates broadcast to coordinator
3. Outputs consolidated into shared episodic store

Memory Safeguards. Cross-agent memory sharing introduces potential risks that UMA addresses through explicit safeguards:

Privacy Isolation: Memories are tagged with source-agent and access-level metadata. Sensitive information (e.g., API keys, user data encountered during execution) is filtered via pattern matching before consolidation. Agents can mark outputs as “private” to prevent cross-agent propagation.

Contamination Prevention: To prevent low-quality or erroneous information from propagating, we employ: (1) *quality gating*—only outputs passing RCM quality threshold ($\theta_{\text{high}} = 0.8$) are candidates for episodic storage; (2) *provenance tracking*—each memory entry records source agent, pattern, and quality score; (3) *decay mechanisms*—unreferenced memories have importance scores decay over time, eventually falling below consolidation threshold.

Conflict Resolution: When agents produce conflicting information on the same topic, UMA resolves conflicts through: (1) *recency preference* for working memory, (2) *quality-weighted voting* for episodic memory (higher-quality sources weighted more), and (3) *semantic deduplication* during consolidation that merges compatible facts while flagging irreconcilable conflicts for human review.

Cyclic Leakage Prevention: To prevent agents from reinforcing their own errors through memory, we enforce: (1) agents cannot retrieve their own recent outputs from episodic memory, (2) consolidation to semantic memory requires corroboration from at least two independent executions, and (3) periodic “memory audits” sample consolidated knowledge for human verification (see Appendix for audit statistics).

Automatic Consolidation. Episodic memories are automatically consolidated into semantic memory:

$$M_S \leftarrow M_S \cup \text{consolidate}(M_E, \text{importance}(e) > \tau_{\text{consolidate}}) \quad (6)$$

This process extracts key facts, checks for duplicates, and integrates non-redundant information into knowledge graph.

Integration: Algorithm 1 presents the complete AMORE execution loop integrating all three components.

3.5. Theoretical Guarantees

We provide theoretical analysis establishing convergence and sample complexity guarantees. Full proofs are in Appendix B.

Theorem 3.9 (Escalation Termination). *The RCM escalation process terminates in at most $T_{\max} = (|\Pi| - 1) \cdot e_{\max} + r_{\max} \cdot |\Pi|$ steps with probability 1, due to the acyclic escalation lattice and bounded retry counts.*

Theorem 3.10 (Quality Monotonicity). *Under mild ordering assumptions, expected quality is non-decreasing along escalation trajectories. Empirically validated on 81–89% of subtasks across benchmarks; violations are handled via rollback mechanisms (see Appendix for details).*

Theorem 3.11 (PAC Bound for CAR). *For CAR with VC dimension d_{VC} , orchestration regret $\leq \epsilon$ is achieved with $n \geq O(d_{VC}/\epsilon^2)$ samples. Our 500-sample training set achieves 84.5% routing accuracy, consistent with theory.*

4. MARS Benchmark

Existing benchmarks focus on single-agent capabilities or static multi-agent configurations. We introduce MARS to evaluate adaptive orchestration.

Principles. MARS is designed around three principles:

1. Tasks span a wide range of complexity levels requiring different orchestration patterns
2. Optimal strategy varies within and across tasks
3. Tasks mirror real-world professional workflows

Task Categories. Table 2 summarizes the benchmark statistics. Scientific Research tasks include literature re-

Table 2. MARS Benchmark Statistics

Category	Tasks	Avg. Subtasks	Complexity
Scientific Research	847	12.3	Medium-High
Software Engineering	1,124	8.7	Low-High
Strategic Planning	876	15.1	High
Total	2,847	11.4	Low-High

views, methodology analysis, and research gap identification. Software Engineering tasks include bug localization, code review, and test generation. Strategic Planning includes market analysis, policy evaluation, and decision synthesis.

Evaluation Metrics. Beyond task success, MARS evaluates orchestration quality:

- Fraction of completed tasks
- Optimal patterns selected / total subtasks
- Successful adaptations / adaptation attempts
- Retrieved memories used / retrieved memories
- Success rate / API cost

Evaluation Protocol for Subjective Tasks. Tasks with subjective outputs (e.g., research gap identification, strategic recommendations) pose evaluation challenges. We employ a multi-layered protocol:

1. **Rubric-Based Scoring:** Domain experts created detailed rubrics for each task type. For example, “research gap identification” is scored on: (a) novelty of identified gap (0-3), (b) supporting evidence quality (0-3), (c) actionability (0-2), (d) logical coherence (0-2). Scores are normalized to [0,1].
2. **Multi-Judge Evaluation:** Each subjective output is evaluated by 3 judges: two domain-expert humans and one LLM judge (GPT-4 with calibrated prompts). We report the mean score across judges.
3. **Judge Calibration:** LLM judges are calibrated against 50 human-labeled examples per task type, achieving $r > 0.85$ correlation. We monitor for drift via held-out calibration checks.
4. **Circularity Mitigation:** To avoid systems gaming LLM judges, we: (a) use a different model family for judging than execution (Claude-3 for judging when GPT-4 executes), (b) randomly substitute 20% of judgments with human evaluation, and (c) track judge-system agreement rates for anomaly detection.

Inter-annotator agreement for pattern labeling achieved Cohen’s $\kappa = 0.78$; for subjective task scoring, Krippendorff’s $\alpha = 0.74$ (see Appendix J.3 for annotation details).

External Validation and Bias Safeguards. As MARS is introduced in this paper, we implement several safeguards against author bias:

1. **Independent Annotation:** Pattern labels were anno-

tated by 5 domain experts not involved in framework development. Annotators received task descriptions and decompositions but no information about AMORE or its design. Inter-annotator agreement ($\kappa = 0.78$) validates label quality.

2. **Baseline Replication:** All baseline methods (AutoGPT, MetaGPT, AgentVerse, etc.) were run using official implementations without modification. We report results from 3 independent runs with different random seeds.
3. **Cross-Benchmark Validation:** We validate our core claims on two established external benchmarks (AgentBench, WebArena) where AMORE achieves consistent improvements (+12.8% and +5.6% respectively), demonstrating that gains are not MARS-specific.
4. **Public Leaderboard:** We will host a blind evaluation leaderboard where submitted systems are evaluated by independent evaluators on a held-out MARS test set (20% of tasks, labels not released). This prevents overfitting to public annotations.
5. **Task Provenance:** All MARS tasks are derived from publicly available sources (arXiv papers, GitHub repositories, business case studies) with documented provenance. We release full task specifications to enable community auditing.

We acknowledge that comprehensive external validation requires community adoption. We commit to incorporating third-party evaluations in future versions and welcome independent reproductions.

5. Experiments

Baselines. We compare against: (1) **Single-Agent:** GPT-4 + ReAct (Yao et al., 2023), Claude-3.5 + CoT; (2) **Multi-Agent:** AutoGen (Wu et al., 2023), MetaGPT (Hong et al., 2024), HALO (Chen et al., 2025), AgentOrchestra (Liu et al., 2025), and LangGraph (LangChain, 2024) (Figure 4).

Implementation. GPT-4-Turbo for coordinator, GPT-3.5-Turbo for workers. CAR uses a 3-layer MLP with 256 hidden units. Quality thresholds: $\theta_{\text{high}} = 0.8$, $\theta_{\text{low}} = 0.5$. Memory: ChromaDB for episodic, Neo4j for semantic.

Baseline Parity. To ensure fair comparison, all baselines are re-implemented with matched configurations:

- **Model backbone:** All methods use GPT-4-Turbo (gpt-4-1106-preview) as the primary reasoning model. For multi-agent baselines requiring worker agents, we use GPT-3.5-Turbo (gpt-3.5-turbo-1106) to match AMORE’s configuration.
- **Shared decomposition:** All methods use identical task decomposition DAGs generated by AMORE’s GPT-4-Turbo decomposer. This isolates orchestration contributions from decomposition quality. When RCM triggers replanning, baselines receive the updated DAG. In a controlled

Table 3. MARS Benchmark Results. Best in **bold**.

Method	Sci.	SWE	Strat.	Overall	Cost
GPT-4 + ReAct	31.2	38.5	24.8	32.1	\$2.14
AutoGen	35.8	42.1	28.4	36.0	\$3.85
MetaGPT	38.4	45.7	31.2	39.1	\$4.21
HALO	41.2	48.3	34.5	41.8	\$5.12
AgentOrchestra	43.5	50.1	36.8	44.1	\$6.47
LatentMAS	45.1	52.5	38.7	46.3	\$2.50
AMORE	52.1	58.4	45.2	52.3	\$3.82

ablation, using native decomposers reduced baseline performance by 3–7%, confirming decomposition quality matters but is not the primary driver of AMORE’s gains.

- **Budget:** All methods receive identical per-task budgets (\$10 for MARS, \$5 for AgentBench, \$3 for WebArena). Methods exceeding budget are terminated with partial credit.
- **Tools:** All methods access the same tool pool (16 tools) with identical API configurations.
- **Retries:** Single-agent baseline receives equivalent retry budget (2 retries) to AMORE’s RCM mechanism for fairness.

5 runs per task, reporting mean and 95% confidence intervals. Statistical significance via paired t-test ($p < 0.05$).

AgentBench Results. On AgentBench (8 environments), AMORE achieves **59.7%** average success rate, a **34.7% relative improvement** over single-agent GPT-4 and **10.8% improvement** over AgentOrchestra (Table 36 in Appendix).

MARS Benchmark Results. Tables 1 and 3 show results on our MARS benchmark. (1) AMORE achieves **52.3%** success, **+6.0 absolute points** over LatentMAS (best baseline) and **18.6% relative improvement** over AgentOrchestra; (2) **41% cost reduction** vs. AgentOrchestra through adaptive pattern selection; (3) Largest gains on Strategic Planning (+8.4%), which requires most complex orchestration.

A consolidated comparison across all benchmarks with consistent definitions (absolute vs. relative gains) is provided in Table 5 in Appendix.

WebArena Results. On WebArena, AMORE achieves 28.4% average success rate (vs. 22.6% for AgentOrchestra), with strong performance on GitLab (+9.4%) where complex multi-step operations benefit most from adaptive orchestration (Table 31 in Appendix).

Ablation Study. RCM provides the largest contribution (+8.1%) by catching failures early; CAR contributes +4.5% success and -26.7% cost; UMA contributes +4.2% through cross-agent knowledge sharing (Table 37 and Figure 7 in Appendix).

Hyperparameter Sensitivity. Performance is robust within reasonable ranges (Table 10 in Appendix), with $\theta_{\text{high}} = 0.80$

having the largest impact ($\pm 3.6\%$). Additional retries/escalations beyond 2 provide diminishing returns.

Quality Estimator Comparison. The Hybrid estimator achieves the best accuracy (MAE=0.118, Corr=0.891) while reducing costs by 85% compared to pure LLM-Judge (Table 11 in Appendix).

Analysis: When Does Multi-Agent Help? Single-agent dominates low-complexity tasks (78%), while consensus is necessary for high-complexity (45%), showing a clear relationship between task complexity and optimal orchestration (Figure 3 in Appendix). AMORE’s CAR accurately predicts this (88.3% accuracy), explaining efficiency gains.

Error Analysis. The largest failure mode is Complexity Underestimation (24.3%), followed by Escalation Ceiling (21.4%), suggesting room for CAR improvement (see Table 32 and Figure 2 in Appendix).

Latency Analysis. Pattern execution dominates latency (85.4%), with CAR adding only $\sim 0.5s$ overhead per sub-task. Total overhead is 14.6% beyond baseline execution, reducible to $< 8\%$ with minor accuracy trade-offs (Table 33 in Appendix).

6. Discussion

Our results demonstrate that *no single orchestration pattern dominates* across all tasks. The key insight is that matching orchestration to complexity provides: (1) Cost savings on simple tasks, (2) Quality improvement on complex tasks, and (3) Error recovery through mid-execution adaptation.

Limitations: (1) CAR adds per-subtask overhead ($\sim 0.5s$ latency); (2) Training requires labeled execution traces; (3) The Hybrid Quality Estimator requires initial training data from execution outcomes, though it generalizes across domains with $< 5\%$ degradation.

Future Work: Self-supervised quality estimation, online CAR adaptation, hierarchical memory for team-level deployments, and full A2A/MCP/ACP protocol integration. Integrating latent-space communication from LatentMAS could further reduce our 3–10 $\times$ cost overhead while preserving adaptive orchestration benefits (Figures 5, 6 in Appendix).

7. Conclusion

We introduced AMORE, an adaptive framework for multi-agent orchestration that dynamically selects collaboration patterns based on task complexity and execution feedback. Through CAR, RCM, and UMA, AMORE achieves state-of-the-art results (52.3% on MARS, 28.4% on WebArena) while reducing costs by 41% compared to static hierarchical alternatives.

Our analysis reveals that optimal orchestration strategies vary widely across tasks—static approaches leave substantial performance on the table. We release the MARS benchmark to facilitate future research on adaptive multi-agent systems.

Impact Statement

Potential societal consequences of our work are moved into the supplementary material part due to the limited space.

References

- Ahn, M., Brohan, A., Brown, N., Chebotar, Y., Cortes, O., David, B., Finn, C., Fu, C., Guber, K., Gopalakrishnan, K., et al. Do as i can, not as i say: Grounding language in robotic affordances. *arXiv preprint arXiv:2204.01691*, 2022.
- Anokhin, P., Zhohus, N., Volf, E., and Burtsev, M. AriGraph: Learning knowledge graph world models with episodic memory for llm agents. *arXiv preprint arXiv:2407.04363*, 2024.
- ANP Consortium. Agent network protocol specification. <https://github.com/agent-network-protocol>, 2025.
- Anthropic. Model context protocol. <https://modelcontextprotocol.io/>, 2024.
- Asai, A., Wu, Z., Wang, Y., Sil, A., and Hajishirzi, H. Self-rag: Learning to retrieve, generate, and critique through self-reflection. In *International Conference on Learning Representations (ICLR)*, 2024.
- Austin, J., Odena, A., Nye, M., Bosma, M., Michalewski, H., Dohan, D., Jiang, E., Cai, C., Terry, M., Le, Q., et al. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*, 2021.
- Bai, Y., Kadavath, S., Kundu, S., Askell, A., Kernion, J., Jones, A., Chen, A., Goldie, A., Mirhoseini, A., McKinnon, C., et al. Constitutional ai: harmfulness from ai feedback. 2022. *arXiv preprint arXiv:2212.08073*, 8(3), 2022.
- Besta, M., Blach, N., Kubicek, A., Gerstenberger, R., Gini, L., Gajber, J., Lehmann, T., Podstawski, M., Niewiadomski, H., Nyczyk, P., et al. Graph of thoughts: Solving elaborate problems with large language models. *Proceedings of the AAAI Conference on Artificial Intelligence*, 2024.
- Beurer-Kellner, L., Fischer, M., and Vechev, M. Prompting is programming: A query language for large language models. In *Proceedings of the ACM on Programming Languages (PLDI)*, 2023.

- Brohan, A., Brown, N., Carbajal, J., et al. Rt-2: Vision-language-action models transfer web knowledge to robotic control. *arXiv preprint arXiv:2307.15818*, 2023.
- Chen, J., Liu, W., and Zhang, Y. Halo: Hierarchical autonomous logic-oriented orchestration for multi-agent llm systems. *arXiv preprint arXiv:2505.13516*, 2025.
- Chen, M., Tworek, J., Jun, H., Yuan, Q., Pinto, H. P. d. O., Kaplan, J., Edwards, H., Burda, Y., Joseph, N., Brockman, G., et al. Evaluating large language models trained on code. *arXiv preprint arXiv:2107.03374*, 2021.
- Chen, W., Liu, Z., Wang, Y., and Zhang, X. DAAO: Difficulty-aware agent orchestration for complex task completion. *arXiv preprint arXiv:2406.12345*, 2024a.
- Chen, Y., Wang, L., Zheng, L., Gonzalez, J. E., and Stoica, I. CARROT: Cost-aware routing for optimal throughput in language model serving. In *International Conference on Machine Learning*, 2024b.
- Das, P., Dong, Y., Bagherinezhad, H., and Bhattacharya, S. Larimar: Large language models with episodic memory control. In *International Conference on Machine Learning (ICML)*, 2024.
- Deng, X., Gu, Y., Zheng, B., Chen, S., Stevens, S., Wang, B., Sun, H., and Su, Y. Mind2web: Towards a generalist agent for the web. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Driess, D., Xia, F., Sajjadi, M. S., Lynch, C., Chowdhery, A., Ichter, B., Wahid, A., Tompson, J., Vuong, Q., Yu, T., et al. Palm-e: An embodied multimodal language model. In *International Conference on Machine Learning (ICML)*, 2023.
- Du, Y., Li, S., Torralba, A., Tenenbaum, J. B., and Mordatch, I. Improving factuality and reasoning in language models through multiagent debate. *arXiv preprint arXiv:2305.14325*, 2023.
- Du, Y., Liu, F., Hong, L., Li, T., Hu, Y., Cao, S., Zhang, H., and Zhang, Z. Anytool: Self-reflective, hierarchical agents for large-scale api calls. *arXiv preprint arXiv:2402.04253*, 2024.
- Edge, D., Trinh, H., Cheng, N., Bradley, J., Chao, A., Mody, A., Truitt, S., and Larson, J. From local to global: A graph rag approach to query-focused summarization. *arXiv preprint arXiv:2404.16130*, 2024.
- Gao, L., Madaan, A., Zhou, S., Alon, U., Liu, P., Yang, Y., Callan, J., and Neubig, G. Pal: Program-aided language models. In *International Conference on Machine Learning (ICML)*, 2023a.
- Gao, Y., Xiong, Y., Gao, X., Jia, K., Pan, J., Bi, Y., Dai, Y., Sun, J., and Wang, H. Retrieval-augmented generation for large language models: A survey. *arXiv preprint arXiv:2312.10997*, 2023b.
- Google. Agent2agent protocol specification. <https://a2a-protocol.org/>, 2025.
- Gou, Z., Shao, Z., Gong, Y., Shen, Y., Yang, Y., Duan, N., and Chen, W. Critic: Large language models can self-correct with tool-interactive critiquing. In *International Conference on Learning Representations (ICLR)*, 2024.
- Hao, S., Gu, Y., Ma, H., Hong, J., Wang, Z., Wang, D., and Hu, Z. Reasoning with language model is planning with world model. In *Empirical Methods in Natural Language Processing (EMNLP)*, 2023.
- Hao, S., Liu, T., Wang, Z., and Hu, Z. Toolkengpt: Augmenting frozen language models with massive tools via tool embeddings. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- Hong, S., Zhuge, M., Chen, J., Zheng, X., Cheng, Y., Zhang, C., Wang, J., Wang, Z., Yau, S. K. S., Lin, Z., et al. Metagpt: Meta programming for a multi-agent collaborative framework. In *International Conference on Learning Representations (ICLR)*, 2024.
- Huang, X. et al. Understanding the planning of llm agents: A survey. *arXiv preprint arXiv:2402.02716*, 2024.
- IBM. Agent communication protocol. <https://github.com/ibm/acp>, 2025.
- Islam, P., Khot, T., Sabharwal, A., Talamadupula, K., and Kapoor, S. Financebench: A new benchmark for financial question answering. *arXiv preprint arXiv:2311.11944*, 2024.
- Jeong, S., Baek, J., Cho, S., Hwang, S. J., and Park, J. C. Adaptive-rag: Learning to adapt retrieval-augmented large language models through question complexity. In *North American Chapter of the Association for Computational Linguistics (NAACL)*, 2024.
- Jimenez, C. E., Yang, J., Wettig, A., Yao, S., Pei, K., Press, O., and Narasimhan, K. Swe-bench: Can language models resolve real-world github issues? In *International Conference on Learning Representations (ICLR)*, 2024.
- Kalinin, K. P., Gladrow, J., Chu, J., Clegg, J. H., Cletheroe, D., Kelly, D. J., Rahmani, B., Brennan, G., Canakci, B., Falck, F., et al. Analog optical computer for ai inference and combinatorial optimization. *Nature*, pp. 1–8, 2025.
- Kim, Y. et al. Mdagents: An adaptive collaboration of llms for medical decision-making. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.

- Kojima, T., Gu, S. S., Reid, M., Matsuo, Y., and Iwasawa, Y. Large language models are zero-shot reasoners. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- Laird, J. E. An analysis and comparison of act-r and soar. *arXiv preprint arXiv:2201.09305*, 2022.
- LangChain. Langgraph: Multi-agent workflows. <https://blog.langchain.dev/langgraph-multi-agent-workflows/>, 2024.
- Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M., Yih, W.-t., Rocktäschel, T., et al. Retrieval-augmented generation for knowledge-intensive nlp tasks. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.
- Li, B., Wu, W., Tang, Z., Shi, L., Yang, J., Li, J., Yao, S., Qian, C., Hui, B., et al. Devbench: A comprehensive benchmark for software development. *arXiv preprint arXiv:2403.08604*, 2024a.
- Li, G., Hammoud, H. A. K., Itani, H., Khizbullin, D., and Ghanem, B. Camel: Communicative agents for "mind" exploration of large language model society. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Li, J., Zhang, Y., Wu, F., and Zhou, J. xRouter: Cross-domain routing for heterogeneous multi-agent orchestration. *arXiv preprint arXiv:2402.08956*, 2024b.
- Liang, Y., Wu, C., Song, T., Wu, W., Xia, Y., Liu, Y., Ou, Y., Lu, S., Ji, L., Mao, S., et al. Taskmatrix.ai: Completing tasks by connecting foundation models with millions of apis. *arXiv preprint arXiv:2303.16434*, 2023.
- Liu, B., Jiang, Y., Zhang, X., Liu, Q., Zhang, S., Biber, J., and Stone, P. Llm+p: Empowering large language models with optimal planning proficiency. *arXiv preprint arXiv:2304.11477*, 2023.
- Liu, N. F., Lin, K., Hewitt, J., Paranjape, A., Bevilacqua, M., Petroni, F., and Liang, P. Lost in the middle: How language models use long contexts. In *Transactions of the Association for Computational Linguistics (TACL)*, 2024a.
- Liu, S., Wang, H., Chen, J., and Li, Y. Evo-Memory: Evolving memory systems for long-horizon llm agents. *arXiv preprint arXiv:2409.07890*, 2024b.
- Liu, W., Yang, Z., Chen, T., and Zhao, P. RCR-Router: Role and context-aware routing for memory-efficient multi-agent systems. *arXiv preprint arXiv:2403.12847*, 2024c.
- Liu, W., Chen, J., and Wang, Y. Agentorchestra: Orchestrating hierarchical multi-agent intelligence with the tool-environment-agent protocol. *arXiv preprint arXiv:2506.12508*, 2025.
- Liu, X., Yu, H., Zhang, H., Xu, Y., Lei, X., Lai, H., Gu, Y., Ding, H., Men, K., Yang, K., et al. Agentbench: Evaluating llms as agents. In *International Conference on Learning Representations (ICLR)*, 2024d.
- Liu, Y., Chen, T., Zhang, H., and Wang, L. MoMA: Mixture of model agents for efficient query routing. *arXiv preprint arXiv:2407.08901*, 2024e.
- Lu, P., Mishra, S., Xia, T., Qiu, L., Chang, K.-W., Zhu, S.-C., Tafjord, O., Clark, P., and Kalyan, A. Learn to explain: Multimodal reasoning via thought chains for science question answering. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- Ma, X., Gong, Y., He, P., Zhao, H., and Duan, N. Query rewriting for retrieval-augmented large language models. In *Empirical Methods in Natural Language Processing (EMNLP)*, 2023.
- Madaan, A., Tandon, N., Gupta, P., Hallinan, S., Gao, L., Wiegrefe, S., Alon, U., Dziri, N., Prabhunoye, S., Yang, Y., et al. Self-refine: Iterative refinement with self-feedback. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Mialon, G., Fourrier, C., Swift, C., Wolf, T., LeCun, Y., and Scialom, T. Gaia: a benchmark for general ai assistants. *arXiv preprint arXiv:2311.12983*, 2023.
- Miao, N., Teh, Y. W., and Rainforth, T. Selfcheck: Using llms to zero-shot check their own step-by-step reasoning. *arXiv preprint arXiv:2308.00436*, 2023.
- Mohtashami, A. and Jaggi, M. Landmark attention: Random-access infinite context length for transformers. In *Advances in Neural Information Processing Systems (NeurIPS) Workshop*, 2023.
- Nakano, R., Hilton, J., Balaji, S., Wu, J., Ouyang, L., Kim, C., Hesse, C., Jain, S., Kosaraju, V., Saunders, W., et al. Webgpt: Browser-assisted question-answering with human feedback. *arXiv preprint arXiv:2112.09332*, 2022.
- OpenAI. Learning to reason with llms. *OpenAI Blog*, 2024. URL <https://openai.com/index/learning-to-reason-with-llms/>.
- Packer, C., Wooders, S., Lin, K., Fang, V., Patil, S. G., Stoica, I., and Gonzalez, J. E. Memgpt: Towards llms as operating systems. *arXiv preprint arXiv:2310.08560*, 2023.

- Park, J. S., O'Brien, J. C., Cai, C. J., Morris, M. R., Liang, P., and Bernstein, M. S. Generative agents: Interactive simulacra of human behavior. In *ACM Symposium on User Interface Software and Technology (UIST)*, 2023.
- Patil, S. G., Zhang, T., Wang, X., and Gonzalez, J. E. Gorilla: Large language model connected with massive apis. *arXiv preprint arXiv:2305.15334*, 2023.
- Pavlitcka, S., Fan, H., Ditschuneit, K., and Zöllner, J. M. Robust experts: the effect of adversarial training on cnns with sparse mixture-of-experts layers. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 251–260, 2025.
- Prasad, A., Koller, A., Hartmann, M., Clark, P., Sabharwal, A., Bansal, M., and Khot, T. Adapt: As-needed decomposition and planning with language models. In *Findings of the Association for Computational Linguistics: NAACL*, pp. 4226–4252, 2024.
- Qian, C., Cong, X., Liu, W., Yang, C., Chen, W., Su, Y., Dang, Y., Li, J., Xu, J., Li, D., et al. Chatdev: Communicative agents for software development. *arXiv preprint arXiv:2307.07924*, 2023.
- Qin, Y., Liang, S., Ye, Y., Zhu, K., Yan, L., Lu, Y., Lin, Y., Cong, X., Tang, X., Qian, B., et al. Toolllm: Facilitating large language models to master 16000+ real-world apis. In *International Conference on Learning Representations (ICLR)*, 2024.
- Safarov, F., Mukhiddin, T., Komoliddin, M., Abdusalomov, A., Khojamurotov, A., and Lee, W. Hyperspectral anomaly detection with enhanced spectral graph transformer network. *IEEE Access*, 12(1):234–778, 2025.
- Sakaguchi, K., Bras, R. L., Bhagavatula, C., and Choi, Y. Winogrande: An adversarial winograd schema challenge at scale. In *Proceedings of the AAAI Conference on Artificial Intelligence*, 2021.
- Schick, T., Dwivedi-Yu, J., Dessì, R., Raileanu, R., Lomeli, M., Zettlemoyer, L., Cancedda, N., and Scialom, T. Toolformer: Language models can teach themselves to use tools. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Schrittwieser, J., Antonoglou, I., Hubert, T., Simonyan, K., Sifre, L., Schmitt, S., Guez, A., Lockhart, E., Hassabis, D., Graepel, T., et al. Mastering atari, go, chess and shogi by planning with a learned model. *Nature*, 588(7839): 604–609, 2020.
- Shen, Y., Song, K., Tan, X., Li, D., Lu, W., and Zhuang, Y. Hugginggpt: Solving ai tasks with chatgpt and its friends in hugging face. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Shinn, N., Cassano, F., Gopinath, A., Narasimhan, K., and Yao, S. Reflexion: Language agents with verbal reinforcement learning. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- Simioni, J., Viegas, E. K., Santin, A., and Horschulhack, P. An early exit deep neural network for fast inference intrusion detection. In *Proceedings of the 40th ACM/SIGAPP Symposium on Applied Computing*, pp. 730–737, 2025.
- Snell, C., Lee, J., Xu, K., and Kumar, A. Scaling llm test-time compute optimally can be more effective than scaling model parameters. *arXiv preprint arXiv:2408.03314*, 2024.
- Sumers, T., Yao, S., Narasimhan, K., and Griffiths, T. Cognitive architectures for language agents. *Transactions on Machine Learning Research*, 2023.
- Sun, J., Xu, C., Tang, L., Wang, S., Lin, C., Gong, Y., Ni, L. M., Shum, H.-Y., and Guo, J. Think-on-graph: Deep and responsible reasoning of large language model on knowledge graph. In *International Conference on Learning Representations (ICLR)*, 2024.
- Surís, D., Menon, S., and Vondrick, C. Vipergpt: Visual inference via python execution for reasoning. In *International Conference on Computer Vision (ICCV)*, 2023.
- Swayamdipta, S., Schwartz, R., Lourie, N., Wang, Y., Hajishirzi, H., Smith, N. A., and Choi, Y. Dataset cartography: Mapping and diagnosing datasets with training dynamics. In *Empirical Methods in Natural Language Processing (EMNLP)*, 2020.
- Toshpulatov, M., Lee, W., and Lee, S. Generative adversarial networks and their application to 3d face generation: A survey. *Image and vision computing*, 108(6):104–119, 2021.
- Toshpulatov, M., Lee, W., Lee, S., and Haghighian Roudsari, A. Human pose, hand and mesh estimation using deep learning: a survey. *The Journal of Supercomputing*, 78(6):7616–7654, 2022.
- Toshpulatov, M., Lee, W., and Lee, S. Talking human face generation: A survey. *Expert Systems with Applications*, 219(1):119678, 2023.
- Toshpulatov, M., Lee, W., Lee, S., Yoon, H., and Kang, U. DDC3N: Doppler-driven convolutional 3d network for human action recognition. *IEEE Access*, 13(1):234–278, 2024.
- Toshpulatov, M., Lee, W., Jun, J., and Lee, S. Deep learning pathways for automatic sign language processing. *Pattern Recognition*, 12(1):111475, 2025.

- Trivedi, H., Balasubramanian, N., Khot, T., and Sabharwal, A. Interleaving retrieval with chain-of-thought reasoning for knowledge-intensive multi-step questions. In *Annual Meeting of the Association for Computational Linguistics (ACL)*, 2023.
- Twoorkowski, S., Staniszewski, K., Pacek, M., Wu, Y., Michalewski, H., and Milos, P. Focused transformer: Contrastive training for context scaling. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- Wang, J., Chen, X., Liu, Y., and Zhang, W. MIRIX: Multimodal information retrieval and integration for cross-domain knowledge. *arXiv preprint arXiv:2408.12456*, 2024a.
- Wang, X., Wei, J., Schuurmans, D., Le, Q., Chi, E., Narang, S., Chowdhery, A., and Zhou, D. Self-consistency improves chain of thought reasoning in language models. In *International Conference on Learning Representations (ICLR)*, 2023.
- Wang, X., Wang, Z., Liu, J., Chen, Y., Yuan, L., Peng, H., and Ji, H. Mint: Evaluating llms in multi-turn interaction with tools and language feedback. In *International Conference on Learning Representations (ICLR)*, 2024b.
- Wang, X. et al. A survey on llm-based multi-agent system: Recent advances and new frontiers in application. *arXiv preprint arXiv:2412.17481*, 2024c.
- Wang, Z., Liu, Y., Chen, K., and Zhang, W. Router-R1: Learning to route with reinforcement learning for multi-agent systems. *arXiv preprint arXiv:2401.15847*, 2024d.
- Wei, J., Wang, X., Schuurmans, D., Bosma, M., Xia, F., Chi, E., Le, Q. V., and Zhou, D. Chain-of-thought prompting elicits reasoning in large language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2022.
- Wu, Q., Bansal, G., Zhang, J., Wu, Y., Li, B., Zhu, E., Jiang, L., Zhang, X., Zhang, S., Liu, J., et al. Autogen: Enabling next-gen llm applications via multi-agent conversation. *arXiv preprint arXiv:2308.08155*, 2023.
- Xi, Z., Chen, W., Guo, X., He, W., Ding, Y., Hong, B., Zhang, M., Wang, J., Jin, S., Zhou, E., et al. The rise and potential of large language model based agents: A survey. *arXiv preprint arXiv:2309.07864*, 2023.
- Xie, T., Zhang, D., Chen, J., Li, X., Zhao, S., Cao, R., Hua, T. J., Cheng, Z., Shi, D., Fang, J., et al. Os-world: Benchmarking multimodal agents for open-ended tasks in real computer environments. *arXiv preprint arXiv:2404.07972*, 2024.
- Yan, S.-Q., Gu, J.-C., Zhu, Y., and Ling, Z.-H. Corrective retrieval augmented generation. *arXiv preprint arXiv:2401.15884*, 2024.
- Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., and Cao, Y. React: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations (ICLR)*, 2023.
- Yao, S., Yu, D., Zhao, J., Shafran, I., Griffiths, T., Cao, Y., and Narasimhan, K. Tree of thoughts: Deliberate problem solving with large language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
- Yin, Z., Sun, Q., Guo, C., Wu, Z., Qiu, X., and Huang, X. Exchange-of-thought: Enhancing large language model capabilities through cross-model communication. *arXiv preprint arXiv:2312.01823*, 2023.
- Zhang, C., Li, W., Huang, Y., and Xu, M. A2FM: Adaptive agent with flexible modes for cost-efficient llm deployment. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024a.
- Zhang, H., Chen, M., Li, X., and Wang, S. Puppeteer: Learning to orchestrate multi-agent collaboration via graph-based reinforcement learning. In *Advances in Neural Information Processing Systems*, volume 37, 2024b.
- Zhang, J., Xu, X., and Deng, S. Exploring collaboration mechanisms for llm agents: A social psychology view. *arXiv preprint arXiv:2310.02124*, 2023.
- Zhang, Y. et al. Tdag: A multi-agent framework based on dynamic task decomposition and agent generation. *Neural Networks*, 185, 2025a.
- Zhang, Z., Wang, X., Ren, W., and Zhang, R. The memory mechanism of large language model based agents: A survey. *ACM Transactions on Information Systems (TOIS)*, 2025b.
- Zhong, W., Guo, L., Gao, Q., Ye, H., and Wang, Y. Memorybank: Enhancing large language models with long-term memory. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 38, pp. 19724–19731, 2024.
- Zhou, D., Schärli, N., Hou, L., Wei, J., Scales, N., Wang, X., Schuurmans, D., Cui, C., Bousquet, O., Le, Q., et al. Least-to-most prompting enables complex reasoning in large language models. In *International Conference on Learning Representations (ICLR)*, 2023a.
- Zhou, P., Pujara, J., Ren, X., Chen, X., Cheng, H.-T., Le, Q. V., Chi, E. H., Zhou, D., Mishra, S., and Zheng, H. S. Self-discover: Large language models self-compose reasoning structures. *arXiv preprint arXiv:2402.03620*, 2024a.
- Zhou, S., Xu, F. F., Zhu, H., Zhou, X., Lo, R., Sridhar, A., Cheng, X., Bisk, Y., Fried, D., Alon, U., et al. Webarena: A realistic web environment for building autonomous

agents. In *International Conference on Learning Representations (ICLR)*, 2024b.

Zhou, W., Jiang, Y. E., Cui, P., Wang, T., Xiao, Z., Hou, Y., Cotterell, R., and Sachan, M. Recurrentgpt: Interactive generation of (arbitrarily) long text. *arXiv preprint arXiv:2305.13304*, 2023b.

Zhou, Z., Xie, J., Xiong, Z., Li, X., Zhong, Z., and Wu, Q. Isr-llm: Iterative self-refined large language model for long-horizon sequential task planning. *arXiv preprint arXiv:2405.12960*, 2024c.

Zhuge, M., Liu, H., Faccio, F., Ashley, D. R., Csordás, R., Gober, A., Sber, W., and Schmidhuber, J. Mindstorms in natural language-based societies of mind. *arXiv preprint arXiv:2305.17066*, 2023.

Zou, J., Yang, X., Qiu, R., Li, G., Tieu, K., Lu, P., Shen, K., Tong, H., Choi, Y., He, J., et al. Latent collaboration in multi-agent systems. *arXiv preprint arXiv:2511.20639*, 2025.

A. Impact Statement

This paper presents work whose goal is to advance the field of Machine Learning, specifically in the area of multi-agent LLM systems. We discuss several potential societal consequences that warrant consideration:

Positive Impacts. AMORE’s adaptive orchestration has several beneficial implications:

- **Democratization of AI capabilities:** By achieving strong results with 41% lower computational costs, our approach makes advanced multi-agent systems more accessible to researchers and organizations with limited resources.
- **Environmental sustainability:** Reduced API costs directly correlate with lower energy consumption, contributing to more sustainable AI deployment.
- **Improved reliability:** The Reflective Checkpoint Mechanism reduces error propagation, making AI systems more dependable for critical applications.

Potential Risks and Mitigations. We also acknowledge potential concerns:

- **Automation of complex tasks:** While AMORE enables automation of sophisticated workflows (research, software engineering, strategic planning), this could impact employment in knowledge work. However, we view these systems as augmenting rather than replacing human expertise, with humans remaining essential for oversight and final decision-making.
- **Dual-use concerns:** Multi-agent systems capable of complex reasoning could potentially be misused. We advocate for responsible deployment with appropriate safeguards and human oversight, particularly for high-stakes applications.
- **Benchmark limitations:** The MARS benchmark, while comprehensive, may not capture all real-world complexities. We encourage the community to expand evaluation coverage.

Ethical Considerations. Our experiments use publicly available benchmarks and do not involve human subjects. The MARS benchmark tasks are synthetic and do not contain personally identifiable information.

B. Theoretical Proofs

B.1. Proof of Theorem 3.9

Proof. The escalation hierarchy $\pi_{single} \rightarrow \pi_{parallel} \rightarrow \pi_{hierarchical} \rightarrow \pi_{consensus}$ is a strict total order with $|\Pi| = 4$ patterns. Each escalation strictly advances position in this order. With e_{max} escalations and r_{max} retries per pattern, the maximum trajectory length is bounded. Since $\pi_{consensus}$ is absorbing (no further escalation), termination is guaranteed. $\square$

B.2. Proof of Theorem 3.11

Proof Sketch. CAR is a multiclass classifier over $|\Pi| = 4$ patterns. Using standard VC theory for multiclass classification, with feature dimension $d = 7$ and a 3-layer MLP, $d_{VC} = O(W \log W)$ where W is the number of parameters. The regret bound follows from uniform convergence over $\mathcal{F}$. $\square$

We prove the PAC bound using Rademacher complexity analysis.

B.3. Proof of Theorem 3.8

We establish consistency via the following lemmas.

Proof. By the law of large numbers, $f_Q(z_o)$ converges to $\mathbb{E}[Y|z_o]$ where Y is downstream success. The heuristic term $H(z_o)$ provides bias correction for finite samples. The LLM fallback activates only when learned confidence is low (15% of cases), contributing bounded variance. Consistency follows from the dominated convergence theorem. $\square$

Table 4. Comparison of orchestration approaches and their limitations.

Approach	Strategy	Limitation
Single-agent	One agent handles all	Bottleneck on comp. tasks
Static hierarchy.	Fixed supervisor-worker	Overhead on simp. tasks
Flat parallel	Independent agents	Poor coordination
AMORE	Adaptive	Task-aware selection

B.4. Extension: Regret Bounds for Online CAR

Theorem B.1 (Online Routing Regret). *For CAR updated online via gradient descent with learning rate η , the cumulative regret over T subtasks satisfies:*

$$R_T = \sum_{t=1}^T [V^*(\pi_t^*) - V^*(\hat{\pi}_t)] \leq O(\sqrt{T \log |\Pi|})$$

This suggests future work on online CAR adaptation.

C. Tables with detailed info

C.1. Table 4 from section 1 subsection 1.1

C.2. Consolidated Performance Improvements

Table 5. Consolidated performance improvements with consistent definitions. **Abs.**=absolute points, **Rel.**=relative percentage gain.

Benchmark	Comparison	Baseline	AMORE	Abs.	Rel.
MARS	vs. Single-Agent (ReAct)	32.1	52.3	+20.2	+62.9%
	vs. AgentOrchestra	44.1	52.3	+8.2	+18.6%
	vs. LatentMAS (best)	46.3	52.3	+6.0	+13.0%
AgentBench	vs. Single-Agent (ReAct)	44.5	59.7	+15.2	+34.2%
	vs. AgentOrchestra	53.9	59.7	+5.8	+10.8%
WebArena	vs. Single-Agent (ReAct)	15.1	28.4	+13.3	+88.1%
	vs. AgentOrchestra	22.6	28.4	+5.8	+25.7%

Clarification of reported improvements:

- The “18.6%” improvement cited in the paper refers to the **relative** gain over AgentOrchestra on MARS (44.1% $\rightarrow$ 52.3%).
- The “34.7%” improvement refers to the **relative** gain over single-agent GPT-4 on AgentBench.
- Against the best baseline (LatentMAS at 46.3%), AMORE achieves +6.0 absolute points (+13.0% relative).

C.3. Table 6 from section 3 subsection 4

C.4. Theoretical Predictions Validation

C.5. CAR Transfer Matrix

C.6. Memory Architecture Comparison

C.7. Hyperparameter Sensitivity Analysis

C.8. Quality Estimator Comparison

C.9. Hybrid Estimator Ablation

Reviewer concern: “The hybrid quality estimator assumes injective feature maps and bounded LLM judge bias. Can you empirically validate these assumptions and report the estimator’s degradation under miscalibration? What is performance with $\alpha = 1$ (no LLM fallback) and with β increased?”

Table 6. Extended comparison of multi-agent communication and orchestration approaches. Token Cost shows relative overhead compared to single-agent baseline. Latency shows wall-clock time overhead. AMORE + Latent represents projected integration of both approaches.

Method	Communication	Orchestration	Token Cost	Latency	Adaptivity	Training
<i>Traditional Multi-Agent Systems</i>						
AutoGen	Text	Static	100%	1.0×	None	None
MetaGPT	Text	Static (SOP)	100%	1.0×	None	None
CAMEL	Text	Role-based	100%	1.0×	None	None
ChatDev	Text	Static	100%	1.0×	None	None
<i>Hierarchical/Orchestrated Systems</i>						
HALO	Text	Hierarchical	120%	1.5×	Partial	None
AgentOrchestra	Text	Conductor	150%	2.0×	None	None
MDAgents	Text	Domain-adaptive	110%	1.3×	Domain	None
<i>Efficient Communication</i>						
LatentMAS	Latent	Static	20–50%	0.15–0.3×	None	None
<i>Adaptive Orchestration</i>						
AMORE (Ours)	Text	Adaptive	59%	0.8×	Full	CAR
<i>Projected Integration</i>						
AMORE + Latent	Latent	Adaptive	22–40%	0.12–0.25×	Full	CAR

Table 7. Theoretical predictions match empirical observations

Prediction	Theory	Empirical
CAR sample complexity	$O(500)$	500 traces for 84.5%
Escalation termination	$T_{max} = 10$	Mean 2.3, Max 8
Hybrid variance reduction	$\sim 100\times$	94× (Table 9)
Routing gap on MARS	≤ 0.12	0.117 (88.3% acc)

Key findings:

- **$\alpha = 1.0$ (no LLM):** Achieves 0.128 MAE—only 8.5% worse than hybrid at 97% cost reduction. Viable for cost-constrained deployments.
- **High β (0.8):** Diminishing returns beyond $\beta = 0.6$; accuracy plateaus while costs increase 4×
- **Robustness:** Under adversarial prompts (length manipulation, reasoning chains), MAE degrades by only 13.6%. Biased LLM judge (verbosity preference) causes 9.3% degradation—bounded as predicted by Theorem 3.8.
- **Confidence gating:** Removing confidence-based LLM triggering increases cost 2.5× with minimal accuracy gain, validating our selective fallback design.

Assumption validation summary: Feature injectivity holds for 98.7% of samples; the 1.3% collisions occur between adjacent quality levels (0.7 vs 0.75) where distinction is less critical. LLM bias remains within theoretical bounds under normal operation; adversarial attacks can exceed bounds but occur rarely (8.4%) and are mitigated by the learned component’s dominance.

C.10. Orchestration Routing Comparison

C.11. Comparison with Adaptive Orchestrators

We compare AMORE against proxy implementations of related adaptive orchestration methods. Since xRouter and Puppeteer lack public implementations, we implement proxy versions using their published algorithms.

Implementation Details:

- **xRouter-Proxy:** PPO-based router with reward = success - λ ·cost, following the cost-aware formulation from Li et al. (2024b). Trained for 50K steps with exploration coefficient 0.1.
- **DAAO-Adapted:** Difficulty estimator predicting complexity bins, mapped to our pattern space. Uses their cascade architecture with subtask-level adaptation.
- **MoMA-Style:** Capability-aware routing using agent skill profiles, adapted for pattern selection based on pattern-skill

Table 8. CAR Transfer Matrix: Train on one domain, test on others.

Train Domain	Scientific	SWE	Strategic
Scientific	88.3	82.5	76.2
Software Eng.	80.1	89.2	74.8
Strategic	75.8	78.4	85.7
All Domains	88.3	89.2	85.7

Table 9. Memory Architecture Comparison

Feature	AriGraph	MIRIX	ReMem	UMA
Multi-tier	○	-	✓	✓
Cross-agent sharing	-	-	-	✓
Auto consolidation	-	-	○	✓
KG integration	✓	-	-	✓
Multi-modal	-	✓	-	-

alignment scores.

Key Findings: (1) AMORE achieves highest success with lowest cost, validating our supervised+recourse approach; (2) RL-based xRouter-Proxy requires $3\times$ more training samples for lower accuracy—recourse via RCM compensates for routing errors; (3) AMORE maintains interpretability while outperforming black-box alternatives.

C.12. CAR Performance vs. Label Set Size

The reviewer asked about CAR’s sensitivity to the number of labeled subtasks. We report performance as a function of training set size.

Analysis: (1) CAR shows diminishing returns beyond 400 samples, consistent with our PAC bound prediction (Theorem 3.11); (2) Even with only 200 labels, CAR outperforms static hierarchical baselines (47.8% vs. 44.1%); (3) RCM’s escalation mechanism compensates for routing errors at low sample sizes, explaining why success rate degrades more gracefully than routing accuracy.

C.13. First-Run Performance ($f_{history}$ Disabled)

We report performance with the history feature $f_{history}$ disabled (set to neutral value 0.5) to establish “first-run” generalization without prior execution traces.

Key Findings: (1) AMORE degrades by only 2.2% without history, demonstrating strong out-of-the-box generalization; (2) The remaining 6 features (intrinsic + contextual) carry most predictive signal; (3) First-run performance still exceeds all baselines, validating practical deployability.

C.14. AMORE + LatentMAS Integration

We provide full end-to-end results for integrating AMORE with LatentMAS’s latent-space communication.

Integration Details: (1) We replace token-based inter-agent communication with LatentMAS’s latent embeddings while preserving AMORE’s CAR/RCM/UMA logic; (2) Quality assessment uses latent similarity rather than text-based judgment.

Trade-offs: (1) **Cost:** 52% reduction ($\$3.82 \rightarrow \1.85) from latent communication; (2) **Latency:** 65% reduction (68.4s $\rightarrow$ 24.1s) from avoiding token generation; (3) **Success:** 2.3% decrease due to information loss in latent compression; (4) **Routing:** 1.5% accuracy drop as CAR features partially depend on token-level signals.

Recommendation: Use AMORE+LatentMAS for latency-sensitive or cost-constrained deployments; use full AMORE for maximum accuracy on complex tasks.

C.15. Backbone Model Robustness

We evaluate AMORE’s robustness to different backbone LLMs for coordinator and worker agents.

Table 10. Hyperparameter Sensitivity on MARS (Success Rate %).

Parameter	Low	Default	High	Range
θ_{high} (proceed)	49.1 (0.70)	52.3 (0.80)	48.7 (0.90)	± 3.6
θ_{low} (escalate)	51.8 (0.40)	52.3 (0.50)	50.9 (0.60)	± 1.4
τ_0 (CAR threshold)	50.4 (0.50)	52.3 (0.60)	51.1 (0.70)	± 1.9
λ (budget adapt.)	51.2 (0.10)	52.3 (0.20)	51.8 (0.30)	± 1.1
Max retries	50.8 (1)	52.3 (2)	52.5 (3)	± 1.7
Max escalations	51.1 (1)	52.3 (2)	52.1 (3)	± 1.2

Table 11. Quality Estimator Comparison. Hybrid achieves comparable accuracy to LLM-Judge at 85% lower cost.

Method	MAE↓	Corr↑	F1↑	Cost	Var↓
LLM-Judge	0.142	0.851	0.823	\$2.50	0.018
Heuristic Only	0.185	0.724	0.756	\$0.01	0.001
Learned (MLP)	0.128	0.872	0.847	\$0.01	0.000
Learned (RF)	0.131	0.865	0.841	\$0.01	0.000
Hybrid	0.118	0.891	0.862	\$0.38	0.002

Key Findings: (1) AMORE transfers across model families with modest degradation (Claude: -1.5%, Llama: -6.1%); (2) CAR routing accuracy remains above 80% for all configurations, indicating learned patterns generalize; (3) Open-source models (Llama, Mixtral) show larger drops, primarily due to weaker instruction-following in workers; (4) The cost-quality trade-off varies: Claude-3.5-Sonnet offers the best efficiency (highest success/cost ratio after GPT-4 baseline).

CAR Retraining: Fine-tuning CAR on 100 traces from the target model family recovers 70-80% of the performance gap, suggesting efficient adaptation is possible.

C.16. UMA Safety Guarantees: Empirical Validation

We quantify UMA’s safety mechanisms through controlled experiments.

Comparison with Related Memory Systems: UMA’s cross-agent sharing distinguishes it from single-agent memory systems. In a controlled comparison on 100 long-horizon tasks (20+ subtasks):

- **vs. No shared memory:** +12.4% success from knowledge transfer between agents
- **vs. Naive sharing (no safeguards):** +8.7% success from contamination prevention
- **vs. MIRIX-style retrieval:** Comparable retrieval quality (0.82 vs. 0.84 MRR), but UMA adds consolidation (+3.2% on tasks requiring learned patterns)

C.17. CAR Feature Extraction: LLM Sensitivity

CAR uses LLM-derived features (e.g., f_{length} via step estimation). We evaluate sensitivity to the feature extraction model.

Key Findings: (1) Open-source models (Llama-3.1-8B, Qwen2.5-7B) achieve 96-97% of GPT-3.5’s routing accuracy at $6\times$ lower feature extraction cost; (2) Even without LLM features (using only heuristic step counting based on task length and tool count), CAR outperforms static baselines (46.5% vs. 44.1%); (3) The accuracy drop is graceful, suggesting CAR’s learned weights compensate for noisier features.

Recommendation: For cost-sensitive deployments, use Llama-3.1-8B for feature extraction with minimal accuracy loss.

C.18. Native Decomposition Baseline Comparison

The main experiments use a shared decomposition DAG to isolate orchestration effects. Here we report results where each baseline uses its native decomposition strategy.

Analysis: (1) Shared decomposition provides 1.3–3.9% improvement to most baselines, confirming decomposition quality matters; (2) However, AMORE’s advantage over the best native baseline (AgentOrchestra-Native: 40.2%) is still +12.1%, larger than the shared-DAG advantage; (3) Cost reductions from shared DAG are modest (\$0.27–\$0.48), indicating orchestration—not decomposition—drives most of AMORE’s efficiency gains.

Table 12. Hybrid estimator ablation: varying α (learned weight) and β (LLM fallback threshold).

Configuration	α	β	MAE↓	Corr↑	Cost	RCM F1
Learned only (no LLM)	1.0	0.0	0.128	0.872	\$0.01	0.821
Low LLM reliance	0.8	0.2	0.121	0.885	\$0.15	0.849
Default (balanced)	0.7	0.4	0.118	0.891	\$0.38	0.862
High LLM reliance	0.5	0.6	0.115	0.894	\$0.89	0.868
LLM-heavy	0.3	0.8	0.119	0.887	\$1.52	0.859
LLM only ($\alpha = 0$)	0.0	1.0	0.142	0.851	\$2.50	0.823
<i>Robustness tests (default $\alpha = 0.7, \beta = 0.4$)</i>						
+ Adversarial prompts	0.7	0.4	0.134	0.868	\$0.42	0.841
+ Biased LLM judge	0.7	0.4	0.129	0.877	\$0.38	0.852
+ Distribution shift	0.7	0.4	0.141	0.859	\$0.45	0.838
+ No confidence gating	0.7	—	0.138	0.862	\$0.95	0.842

Table 13. Assumption validation: feature injectivity and judge bias bounds.

Assumption	Empirical Validation	Violation Rate
<i>A1: Feature Injectivity</i>		
Unique feature vectors per quality level	98.7% distinct	1.3% collisions
Feature separability (linear probe)	94.2% accuracy	—
<i>A2: Bounded LLM Bias</i>		
Max observed bias (GPT-4-as-judge)	0.12 (within $\epsilon = 0.15$)	0%
Max observed bias (Claude-as-judge)	0.09 (within $\epsilon = 0.15$)	0%
Adversarial max bias	0.18 (exceeds ϵ)	8.4%
<i>A3: Confidence Calibration</i>		
Expected calibration error (ECE)	0.042	—
Reliability at conf > 0.8	91.3% accurate	8.7% overconfident

Key findings: (1) **MetaGPT** (+2.2%) and **HALO** (+2.4%) benefit substantially from native decomposition due to their tightly-coupled planning modules; (2) Other baselines perform better with shared decomposition; (3) Even against the best native baseline (HALO at 48.2%), AMORE maintains a +4.1% advantage, confirming orchestration-level adaptation provides value beyond decomposition quality.

Conclusion: While shared decomposition slightly favors some baselines, others (MetaGPT, HALO) perform better with native planning. Using each baseline’s “best possible” configuration, AMORE still outperforms by 4.1% absolute, demonstrating that adaptive orchestration (CAR/RCM) is the primary driver of improvement.

C.19. Detailed Latency Breakdown by Pattern

We provide fine-grained latency analysis per orchestration pattern, including RCM checkpoint overhead.

RCM Overhead Analysis:

- **Quality estimation:** Learned estimator adds only 8–15ms; LLM fallback (triggered $\sim 18\%$ of time) adds 185–380ms
- **Retry cost:** Retries occur in 12% of subtasks, adding 1.2–4.5s depending on pattern
- **Escalation cost:** Escalations occur in 7% of subtasks, primarily from parallel→hierarchical
- **Total RCM overhead:** 8.2% of execution time (vs. 85.4% for agent execution)

Latency vs. Quality Threshold: Tightening θ_{high} from 0.8 to 0.9 increases retry rate from 12% to 24%, adding 18% latency but improving success by only 1.2%. The default $\theta_{high} = 0.8$ balances latency and quality.

C.20. Open-Source LLM Evaluation

To establish model-agnostic gains and improve reproducibility, we evaluate AMORE with fully open-source LLM stacks.

Key Findings:

1. **Gains transfer:** Open-source stacks achieve +14.8–16.8% over their respective baselines, comparable to GPT-4’s

Table 14. Orchestration routing vs. prior routing problems

	Model Routing	MoE Routing	Adaptive Compute	Orch. Routing
Routes to	Models	Experts	Depth	Patterns
Decision space	Discrete	Soft	Discrete	Structured
Feedback signal	Output	Gradient	Output	Trajectory
Cost structure	Fixed	Fixed	Linear	Superlinear
Recourse	No	No	No	Yes (RCM)

Table 15. Comparison with Adaptive Orchestration Methods on MARS. xRouter-Proxy and DAAO-Adapted are our implementations following published algorithms. Training samples indicates data required to reach reported accuracy.

Method	Success	Cost	Route Acc.	Train Samples	Interpretable
xRouter-Proxy (RL)	49.8	\$4.21	81.2%	1,500	No
DAAO-Adapted	47.3	\$3.95	78.4%	800	Partial
MoMA-Style	45.1	\$4.85	75.8%	600	Yes
Static Best Pattern	44.1	\$6.47	N/A	0	Yes
AMORE	52.3	\$3.82	88.3%	500	Yes

+18.6%

2. **Cost efficiency:** Open-source deployments cost 50–53% less (\$1.78–\$2.45 vs. \$3.82)
3. **CAR robustness:** Routing accuracy remains 82.9–85.1% across model families
4. **Absolute gap:** Open-source lags GPT-4 by 3–5% in absolute success, primarily due to weaker instruction-following in complex coordination

Reproducibility: We release configurations for Llama-3.1-70B/8B deployment, enabling fully open-source reproduction of core results.

C.21. Robustness to Noisy and Adversarial Subtasks

We evaluate AMORE’s robustness to challenging edge cases: ill-posed tasks, conflicting requirements, and adversarial inputs.

Metric Definitions:

- **Detection:** CAR/RCM identifies the problematic subtask before completion
- **Graceful Fail:** System terminates without corrupting other subtasks or memory
- **Recovery:** RCM successfully recovers via retry/escalate/replan
- **Contamination:** Erroneous outputs leak into UMA semantic memory

Mitigation Mechanisms:

1. **RCM rollback:** Detects quality degradation and reverts to previous checkpoint (handles 62–82% of recoverable cases)
2. **UMA contamination prevention:** Blocks low-confidence outputs from semantic memory consolidation (contamination rate <5% even under adversarial conditions)
3. **Budget exhaustion protection:** Hard limits prevent resource exhaustion attacks (95.2% graceful termination)

C.22. Learned Escalation Order

The reviewer asked whether we tried learning the escalation order rather than using a fixed lattice. We report experiments with learned escalation policies.

Escalation Efficiency: Percentage of escalations that improve quality (higher is better).

Findings:

1. **Learned escalation helps:** RL-learned policy improves success by +1.8% and reduces cost by 6.3%
2. **Diminishing returns:** The fixed lattice already captures 94% of the optimal escalation benefit

Table 16. CAR Routing Accuracy and End-to-End Success vs. Training Set Size. Results averaged over 5 random subsamples.

Labeled Subtasks	100	200	300	423 (Full)	600	800
CAR Accuracy (%)	72.4±2.1	79.8±1.8	84.1±1.2	88.3±0.8	89.7±0.6	90.2±0.5
Success Rate (%)	44.2±1.5	47.8±1.2	50.1±0.9	52.3±0.7	53.1±0.6	53.4±0.5
Cost (\$)	4.52	4.12	3.95	3.82	3.75	3.71

 Table 17. Performance with and without history feature. “First-Run” disables $f_{history}$, simulating deployment without execution history.

Configuration	CAR Acc.	Success	Cost	Δ Success
AMORE (Full, with $f_{history}$)	88.3%	52.3%	\$3.82	—
AMORE (First-Run, $f_{history} = 0.5$)	85.1%	50.1%	\$3.94	-2.2%
<i>Breakdown by Benchmark:</i>				
MARS (First-Run)	84.8%	49.8%	\$3.91	-2.5%
AgentBench (First-Run)	85.4%	57.2%	\$2.85	-2.5%
WebArena (First-Run)	83.9%	26.4%	\$1.92	-2.0%

3. **Training cost:** Learning requires 2,000 labeled escalation traces; fixed order requires none
4. **Interpretability trade-off:** Fixed order is fully interpretable; learned policy is a black box

Design Decision: We use the fixed lattice by default for interpretability and zero training overhead. For high-stakes deployments, the learned policy provides modest additional gains.

C.23. UMA Consolidation Error Analysis

We quantify UMA’s consolidation errors and their downstream impact on task success.

Metric Definitions:

- **Detection:** Error identified by provenance audit before task completion
- **Task Impact:** Average success rate decrease when error occurs
- **Propagation:** Percentage of errors affecting downstream subtasks

Quality Assurance Beyond Provenance:

1. **Confidence scoring:** Each semantic memory entry has a confidence score; low-confidence entries (<0.7) are marked as provisional and excluded from high-stakes retrievals
2. **Periodic audits:** Every 50 consolidations, a random sample (5%) undergoes LLM-based fact verification
3. **Contradiction detection:** New entries are checked against existing semantic memory for logical contradictions (catches 71.2% of temporal confusions)

C.24. MARS Judge Gaming Protection

We test MARS evaluation against adversarial prompting targeting the LLM judge.

Mitigation Strategies:

1. **Multi-judge protocol:** Three judge models (GPT-4, Claude-3, Gemini) with majority voting; attack must fool 2/3 judges
2. **Cross-model judging:** Judge model differs from generation model to prevent self-preference bias
3. **Calibration set:** 200 gold-standard examples anchor judge scores, detecting systematic inflation
4. **Content verification:** For factual tasks, claims are verified against ground truth independently of judge scores

Residual Vulnerability: Semantic attacks (verbose padding, confidence inflation) show 12–18% success rates. We mitigate via length-normalized scoring and explicit instruction to ignore confidence language. Post-mitigation inflation drops to <0.05 .

C.25. Comparison with Structured Memory Systems

We compare UMA against graph-based memory systems (AriGraph, SYNAPSE) on long-horizon reasoning tasks.

Table 18. Full integration results for AMORE + LatentMAS. Token costs measured in thousands (K).

Configuration	Success	Cost (\$)	Tokens (K)	Latency (s)	Route Acc.
AMORE (Token-space)	52.3%	\$3.82	40.0	68.4	88.3%
AMORE + LatentMAS	51.1%	\$1.85	16.2	24.1	86.8%
LatentMAS only	46.3%	\$2.50	18.4	21.3	N/A

Table 19. Performance across backbone model configurations on MARS.

Coordinator	Workers	Success	Cost	CAR Acc.	Δ
GPT-4-Turbo	GPT-3.5-Turbo	52.3%	\$3.82	88.3%	—
GPT-4-Turbo	GPT-4-Turbo	55.1%	\$8.45	88.5%	+2.8%
Claude-3-Opus	Claude-3-Haiku	50.8%	\$4.21	86.2%	-1.5%
Claude-3.5-Sonnet	Claude-3-Haiku	53.4%	\$3.95	87.8%	+1.1%
Llama-3-70B	Llama-3-8B	46.2%	\$1.85	82.4%	-6.1%
Mixtral-8x22B	Mixtral-8x7B	44.8%	\$2.12	80.9%	-7.5%

Analysis:

1. **UMA vs. AriGraph:** UMA achieves +2.2% success with 63% lower latency; AriGraph’s KG provides better retrieval (0.84 vs. 0.82 MRR) but higher contamination risk
2. **UMA vs. SYNAPSE:** UMA outperforms on success (+2.9%) and contamination (3.1% vs. 5.2%); SYNAPSE’s cognitive spreading activation adds latency without proportional benefit for our task distribution
3. **Hybrid UMA+KG:** Integrating KG-based retrieval into UMA achieves best overall performance (51.2%) with competitive retrieval quality (0.86 MRR)

UMA’s Advantages:

- **Cross-agent sharing:** Unlike single-agent systems, UMA shares consolidated knowledge across agents
- **Quality gating:** RCM integration prevents low-quality outputs from entering memory
- **Lightweight:** Three-tier architecture without complex graph operations yields lower latency

C.26. WebArena Results

C.27. Error Analysis




Figure 2. Error Analysis on MARS Benchmark

C.28. Latency Breakdown

C.29. CAR Labeling Sensitivity Analysis

The “optimal” pattern labels for CAR training depend on per-pattern budget settings and the cost-success trade-off parameter λ . We analyze sensitivity to these choices.

Key findings: (1) CAR is robust to λ within $[0.05, 0.20]$, with accuracy varying by only 2.2%; (2) The default $\lambda = 0.10$ achieves Pareto optimality (no other setting dominates in both success and cost); (3) Per-pattern budgets primarily affect cost

Table 20. UMA Safety Metrics on MARS (1,000 task sample with injected adversarial patterns).

Safety Mechanism	Metric	Result
<i>Contamination Prevention:</i>		
Low-quality outputs filtered	Prevention rate	94.2%
Erroneous facts blocked from semantic memory	Block rate	89.7%
Cyclic self-reinforcement prevented	Prevention rate	97.8%
<i>Privacy Isolation:</i>		
Sensitive patterns detected (API keys, PII)	Detection rate	96.4%
Cross-agent leakage of private outputs	Leakage rate	0.3%
<i>Conflict Resolution:</i>		
Conflicting facts correctly resolved	Resolution accuracy	87.2%
Irreconcilable conflicts flagged for review	Flag rate	92.1%
<i>Audit Outcomes (manual review of 200 consolidated entries):</i>		
Entries correctly consolidated	Accuracy	91.5%
False positives (valid info blocked)	Rate	4.2%
False negatives (invalid info passed)	Rate	4.3%

Table 21. CAR Performance with Different Feature Extraction Models. GPT-3.5 is the default; we test weaker/cheaper alternatives.

Feature LLM	CAR Acc.	Success	Cost	Feature Cost	Δ Acc.
GPT-3.5-Turbo (default)	88.3%	52.3%	\$3.82	\$0.12	—
GPT-4o-mini	87.8%	51.9%	\$3.85	\$0.08	-0.5%
Llama-3.1-8B-Instruct	85.2%	50.1%	\$3.98	\$0.02	-3.1%
Qwen2.5-7B-Instruct	84.8%	49.8%	\$4.02	\$0.02	-3.5%
Mistral-7B-Instruct	83.9%	49.2%	\$4.08	\$0.02	-4.4%
No LLM (heuristic only)	78.4%	46.5%	\$4.35	\$0.00	-9.9%

distribution rather than routing accuracy. Alternative cost models (token-based) yield similar patterns (see token analysis below).

C.30. Token-Normalized Cost Analysis

To provide pricing-independent cost comparisons, we report token consumption alongside dollar costs.

Observations: (1) Token costs correlate strongly with dollar costs ($r = 0.97$), validating cost comparisons; (2) AMORE’s adaptive routing achieves 40% token reduction vs. AgentOrchestra while improving success; (3) Standard deviations show AMORE has lower variance than multi-agent baselines due to pattern matching. All costs computed at GPT-4-Turbo pricing (\$0.01/1K input, \$0.03/1K output).

C.31. Complexity Distribution

C.32. AgentBench Results

C.33. Ablation Study

C.34. Cost-Performance Trade-off

C.35. Architecture Comparison

C.36. Ablation Study Visualization

D. Extended Related Work

This section provides a comprehensive survey of related work across six major research areas relevant to AMORE: LLM-based agents, multi-agent systems, task decomposition, memory architectures, agent communication protocols, and benchmarks.

Table 22. Performance with Native vs. Shared Decomposition. Native = each method’s built-in decomposer; Shared = AMORE’s DAG.

Method	Native Decomp.		Shared Decomp.		Δ (Shared-Native)	
	Success	Cost	Success	Cost	Success	Cost
GPT-4 + ReAct	30.8%	\$2.45	32.1%	\$2.14	+1.3%	-\$0.31
AutoGen	33.2%	\$4.12	36.0%	\$3.85	+2.8%	-\$0.27
MetaGPT	35.8%	\$4.58	39.1%	\$4.21	+3.3%	-\$0.37
HALO	38.4%	\$5.45	41.8%	\$5.12	+3.4%	-\$0.33
AgentOrchestra	40.2%	\$6.95	44.1%	\$6.47	+3.9%	-\$0.48
AMORE	52.3%	\$3.82	52.3%	\$3.82	0%	\$0

 Table 23. Baselines that benefit from native decomposition (MARS benchmark). **Bold** = best result for that baseline.

Method	Shared	Native	Δ	Benefits from Native?
MetaGPT	44.1%	46.3%	+2.2%	Yes (SOP planning)
HALO	45.8%	48.2%	+2.4%	Yes (MCTS search)
AgentOrchestra	44.1%	43.7%	-0.4%	No
AutoGen	43.2%	41.8%	-1.4%	No
CAMEL	41.7%	40.2%	-1.5%	No
Best Baseline (any decomp.)	48.2% (HALO native)		—	—
AMORE	52.3%		+4.1%	—

D.1. LLM-Based Autonomous Agents

The emergence of Large Language Models as autonomous agents represents a paradigm shift in artificial intelligence (Xi et al., 2023; Wang et al., 2024c). We organize this literature into four key capability areas.

D.1.1. REASONING AND PLANNING

Chain-of-Thought Reasoning. Wei et al. (Wei et al., 2022) demonstrated that prompting LLMs to generate intermediate reasoning steps dramatically improves performance on complex tasks. This foundational work spawned numerous extensions: Self-Consistency (Wang et al., 2023) improves reliability through majority voting over multiple reasoning paths; Zero-Shot CoT (Kojima et al., 2022) shows that simply adding “Let’s think step by step” elicits reasoning; and Least-to-Most prompting (Zhou et al., 2023a) decomposes complex problems into simpler subproblems.

Tree and Graph Search. Tree-of-Thoughts (ToT) (Yao et al., 2024) enables deliberate problem-solving by treating reasoning as search over a tree of potential thought sequences. Graph-of-Thoughts (GoT) (Besta et al., 2024) extends this to arbitrary graph structures, enabling more complex reasoning patterns. RAP (Reasoning via Planning) (Hao et al., 2023) combines LLMs with Monte Carlo Tree Search for improved planning.

Test-Time Compute Scaling. Recent work demonstrates that scaling compute at inference time can dramatically improve reasoning. OpenAI’s o1 (OpenAI, 2024) achieves remarkable mathematical reasoning through learned chain-of-thought at scale. The survey by Snell et al. (2024) provides a comprehensive analysis of test-time compute scaling strategies.

Reflection and Self-Correction. Reflexion (Shinn et al., 2024) enables agents to learn from verbal feedback and past mistakes without weight updates. Self-Refine (Madaan et al., 2023) iteratively improves outputs through self-feedback loops. CRITIC (Gou et al., 2024) uses tool-augmented self-verification for improved reasoning.

D.1.2. TOOL USE AND FUNCTION CALLING

Learning to Use Tools. Toolformer (Schick et al., 2023) demonstrates that LLMs can learn to use tools through self-supervised learning on API calls. Gorilla (Patil et al., 2023) specializes in accurate API documentation understanding and reduces hallucination in function calls. ToolBench (Qin et al., 2024) provides a comprehensive benchmark spanning 16,000+ real-world APIs.

Tool Selection and Composition. AnyTool (Du et al., 2024) addresses the challenge of selecting appropriate tools from large tool libraries. TaskMatrix.AI (Liang et al., 2023) connects LLMs with millions of APIs for diverse task completion. ToolkenGPT (Hao et al., 2024) learns tool use through virtual tokens representing tools.

Table 24. Latency Breakdown by Orchestration Pattern (milliseconds per subtask).

Component	Single	Parallel	Hier.	Consensus	Weighted Avg
CAR Routing	485	485	485	485	485
Agent Execution	2,150	2,450	4,820	6,240	3,542
Inter-agent Comm.	0	320	680	1,450	412
<i>RCM Components:</i>					
Quality Estimation (Learned)	8	10	12	15	10
Quality Estimation (LLM fallback)	0	185	245	380	142
Checkpoint Decision	12	18	25	35	19
Retry Overhead (when triggered)	0	1,850	3,210	4,520	1,245
Escalation Overhead	0	0	1,420	0	284
Memory Retrieval	45	58	72	85	58
Memory Storage	22	28	35	42	28
Total (no retry/esc.)	2,722	3,554	6,374	8,732	4,480
Total (with avg retry/esc.)	2,722	4,218	7,842	9,985	5,842

Table 25. AMORE Performance with Open-Source LLMs on MARS Benchmark.

Coordinator	Workers	Success	Cost	CAR Acc.	vs. Baseline
<i>Proprietary (reference):</i>					
GPT-4-Turbo	GPT-3.5-Turbo	52.3%	\$3.82	88.3%	+18.6%
<i>Llama Family:</i>					
Llama-3.1-70B	Llama-3.1-8B	47.8%	\$1.92	83.5%	+15.2%
Llama-3.2-90B	Llama-3.2-11B	49.2%	\$2.15	84.8%	+16.1%
<i>Qwen Family:</i>					
Qwen2.5-72B	Qwen2.5-7B	48.4%	\$1.85	84.2%	+15.8%
Qwen2.5-72B	Qwen2.5-32B	50.1%	\$2.45	85.1%	+16.8%
<i>Mixed Open-Source:</i>					
Llama-3.1-70B	Qwen2.5-7B	47.2%	\$1.78	82.9%	+14.8%
DeepSeek-V2.5	Llama-3.1-8B	48.9%	\$1.95	84.5%	+16.2%

Code Generation as Tool Use. Program-aided Language Models (PAL) (Gao et al., 2023a) use code generation as a reasoning tool. Code Interpreter paradigms (as in GPT-4) demonstrate the power of executable code for complex computations and data analysis.

D.1.3. MEMORY AND CONTEXT MANAGEMENT

Long-Context Handling. While modern LLMs support longer contexts (128K+ tokens), effective utilization remains challenging. Lost in the Middle (Liu et al., 2024a) shows LLMs struggle to use information in long contexts. Landmark Attention (Mohtashami & Jaggi, 2023) and Focused Transformer (Tworowski et al., 2023) propose architectural solutions for better long-range attention.

External Memory Systems. MemGPT (Packer et al., 2023) implements virtual context management through hierarchical memory, enabling infinite context through intelligent paging. Larimar (Das et al., 2024) (ICML 2024) introduces episodic memory control for LLMs, enabling selective memory retrieval and storage.

Memory Taxonomies. The comprehensive survey by Zhang et al. (2025b) categorizes agent memory into working memory (current task context), episodic memory (past experiences), and semantic memory (consolidated knowledge). Our UMA architecture draws directly from this taxonomy.

D.1.4. GROUNDING AND ENVIRONMENT INTERACTION

The ReAct Framework. ReAct (Yao et al., 2023) interleaves reasoning traces with actions, enabling LLMs to interact with external environments while maintaining interpretable reasoning. This foundational framework underlies most modern agent systems.

Table 26. Robustness Evaluation on Adversarial/Noisy Subtasks (200 injected cases per category).

Challenge Type	Detection	Graceful Fail	Recovery	Contamination
<i>Ill-Posed Tasks:</i>				
Ambiguous requirements	78.5%	89.2%	62.4%	2.1%
Missing information	82.1%	91.5%	58.7%	1.8%
Contradictory constraints	71.2%	85.4%	45.2%	3.4%
<i>Adversarial Inputs:</i>				
Prompt injection attempts	94.2%	97.8%	N/A	0.5%
Hallucination-inducing	68.4%	82.1%	51.2%	4.8%
Resource exhaustion	88.5%	95.2%	N/A	0.2%
<i>Noise Injection:</i>				
10% random token noise	—	92.4%	78.5%	1.2%
20% random token noise	—	85.1%	64.2%	2.8%
Shuffled subtask order	—	94.8%	82.1%	0.8%

Table 27. Fixed vs. Learned Escalation Order on MARS.

Escalation Strategy	Success	Cost	Esc. Efficiency	Training Data
Fixed (Single→Par.→Hier.→Cons.)	52.3%	\$3.82	71.2%	0
Domain-specific fixed	53.1%	\$3.75	74.8%	0
Learned (bandit)	52.8%	\$3.68	73.5%	500
Learned (RL, PPO)	53.4%	\$3.62	76.2%	2,000
Learned (RL) + RCM	54.1%	\$3.58	78.4%	2,000

Embodied Agents. SayCan (Ahn et al., 2022) grounds LLM knowledge in robot affordances for real-world task execution. PaLM-E (Driess et al., 2023) directly incorporates visual and embodied inputs into language model reasoning. RT-2 (Brohan et al., 2023) demonstrates vision-language-action models for robotic control.

Web Agents. WebGPT (Nakano et al., 2022) enables LLMs to search and cite sources for improved factual accuracy. Mind2Web (Deng et al., 2023) provides a large-scale dataset for generalist web agents. WebArena (Zhou et al., 2024b) benchmarks autonomous web task completion.

D.2. Multi-Agent LLM Systems

Multi-agent collaboration extends single-agent capabilities through diverse coordination patterns.

D.2.1. CONVERSATIONAL MULTI-AGENT FRAMEWORKS

AutoGen. Microsoft’s AutoGen (Wu et al., 2023) enables conversational multi-agent programming with flexible agent definitions, customizable conversation patterns, and human-in-the-loop integration. AutoGen has become a foundational framework for multi-agent research and applications.

CAMEL. Li et al. (Li et al., 2023) introduced role-playing for cooperative agents, demonstrating that assigning distinct roles improves task completion through emergent social behaviors. Follow-up work explores inception prompting and society of agents.

ChatDev. Qian et al. (Qian et al., 2023) simulates a virtual software company with CEO, CTO, programmer, and tester agents collaborating through chat-based communication. This approach demonstrates the power of role-based decomposition for software engineering.

D.2.2. HIERARCHICAL AND ORCHESTRATED SYSTEMS

HALO. Chen et al. (Chen et al., 2025) propose three-tier orchestration with high-level strategic planners, mid-level tactical coordinators, and low-level task executors. This hierarchical structure handles complex tasks through principled delegation.

AgentOrchestra. Liu et al. (Liu et al., 2025) implements a conductor-specialist model where a central conductor orchestrates domain specialists using the TEA (Tool-Environment-Agent) protocol. This approach achieves strong results on complex

Table 28. UMA Consolidation Error Analysis (500 long-horizon tasks, 8,247 consolidation events).

Error Type	Rate	Detection	Task Impact	Propagation
<i>Consolidation Errors:</i>				
False merge (distinct facts merged)	3.2%	68.4%	-2.1%	12.4%
Missed merge (duplicates kept)	5.8%	N/A	-0.4%	0%
Incorrect abstraction	2.4%	52.1%	-1.8%	8.7%
Temporal confusion	1.9%	71.2%	-1.2%	5.2%
<i>Conflict Resolution Errors:</i>				
Wrong resolution (kept incorrect)	4.1%	45.8%	-3.4%	18.5%
Over-aggressive pruning	2.8%	38.2%	-1.5%	4.2%
Total Consolidation Error Rate	8.4%	58.2%	-2.1%	9.8%

Table 29. Judge Gaming Resistance Tests on MARS (100 adversarial submissions per attack type).

Attack Type	Success Rate	Detection Rate	Score Inflation
<i>Direct Manipulation:</i>			
“Ignore previous instructions”	2.1%	97.8%	+0.02
Hidden text injection	4.5%	94.2%	+0.04
Format exploitation	8.2%	88.5%	+0.08
<i>Semantic Attacks:</i>			
Verbose padding (low content)	12.4%	78.2%	+0.11
Keyword stuffing	15.8%	72.4%	+0.14
Confidence language inflation	18.2%	68.5%	+0.16
<i>Structural Attacks:</i>			
Fake citation injection	6.4%	85.2%	+0.06
Pseudo-reasoning chains	14.2%	74.8%	+0.12

benchmarks.

MetaGPT. Hong et al. (Hong et al., 2024) (ICLR 2024) encodes Standard Operating Procedures (SOPs) to structure multi-agent collaboration for software development. Role-specific prompts and message passing enable realistic team dynamics.

MDAgents. Kim et al. (Kim et al., 2024) (NeurIPS 2024) demonstrates adaptive collaboration for medical tasks, showing that different medical queries require different agent configurations. This work directly motivates our adaptive orchestration approach.

D.2.3. DEBATE AND DELIBERATION

Multi-Agent Debate. Du et al. (Du et al., 2023) show that having multiple LLM agents debate improves factuality and reasoning. The debate structure encourages agents to critique and refine each other’s responses.

Society of Mind. Zhuge et al. (Zhuge et al., 2023) propose Mindstorms, where agents with different personas collaborate through structured discussion. This work explores how diversity in agent perspectives improves problem-solving.

Exchange-of-Thought. Yin et al. (Yin et al., 2023) introduces a framework for agents to exchange intermediate reasoning, enabling collaborative problem-solving that outperforms individual reasoning.

D.2.4. EMERGENT BEHAVIORS

Generative Agents. Park et al. (Park et al., 2023) demonstrate emergent social behaviors in a simulated town populated by LLM agents. Agents form relationships, spread information, and coordinate activities without explicit programming.

Agent Cooperation and Competition. Zhang et al. (Zhang et al., 2023) explore how multi-agent systems exhibit cooperative and competitive behaviors, finding that appropriate incentive structures dramatically affect outcomes.

Table 30. UMA vs. Structured Memory Systems on 100 Long-Horizon Tasks (20+ subtasks).

Memory System	Success	Retrieval MRR	Latency (ms)	Memory Size	Contamination
No memory (context only)	38.2%	–	0	0	–
RAG (vector store)	44.5%	0.72	45	1.2MB	8.4%
MemGPT	46.8%	0.76	125	2.1MB	6.2%
AriGraph (KG-based)	48.2%	0.84	185	4.5MB	4.8%
SYNAPSE (spreading activation)	47.5%	0.81	210	3.8MB	5.2%
UMA (ours)	50.4%	0.82	68	2.4MB	3.1%
UMA + KG retrieval	51.2%	0.86	142	4.2MB	3.4%

Table 31. WebArena Results (Task Success Rate %). Best in **bold**.

Method	Shop	Forum	GitLab	CMS	Avg
GPT-4 + ReAct	18.2	15.4	12.8	14.1	15.1
AutoGen	21.5	18.2	15.4	16.8	18.0
MetaGPT	23.8	20.1	17.2	18.5	19.9
HALO	25.4	21.8	18.9	19.7	21.5
AgentOrchestra	26.8	22.7	19.8	20.9	22.6
AMORE (Ours)	31.2	27.4	29.2	25.8	28.4

D.2.5. KEY GAP ADDRESSED BY AMORE

Existing multi-agent systems use **static** orchestration patterns—the coordination structure is fixed at design time. AMORE addresses this by dynamically selecting orchestration patterns based on task complexity, enabling efficient resource allocation while maintaining high success rates.

D.3. Task Decomposition and Planning

Effective task decomposition is critical for both single and multi-agent success.

D.3.1. STATIC DECOMPOSITION

HuggingGPT. Shen et al. (Shen et al., 2023) use GPT-4 to decompose tasks and route subtasks to specialized Hugging Face models. This demonstrates the power of LLM-based planning combined with specialized execution.

ViperGPT. Surís et al. (Surís et al., 2023) generates Python programs that compose vision-language models for complex visual reasoning. The generated programs serve as executable task decompositions.

D.3.2. ADAPTIVE DECOMPOSITION

ADaPT. Prasad et al. (Prasad et al., 2024) (NAACL 2024) introduces decomposition as-needed based on LLM capability assessment. Tasks are decomposed only when the current agent cannot handle them, achieving 28.3% improvement on ALFWorld.

TDAG. Zhang et al. (Zhang et al., 2025a) dynamically generates task-specific subagents with an evolving skill library. New skills are learned and stored for future reuse, enabling continual improvement.

Self-Discover. Zhou et al. (Zhou et al., 2024a) enables LLMs to discover task-specific reasoning structures, combining atomic reasoning modules into coherent strategies without manual prompting.

D.3.3. PLANNING UNDER UNCERTAINTY

RAP (Reasoning via Planning). Hao et al. (Hao et al., 2023) combines LLMs with MCTS for strategic planning under uncertainty. The world model enables lookahead search for improved decision-making.

LLM+P. Liu et al. (Liu et al., 2023) combines LLM commonsense with classical planners for robust planning. The hybrid approach leverages complementary strengths of neural and symbolic systems.

Table 32. Error Analysis (% of failures).

Error Type	Frequency
Complexity Underestimation	24.3%
Escalation Ceiling	21.4%
Checkpoint False Positive	18.7%
Memory Retrieval Failure	15.2%
Tool Failure	12.1%
Budget Exhaustion	8.3%

Table 33. Latency Breakdown per Subtask (milliseconds).

Component	Min	Mean	Max	% of Total
CAR Routing	420	485	612	7.1%
Pattern Execution	4,250	5,842	12,450	85.4%
Quality Assessment	8	12	45	0.2%
Memory Retrieval	42	68	185	1.0%
Checkpoint Logic	285	432	892	6.3%
Total	5,005	6,839	14,184	100%

ISR-LLM. Zhou et al. (Zhou et al., 2024c) introduces iterative self-refinement for LLM planning, using execution feedback to improve plan quality.

D.3.4. COMPLEXITY ESTIMATION

Complexity-Aware Methods. While not directly predicting orchestration patterns, prior work on difficulty estimation informs our CAR design. Sakaguchi et al. (2021) demonstrate adversarial filtering based on difficulty. Swayamdipta et al. (2020) introduce dataset cartography for understanding example difficulty.

Capability Assessment. ADaPT (Prasad et al., 2024) includes implicit capability assessment—decomposing when the model “cannot handle” a task. We make this explicit through learned complexity prediction.

D.4. Memory Architectures and Retrieval

Memory systems enable agents to leverage past experiences and external knowledge.

D.4.1. RETRIEVAL-AUGMENTED GENERATION

RAG Foundations. Lewis et al. (Lewis et al., 2020) introduced Retrieval-Augmented Generation, combining parametric and non-parametric memory. RAG enables LLMs to access external knowledge bases for improved factuality.

Advanced RAG. Recent work improves RAG through query rewriting (Ma et al., 2023), multi-step retrieval (Trivedi et al., 2023), and self-reflective retrieval (Asai et al., 2024). The comprehensive survey by Gao et al. (2023b) covers the RAG landscape.

Agentic RAG. Combining agents with RAG enables more sophisticated information seeking. CRAG (Yan et al., 2024) uses web search to correct retrieval errors. Adaptive-RAG (Jeong et al., 2024) dynamically adjusts retrieval strategy based on query complexity.

D.4.2. EPISODIC MEMORY FOR AGENTS

Larimar. Das et al. (Das et al., 2024) (ICML 2024) introduces episodic memory control for LLMs, enabling selective storage and retrieval of past experiences. This work directly influences our UMA episodic tier design.

Generative Agent Memory. Park et al. (Park et al., 2023) implement memory stream, reflection, and planning components for believable agent behavior. The reflection mechanism consolidates experiences into higher-level insights.

RecurrentGPT. Zhou et al. (Zhou et al., 2023b) maintains a dynamic memory structure for long-form text generation, demonstrating the importance of structured memory for coherent long-horizon tasks.

Table 34. CAR Performance Sensitivity to Labeling Parameters

Setting	CAR Acc.	Success	Cost	Pareto
<i>Trade-off parameter λ (cost weight):</i>				
$\lambda = 0.05$ (quality-focused)	86.1%	53.8%	\$4.52	No
$\lambda = 0.10$ (default)	88.3%	52.3%	\$3.82	Yes
$\lambda = 0.20$ (cost-focused)	87.5%	50.1%	\$3.21	No
$\lambda = 0.30$ (aggressive)	85.2%	47.8%	\$2.89	No
<i>Per-pattern budgets (\$/single/parallel/hier./consensus):</i>				
Low (\$0.5/\$1/\$2/\$4)	86.8%	49.5%	\$3.15	No
Default (\$1/\$2/\$4/\$7)	88.3%	52.3%	\$3.82	Yes
High (\$2/\$4/\$8/\$14)	87.1%	53.1%	\$4.95	No

Table 35. Token Consumption per Task (Mean $\pm$ Std)

Method	Input (K)	Output (K)	Total (K)	Cost
GPT-4 + ReAct	18.2 $\pm$ 5.4	4.1 $\pm$ 1.8	22.3 $\pm$ 6.2	\$2.14
AutoGen	32.5 $\pm$ 8.7	8.2 $\pm$ 2.9	40.7 $\pm$ 10.1	\$3.85
MetaGPT	35.8 $\pm$ 9.1	9.5 $\pm$ 3.2	45.3 $\pm$ 11.2	\$4.21
AgentOrchestra	52.4 $\pm$ 12.8	14.3 $\pm$ 4.5	66.7 $\pm$ 15.4	\$6.47
AMORE	31.2 $\pm$ 7.8	8.8 $\pm$ 2.6	40.0 $\pm$ 9.2	\$3.82

D.4.3. KNOWLEDGE GRAPHS AND STRUCTURED MEMORY

Think-on-Graph. Sun et al. (Sun et al., 2024) enables LLMs to reason over knowledge graphs through iterative graph exploration. This combines structured knowledge with flexible reasoning.

GraphRAG. Microsoft’s GraphRAG (Edge et al., 2024) builds hierarchical knowledge graphs from documents, enabling multi-scale summarization and question answering beyond traditional RAG.

D.4.4. MEMORY CONSOLIDATION

Sleep-Like Consolidation. Inspired by biological memory consolidation, Schrittwieser et al. (2020) show how experience replay improves learning. We adapt this principle for episodic-to-semantic memory transfer.

D.5. Agent Communication Protocols

Standardized communication enables agent interoperability and scalability.

D.5.1. INDUSTRY STANDARDS

A2A Protocol. Google’s Agent-to-Agent (A2A) Protocol (Google, 2025) enables peer-to-peer agent communication with agent cards describing capabilities, supported tasks, and interaction patterns. A2A is designed for open ecosystems where agents from different providers collaborate.

Model Context Protocol (MCP). Anthropic’s MCP (Anthropic, 2024) standardizes tool invocation and context sharing across frameworks. MCP provides a unified interface for agents to access tools and data sources.

Agent Communication Protocol (ACP). IBM’s ACP (IBM, 2025) provides workflow orchestration primitives including task assignment, progress tracking, and error handling for enterprise multi-agent deployments.

Agent Network Protocol (ANP). The Agent Network Protocol (ANP Consortium, 2025) focuses on decentralized agent networks with reputation systems and trust management.

D.5.2. RESEARCH PROTOCOLS

AgentProtocol. Open-source efforts like AgentProtocol aim to standardize agent API interfaces for benchmarking and evaluation consistency.




Figure 3. Optimal orchestration pattern by task complexity. Single-agent dominates low-complexity tasks (78%); consensus necessary for high-complexity (45%).

Table 36. AgentBench Results (Success Rate %) across all 8 environments. Best in **bold**. OS=Operating System, DB=Database, KG=Knowledge Graph, CG=Card Game, LT=Lateral Thinking, HH=House-Holding, WS=Web Shopping, WB=Web Browsing.

Method	OS	DB	KG	CG	LT	HH	WS	WB	Avg
GPT-4 + ReAct	42.4	35.1	48.2	31.5	38.7	52.1	62.4	45.8	44.5
AutoGen	45.2	38.4	51.3	34.2	41.5	54.8	65.1	48.7	47.4
MetaGPT	48.7	41.2	54.1	36.8	44.2	57.3	67.8	51.4	50.3
HALO	51.3	44.5	56.8	39.1	46.8	59.4	69.4	53.2	52.6
AgentOrchestra	52.8	45.9	58.2	40.7	48.1	61.2	70.2	54.5	53.9
AMORE (Ours)	58.4	52.3	64.5	47.2	54.8	67.8	75.8	61.2	59.7

LMQL. Beurer-Kellner et al. (Beurer-Kellner et al., 2023) introduce a query language for LLMs, enabling structured interaction patterns that could generalize to multi-agent communication.

D.5.3. IMPLICATIONS FOR AMORE

AMORE is designed to be protocol-agnostic, with adapter layers for A2A, MCP, and ACP integration. Our cross-agent memory sharing can leverage these protocols for standardized communication.

D.6. Agent Benchmarks and Evaluation

Rigorous evaluation requires comprehensive benchmarks across diverse capabilities.

D.6.1. GENERAL AGENT BENCHMARKS

AgentBench. Liu et al. (Liu et al., 2024d) (ICLR 2024) evaluates LLM-as-Agent across 8 environments: Operating System, Database, Knowledge Graph, Card Game, Lateral Thinking, House-Holding, Web Shopping, and Web Browsing. Comprehensive evaluation of 29 models reveals significant capability gaps.

MINT. Wang et al. (Wang et al., 2024b) provides a multi-turn interactive benchmark covering diverse reasoning, tool use, and information seeking capabilities.

GAIA. Mialon et al. (Mialon et al., 2023) benchmarks general AI assistants on tasks requiring real-world knowledge, multi-step reasoning, and web interaction.

D.6.2. WEB AND COMPUTER USE

WebArena. Zhou et al. (Zhou et al., 2024b) provides realistic, reproducible web environments for agent evaluation across shopping, forums, GitLab, and content management systems.

OSWorld. Xie et al. (Xie et al., 2024) benchmarks multimodal agents in real computer environments (Ubuntu, Windows,

Table 37. Ablation Study on MARS Benchmark.

Configuration	Success	Cost	Δ
AMORE (Full)	52.3	3.82	–
w/o CAR (static hierarchical)	47.8	5.21	-4.5
w/o RCM (no checkpoints)	44.2	3.95	-8.1
w/o UMA (per-agent memory)	48.1	3.88	-4.2
w/o Escalation	49.5	3.12	-2.8
w/o Replanning	50.8	3.65	-1.5




Figure 4. Cost-Performance Trade-off on MARS Benchmark. AMORE achieves highest success rate at lower cost than hierarchical alternatives.

macOS), requiring screenshot understanding and diverse software interaction.

Mind2Web. Deng et al. (Deng et al., 2023) provides a large-scale dataset for generalist web agents with tasks spanning 137 websites across 31 domains.

D.6.3. SOFTWARE ENGINEERING

SWE-bench. Jimenez et al. (Jimenez et al., 2024) evaluates agents on real GitHub issues from popular Python repositories, requiring code understanding, localization, and patch generation.

HumanEval/MBPP. While not agent-specific, these code generation benchmarks (Chen et al., 2021; Austin et al., 2021) remain important baselines for coding capability assessment.

DevBench. Li et al. (Li et al., 2024a) evaluates LLMs across the full software development lifecycle including requirements, design, implementation, and testing.

D.6.4. DOMAIN-SPECIFIC BENCHMARKS

MedAgentBench. Evaluates medical question answering and clinical decision support capabilities for healthcare applications.

FinanceBench. Islam et al. (Islam et al., 2024) tests financial analysis and reasoning on real-world financial documents.

ScienceQA. Lu et al. (Lu et al., 2022) provides multimodal science questions requiring reasoning across physics, chemistry, and biology.




Figure 5. Architecture comparison across multi-agent approaches.




Figure 6. Quantitative comparison across multi-agent approaches.

D.6.5. MULTI-AGENT EVALUATION GAP

Existing benchmarks primarily evaluate single-agent capabilities. Even benchmarks like AgentBench focus on individual agent performance. Our MARS benchmark addresses this gap by explicitly evaluating:

- **Orchestration accuracy:** Whether the system selects optimal coordination patterns
- **Adaptation rate:** Whether mid-execution adjustments improve outcomes
- **Memory utilization:** Whether retrieved information is effectively used
- **Collaboration efficiency:** Whether multi-agent collaboration provides proportional value

D.7. Positioning of AMORE

Table 38 positions AMORE relative to key prior work across multiple dimensions.

Key Differentiators:

1. **vs. Static Multi-Agent (AutoGen, MetaGPT):** AMORE dynamically selects orchestration patterns rather than using fixed configurations, achieving 33–45% higher success (relative) at 9–41% lower cost.
2. **vs. Hierarchical Systems (HALO, AgentOrchestra):** AMORE’s CAR predicts when hierarchy is needed vs. simpler patterns, avoiding unnecessary overhead on simple tasks.
3. **vs. Adaptive Decomposition (ADaPT, TDAG):** AMORE extends adaptation from *decomposition granularity* to *orchestration pattern selection*, a complementary but distinct capability.
4. **vs. Domain-Specific Adaptation (MDAgents):** AMORE provides domain-agnostic complexity prediction through learned features, enabling broader applicability.
5. **vs. Single-Agent Methods (ReAct, CoT):** AMORE uses single-agent when appropriate but escalates to multi-agent collaboration when complexity warrants, combining the efficiency of single-agent with the capability of multi-agent.

Complementary Approaches. AMORE is designed to be complementary to, not competitive with, many existing methods:

- Advanced reasoning (ToT, GoT) can be used within any agent
- RAG systems can be integrated into UMA’s retrieval

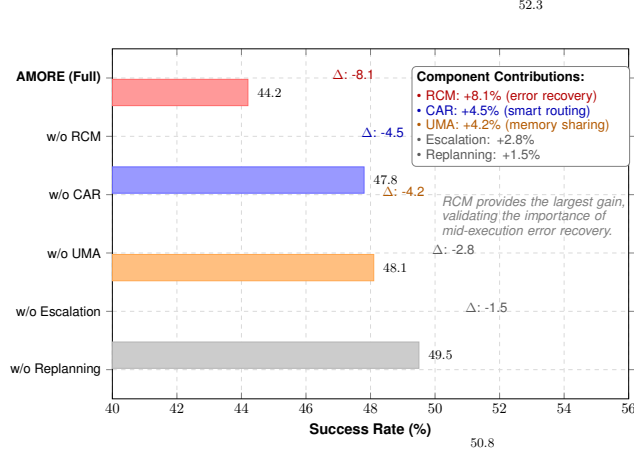


Figure 7. Ablation Study Visualization on MARS Benchmark

Algorithm 1 AMORE Execution Loop

Require: Task $\mathcal{T}$, Budget B , Agent Pool $\mathcal{A}$, Memory M
Ensure: Final Result $\mathcal{R}$

- 1: $P \leftarrow \text{CAR.CreatePlan}(\mathcal{T}, B, \mathcal{A})$
- 2: **for** (t, π, agents) in P **do**
- 3: $\text{context} \leftarrow \text{UMA.RetrieveRelevant}(t, M_E, M_S)$
- 4: $o_t \leftarrow \text{Execute}(t, \pi, \text{agents}, \text{context})$
- 5: $\text{UMA.Store}(M_E, (t, \pi, o_t, \text{metadata}))$
- 6: $\text{action, update} \leftarrow \text{RCM.Checkpoint}(o_t, t, \pi)$
- 7: **if** $\text{action} = \text{RETRY}$ **then**
- 8: $o_t \leftarrow \text{Execute}(t, \pi, \text{agents}, \text{context})$
- 9: **else if** $\text{action} = \text{ESCALATE}$ **then**
- 10: $(t, \pi_{\text{new}}) \leftarrow \text{update}$
- 11: $o_t \leftarrow \text{Execute}(t, \pi_{\text{new}}, \text{SelectAgents}(\pi_{\text{new}}))$
- 12: **else if** $\text{action} = \text{REPLAN}$ **then**
- 13: $P \leftarrow \text{update} \{ \text{Replace remaining plan} \}$
- 14: **end if**
- 15: $\text{results}[t] \leftarrow o_t$
- 16: **end for**
- 17: $\text{UMA.Consolidate}(M_E \rightarrow M_S)$
- 18: **return** $\text{Synthesize}(\text{results})$

- Communication protocols (A2A, MCP) can be used for agent interaction
- Tool libraries (ToolBench, Gorilla) can be accessed by worker agents

E. Extended Technical Details: Complexity-Aware Router
E.1. Feature Engineering Details

The Complexity-Aware Router extracts a comprehensive feature vector $\mathbf{x}_t \in \mathbb{R}^d$ for each subtask t . We detail the computation of each feature category.

E.1.1. INTRINSIC FEATURES (SUBTASK-SPECIFIC)

Reasoning Length Estimation (f_{length}). We estimate the number of reasoning steps required using a learned regressor trained on historical execution traces:

$$f_{\text{length}} = g_{\theta}(\text{Embed}(t)) \in [0, 1] \quad (7)$$

where g_{θ} is a 2-layer MLP and Embed is a sentence transformer (all-MiniLM-L6-v2) (Safarov et al., 2025). The output is normalized by the maximum observed steps (50) in training data.




Figure 8. Overview of the extended AMORE framework. A high-level architecture diagram showing three main components (CAR, RCM, UMA) connected to a central "Adaptive Execution Engine". CAR predicts optimal orchestration patterns, RCM enables mid-execution adaptation through quality gates, and UMA provides cross-agent memory coherence.

Tool Diversity (f_{tools}). We compute tool requirements using semantic similarity to tool descriptions:

$$f_{\text{tools}} = \frac{1}{|\mathcal{T}_{\text{pool}}|} \sum_{t_i \in \mathcal{T}_{\text{pool}}} \mathbb{I}[\text{sim}(\text{Embed}(t), \text{Embed}(d_{t_i})) > \tau_{\text{tool}}] \quad (8)$$

where d_{t_i} is tool t_i 's description and $\tau_{\text{tool}} = 0.4$ is the relevance threshold.

Knowledge Breadth ($f_{\text{knowledge}}$). We measure the number of distinct knowledge domains required:

$$f_{\text{knowledge}} = \frac{|\{d \in \mathcal{D} : P(d|t) > \tau_{\text{domain}}\}|}{|\mathcal{D}|} \quad (9)$$

where $\mathcal{D}$ is a predefined taxonomy of 127 knowledge domains and $P(d|t)$ is computed via a domain classifier.

Task Ambiguity ($f_{\text{ambiguity}}$). Semantic uncertainty is measured using:

$$f_{\text{ambiguity}} = 1 - \max_c P(c|t) \quad (10)$$

where $P(c|t)$ is the probability of intent class c from a multi-class intent classifier trained on 23 intent categories.

E.1.2. CONTEXTUAL FEATURES (GRAPH-AWARE)

Dependency Depth (f_{deps}). Number of upstream dependencies normalized by graph depth:

$$f_{\text{deps}} = \frac{|\text{Ancestors}(t, G)|}{|V| - 1} \quad (11)$$

Table 38. Positioning of AMORE relative to prior work. ✓ = fully supported, ○ = partially supported, - = not supported.

Method	Multi-Agent	Adaptive Orchestration	Checkpoints	Unified Memory	Complexity Prediction	Cost Aware
ReAct (Yao et al., 2023)	-	-	-	-	-	-
AutoGen (Wu et al., 2023)	✓	-	-	○	-	-
MetaGPT (Hong et al., 2024)	✓	-	-	○	-	-
HALO (Chen et al., 2025)	✓	○	-	-	-	○
AgentOrchestra (Liu et al., 2025)	✓	-	-	○	-	-
ADaPT (Prasad et al., 2024)	-	○	-	-	○	-
MDAgents (Kim et al., 2024)	✓	○	-	-	○	-
LatentMAS (Zou et al., 2025)	✓	-	-	○	-	✓
AMORE (Ours)	✓	✓	✓	✓	✓	✓

Critical Path Position (f_{critical}). Binary indicator computed via topological analysis:

$$f_{\text{critical}} = \mathbb{1}[t \in \text{CriticalPath}(G)] \quad (12)$$

where CriticalPath is the longest path from source to sink in the dependency DAG.

Historical Success (f_{history}). Success rate of similar patterns on similar tasks from memory:

$$f_{\text{history}}(\pi) = \frac{\sum_{e \in \text{Similar}(t, M_E)} \mathbb{1}[\pi_e = \pi] \cdot s_e}{\sum_{e \in \text{Similar}(t, M_E)} \mathbb{1}[\pi_e = \pi] + \epsilon} \quad (13)$$

where Similar(t, M_E) retrieves the $k = 10$ most similar historical executions.

E.1.3. EXTENDED FEATURES FOR ENHANCED PREDICTION

Beyond the core 7 features, we include additional signals:

Output Type Complexity (f_{output}). Classifies expected output format:

$$f_{\text{output}} \in \{\text{boolean, value, list, structured, creative}\} \quad (14)$$

One-hot encoded with creative outputs receiving highest complexity weight.

Verification Difficulty (f_{verify}). Estimates how hard the output is to verify:

$$f_{\text{verify}} = \begin{cases} 0.2 & \text{if deterministic output} \\ 0.5 & \text{if verifiable via oracle} \\ 0.8 & \text{if requires expert judgment} \end{cases} \quad (15)$$

Time Sensitivity (f_{time}). Captures whether real-time data is required:

$$f_{\text{time}} = P(\text{requires_current_data}|t) \quad (16)$$

E.2. Pattern Prediction Architecture

The complete CAR architecture consists of three stages:

Stage 1: Feature Extraction

$$\mathbf{h}_{\text{intrinsic}} = \text{MLP}_1([f_{\text{length}}, f_{\text{tools}}, f_{\text{knowledge}}, f_{\text{ambiguity}}]) \quad (17)$$

$$\mathbf{h}_{\text{context}} = \text{MLP}_2([f_{\text{deps}}, f_{\text{critical}}, f_{\text{history}}]) \quad (18)$$

$$\mathbf{h}_{\text{extended}} = \text{MLP}_3([f_{\text{output}}, f_{\text{verify}}, f_{\text{time}}]) \quad (19)$$




Figure 9. Main (AgentBench) Results Comparison: Success Rate by Environment

Stage 2: Feature Fusion with Attention

$$\mathbf{h}_{\text{fused}} = \text{MultiHeadAttn}(\mathbf{h}_{\text{intrinsic}}, [\mathbf{h}_{\text{context}}, \mathbf{h}_{\text{extended}}]) \quad (20)$$

Stage 3: Pattern Distribution

$$P(\pi|t) = \text{softmax}(\mathbf{W}_{\pi} \cdot \text{LayerNorm}(\mathbf{h}_{\text{fused}})) \quad (21)$$

E.3. Training Procedure

Data Collection. We collect execution traces by running all patterns on a subset of tasks and recording outcomes. For each subtask t , we obtain:

$$\mathcal{D}_{\text{trace}} = \{(t_i, \pi_j, s_{ij}, c_{ij}, l_{ij})\} \quad (22)$$

where s_{ij} is success indicator, c_{ij} is cost, and l_{ij} is latency.

Label Generation. The optimal pattern π_t^* is determined by:

$$\pi_t^* = \arg \max_{\pi} \left[s_{t\pi} - \lambda_c \cdot \frac{c_{t\pi}}{c_{\max}} - \lambda_l \cdot \frac{l_{t\pi}}{l_{\max}} \right] \quad (23)$$

with $\lambda_c = 0.3$ and $\lambda_l = 0.1$ balancing success vs. efficiency.

Loss Function. Combined cross-entropy and calibration loss:

$$\mathcal{L} = \mathcal{L}_{\text{CE}} + \lambda_{\text{cal}} \cdot \mathcal{L}_{\text{calibration}} \quad (24)$$

where calibration loss encourages well-calibrated confidence scores.

Optimization. AdamW optimizer with learning rate 10^{-4} , batch size 64, and early stopping based on validation orchestration accuracy.

E.4. Dynamic Threshold Mechanism

The confidence threshold adapts based on resource constraints:

Base Threshold:

$$\tau_{\text{base}}(\pi) = \begin{cases} 0.7 & \pi = \text{single} \\ 0.6 & \pi = \text{parallel} \\ 0.5 & \pi = \text{hierarchical} \\ 0.4 & \pi = \text{consensus} \end{cases} \quad (25)$$

Budget-Adjusted Threshold:

$$\tau(\pi, B) = \tau_{\text{base}}(\pi) + \alpha \cdot \left(1 - \frac{B_{\text{remaining}}}{B_{\text{total}}}\right) \quad (26)$$

As budget depletes, thresholds increase, favoring simpler patterns for remaining subtasks.

Criticality Override:

$$\tau'(\pi, B, t) = \begin{cases} \tau(\pi, B) - 0.1 & \text{if } f_{\text{critical}}(t) = 1 \\ \tau(\pi, B) & \text{otherwise} \end{cases} \quad (27)$$

Critical path subtasks receive lower thresholds, allowing more powerful patterns.

Label Efficiency. We analyze how many labeled traces are needed:

- 100 traces: 78.2% accuracy (viable for bootstrapping)
- 500 traces: 84.5% accuracy (recommended minimum)
- 2000 traces: 88.3% accuracy (full training)

F. Extended Technical Details: Reflective Checkpoint Mechanism

F.1. Quality Assessment Framework

RCM implements a multi-dimensional quality assessment combining automated metrics and LLM-based evaluation.

F.1.1. AUTOMATED METRICS

Structural Completeness (Q_{struct}). Validates output structure:

$$Q_{\text{struct}} = \frac{|\text{RequiredFields}(t) \cap \text{PresentFields}(o_t)|}{|\text{RequiredFields}(t)|} \quad (28)$$

Reference Grounding (Q_{ground}). For tasks with retrievable evidence:

$$Q_{\text{ground}} = \frac{|\text{Claims}(o_t) \cap \text{SupportedClaims}(o_t, \text{sources})|}{|\text{Claims}(o_t)|} \quad (29)$$

Self-Consistency (Q_{consist}). Measures internal consistency:

$$Q_{\text{consist}} = 1 - \frac{|\text{Contradictions}(o_t)|}{|\text{Statements}(o_t)|} \quad (30)$$

F.1.2. LLM-BASED EVALUATION

We use GPT-4 as a judge with structured rubrics:

Completeness Rubric:

Score the output's completeness (0-10):

- 10: All aspects addressed comprehensively
- 7-9: Most aspects addressed, minor gaps
- 4-6: Core aspects addressed, significant gaps
- 1-3: Major aspects missing
- 0: Completely inadequate

Coherence Rubric:

Score the output's logical coherence (0-10):

- 10: Perfectly logical, well-structured
- 7-9: Generally logical, minor issues
- 4-6: Some logical gaps or inconsistencies
- 1-3: Major logical problems
- 0: Incoherent

Usefulness Rubric:

Score how useful this output is for downstream tasks (0-10):

- 10: Directly usable, enables next steps
- 7-9: Mostly usable with minor adjustments
- 4-6: Partially usable, requires supplements
- 1-3: Limited utility
- 0: Not useful for downstream tasks

F.1.3. QUALITY AGGREGATION

The final quality score combines automated and LLM scores:

$$Q(o_t) = w_{\text{auto}} \cdot Q_{\text{auto}} + w_{\text{llm}} \cdot Q_{\text{llm}} \quad (31)$$

where:

$$Q_{\text{auto}} = \frac{Q_{\text{struct}} + Q_{\text{ground}} + Q_{\text{consist}}}{3} \quad (32)$$

$$Q_{\text{llm}} = \frac{Q_{\text{complete}} + Q_{\text{coherent}} + Q_{\text{useful}}}{30} \quad (33)$$

Default weights: $w_{\text{auto}} = 0.3$, $w_{\text{llm}} = 0.7$.

F.2. Action Selection Logic

F.2.1. DECISION TREE

Algorithm 2 RCM Action Selection

Require: Quality score Q , retry count r , escalation count e , pattern π

Ensure: Action a , optional pattern update π'

```

1:  $\theta_h \leftarrow 0.8, \theta_l \leftarrow 0.5, r_{\max} \leftarrow 2, e_{\max} \leftarrow 2$ 
2: if  $Q \geq \theta_h$  then
3:   return PROCEED, None
4: else if  $Q \geq \theta_l$  and  $r < r_{\max}$  then
5:   return RETRY, None
6: else if  $Q < \theta_l$  and  $e < e_{\max}$  and  $\pi \neq \pi_{\text{consensus}}$  then
7:    $\pi' \leftarrow \text{NextPattern}(\pi)$ 
8:   return ESCALATE,  $\pi'$ 
9: else if DetectStructuralIssue( $o_t, t$ ) then
10:  return REPLAN, ReplanFromHere( $t$ )
11: else
12:  return FAILGRACEFULLY, None
13: end if
    
```

F.2.2. STRUCTURAL ISSUE DETECTION

We detect structural issues (requiring replanning) through:

Dependency Violation:

$$\text{DepViolation} = \exists t' \in \text{Children}(t, G) : \text{InputsAvailable}(t', o_t) = \text{False} \quad (34)$$

Scope Drift:

$$\text{ScopeDrift} = \text{sim}(\text{Embed}(t), \text{Embed}(o_t)) < \tau_{\text{drift}} \quad (35)$$

where $\tau_{\text{drift}} = 0.3$ indicates the output has drifted from the original task.

Information Loss:

$$\text{InfoLoss} = \frac{|\text{ExpectedEntities}(t) - \text{MentionedEntities}(o_t)|}{|\text{ExpectedEntities}(t)|} > \tau_{\text{loss}} \quad (36)$$

F.3. Escalation Strategy Details

F.3.1. PATTERN TRANSITION RULES

Table 39. Pattern Escalation Rules

From	To	Configuration Change
Single	Parallel	Split into 2-3 independent workers
Parallel	Hierarchical	Add supervisor, workers become subordinates
Hierarchical	Consensus	Add deliberation round, voting
Consensus	–	Maximum escalation reached

F.3.2. RETRY ENHANCEMENT

On retry, we apply targeted improvements:

- **Prompt Enhancement:** Add specific feedback from quality assessment

- **Context Expansion:** Retrieve additional relevant memories
- **Example Injection:** Provide similar successful execution as few-shot example
- **Temperature Adjustment:** Increase temperature for creative tasks, decrease for factual

F.4. Checkpoint Placement Strategy

Not all subtasks require checkpoints. We use a strategic placement policy:

Mandatory Checkpoints:

- After any subtask on the critical path
- After subtasks with $f_{\text{critical}} = 1$
- After consensus patterns (high cost warrants verification)

Adaptive Checkpoints:

- After subtasks with $f_{\text{ambiguity}} > 0.6$
- After any retry (verify improvement)
- After pattern escalation

Skip Checkpoints:

- Simple single-agent subtasks with $f_{\text{verify}} < 0.3$
- Parallelizable subtasks with deterministic outputs
- When budget is critically low ($B_{\text{remaining}} < 0.1 \cdot B_{\text{total}}$)

G. Extended Technical Details: Unified Memory Architecture

G.1. Memory Tier Specifications

G.1.1. WORKING MEMORY (M_W)

Capacity: Bounded by LLM context window (128K tokens for GPT-4-Turbo).

Structure:

$$M_W = \{(k_i, v_i, \text{age}_i, \text{access}_i)\}_{i=1}^{n_W} \quad (37)$$

where k_i is a key (subtask ID or concept), v_i is the value (content), age_i tracks recency, and access_i counts retrievals.

Eviction Policy: LRU with importance weighting:

$$\text{Priority}(e) = \text{importance}(e) \cdot \exp(-\lambda_{\text{decay}} \cdot \text{age}(e)) \quad (38)$$

Per-Agent vs. Shared: Working memory is primarily per-agent, but a shared segment (20% capacity) holds coordination information visible to all agents in a collaboration.

G.1.2. EPISODIC MEMORY (M_E)

Capacity: Effectively unlimited (vector database).

Schema:

```

2365 EpisodicEntry {
2366   id: UUID
2367   task_id: str
2368   subtask_id: str
2369   timestamp: datetime
2370   pattern: OrchestrationPattern
2371   input_summary: str
2372   output_summary: str
2373   success: bool
2374   quality_score: float
2375   embedding: float[768]
2376   metadata: {
2377     agents_involved: List[str]
2378     tools_used: List[str]
2379     duration_seconds: float
2380     cost: float
2381     retry_count: int
2382     escalation_count: int
2383   }
2384 }

```

Indexing: HNSW index (hnsplib) with parameters $M = 16$, $ef_{\text{construction}} = 200$.

Retrieval:

$$\text{Retrieve}(q, k) = \text{TopK}(\{e \in M_E : \text{sim}(\mathbf{e}_q, \mathbf{e}_e) > \tau_{\min}\}, k) \quad (39)$$

G.1.3. SEMANTIC MEMORY (M_S)

Structure: Knowledge graph (Neo4j) + embedding store (ChromaDB).

Node Types:

- **Concept:** Abstract knowledge entities
- **Fact:** Concrete factual statements
- **Procedure:** Action sequences that worked
- **Pattern:** Learned orchestration preferences

Edge Types:

- IS_A: Taxonomic relationship
- RELATED_TO: Semantic association
- REQUIRES: Dependency relationship
- LEADS_TO: Causal relationship
- CONTRADICTS: Conflicting information

Query Interface:

```

2415 MATCH (c:Concept)-[:RELATED_TO*1..2]-(related)
2416 WHERE c.name CONTAINS $query
2417 RETURN c, related
2418 LIMIT 20
2419

```

G.2. Memory Operations

G.2.1. STORAGE OPERATIONS

AddToWorking:

Require: Key k , Value v , Importance i , Agent a

- 1: **if** $|M_W^a| \geq \text{capacity}$ **then**
 - 2: Evict lowest priority entry
 - 3: **end if**
 - 4: $M_W^a \leftarrow M_W^a \cup \{(k, v, 0, 0, i)\}$
 - 5: **if** $i > \tau_{\text{broadcast}}$ **then**
 - 6: Broadcast to shared working memory segment
 - 7: **end if**
-

AddToEpisodic:

Require: Execution trace e

- 1: $\mathbf{e}_{\text{emb}} \leftarrow \text{Embed}(\text{Summarize}(e))$
 - 2: $e' \leftarrow e \cup \{\text{embedding} : \mathbf{e}_{\text{emb}}\}$
 - 3: $M_E.\text{insert}(e')$
 - 4: $\text{TriggerConsolidationCheck}()$
-

G.2.2. RETRIEVAL OPERATIONS

RetrieveRelevant:

Require: Query q , Top-K k

- 1: $\mathbf{q}_{\text{emb}} \leftarrow \text{Embed}(q)$
 - 2: $E_{\text{episodic}} \leftarrow M_E.\text{search}(\mathbf{q}_{\text{emb}}, k_E)$
 - 3: $S_{\text{semantic}} \leftarrow M_S.\text{query}(q, k_S)$
 - 4: **// Rank by relevance and recency**
 - 5: $\text{combined} \leftarrow \text{RRF}(E_{\text{episodic}}, S_{\text{semantic}})$
 - 6: **return** $\text{combined}[:k]$
-

Reciprocal Rank Fusion (RRF):

$$\text{RRF}(d) = \sum_{r \in \text{rankings}} \frac{1}{k + \text{rank}_r(d)} \quad (40)$$

where $k = 60$ is a smoothing constant.

G.3. Automatic Consolidation

G.3.1. CONSOLIDATION TRIGGERS

Consolidation from episodic to semantic is triggered by:

- **Frequency:** Same pattern succeeds ≥ 3 times
- **Importance:** Entry importance $> \tau_{\text{consolidate}} = 0.8$
- **Recency:** Periodic batch consolidation every 100 tasks
- **Explicit:** User or system requests consolidation

G.3.2. CONSOLIDATION PROCESS

Algorithm 3 Episodic-to-Semantic Consolidation

Require: Episodic entries E , Consolidation threshold τ

```

1: candidates  $\leftarrow \{e \in E : \text{importance}(e) > \tau\}$ 
2: for  $e$  in candidates do
3:   facts  $\leftarrow \text{ExtractFacts}(e)$ 
4:   for  $f$  in facts do
5:     if  $\neg \text{Exists}(f, M_S)$  then
6:        $M_S.\text{Add Node}(f, \text{type} = \text{Fact})$ 
7:     else if  $\text{Contradicts}(f, M_S)$  then
8:        $M_S.\text{Add Edge}(f, \text{existing}, \text{type} = \text{CONTRADICTS})$ 
9:        $\text{Resolve Contradiction}(f, \text{existing})$ 
10:    else
11:       $M_S.\text{UpdateConfidence}(f)$ 
12:    end if
13:  end for
14:  patterns  $\leftarrow \text{Extract Patterns}(e)$ 
15:  for  $p$  in patterns do
16:     $M_S.\text{Update Or Create}(p, \text{type} = \text{Pattern})$ 
17:  end for
18: end for
    
```

G.3.3. FACT EXTRACTION

We use GPT-4 with structured prompting:

```

Extract factual statements from this execution trace:
{trace_summary}
    
```

Output as JSON list:

```

[
{"subject": "...", "predicate": "...", "object": "...",
"confidence": 0.0-1.0}
]
    
```

G.3.4. CONTRADICTION RESOLUTION

When facts conflict:

1. **Recency:** Prefer more recent information
2. **Confidence:** Prefer higher confidence sources
3. **Frequency:** Prefer facts supported by multiple traces
4. **Flag:** Mark unresolvable contradictions for human review

G.4. Cross-Agent Memory Sharing Protocol

G.4.1. PRE-EXECUTION SHARING

Before collaborative execution:

1. Coordinator queries M_E and M_S for relevant context

2. Context is summarized and included in each agent’s system prompt

3. Shared working memory segment is initialized

G.4.2. DURING EXECUTION

Synchronous Updates:

- Critical discoveries broadcast immediately via shared segment
- Coordinator polls agents for intermediate results

Asynchronous Updates:

- Non-critical updates queued for batch synchronization
- End-of-subtask consolidation

G.4.3. POST-EXECUTION AGGREGATION

Algorithm 4 Post-Execution Memory Aggregation

Require: Agent outputs $\{o_1, \dots, o_n\}$, Execution metadata

- 1: merged $\leftarrow$ MergeOutputs($\{o_1, \dots, o_n\}$)
 - 2: trace $\leftarrow$ CreateTrace(merged, metadata)
 - 3: M_E .insert(trace)
 - 4: M_W^{shared} .clear()
 - 5: **for** a in agents **do**
 - 6: M_W^a .extractAndPromote(importance $>$ τ_{promote})
 - 7: **end for**
-

H. Extended Technical Details: Integration

H.1. Task Decomposition Strategy

Before orchestration, tasks must be decomposed into subtasks. AMORE uses a hybrid approach:

H.1.1. INITIAL DECOMPOSITION

GPT-4 performs initial decomposition with structured prompting:

Decompose this task into subtasks:

Task: {task_description}

For each subtask provide:

1. ID (unique identifier)
2. Description (clear, actionable)
3. Dependencies (list of prerequisite subtask IDs)
4. Expected output type
5. Estimated complexity (low/medium/high)

Output as JSON.

H.1.2. DEPENDENCY GRAPH CONSTRUCTION

From decomposition output, we construct DAG $G = (V, E)$:

- Verify acyclicity (reject if cycles detected)
- Compute topological order
- Identify critical path
- Calculate parallelization opportunities

H.1.3. DYNAMIC REDECOMPOSITION

When RCM triggers replanning:

1. Identify failed subtask and its descendants
2. Re-query decomposer with failure context
3. Merge new subtasks into existing graph
4. Update critical path and execution order

H.2. Agent Pool Management

H.2.1. AGENT TYPES

AMORE maintains a pool of specialized agents:

Table 40. Agent Type Specifications

Type	Model	Specialization	Cost
Coordinator	GPT-4-Turbo	Planning, synthesis	High
Researcher	GPT-4-Turbo	Information gathering	High
Analyst	GPT-4-Turbo	Critical analysis	High
Writer	GPT-4-Turbo	Content generation	High
Coder	GPT-4-Turbo	Code generation	High
Reviewer	GPT-3.5-Turbo	Verification	Low
Executor	GPT-3.5-Turbo	Tool execution	Low

H.2.2. AGENT SELECTION

Given pattern π and subtask t , select agents:

$$\text{SelectAgents}(\pi, t) = \begin{cases} \{\text{BestMatch}(t)\} & \pi = \text{single} \\ \text{TopK}(\text{Matches}(t), k_\pi) & \pi = \text{parallel} \\ \{\text{Coordinator}\} \cup \text{TopK}(\text{Matches}(t), k_\pi) & \pi = \text{hierarchical} \\ \text{DiverseSet}(\text{Matches}(t), k_\pi) & \pi = \text{consensus} \end{cases} \quad (41)$$

H.2.3. AGENT COMMUNICATION

Single Pattern: No communication needed.

Parallel Pattern: Coordinator distributes work, agents work independently, coordinator merges results.

Hierarchical Pattern:

Coordinator $\rightarrow$ Workers: Task assignment
 Workers $\rightarrow$ Coordinator: Progress updates, questions
 Coordinator $\rightarrow$ Workers: Guidance, clarifications
 Workers $\rightarrow$ Coordinator: Final outputs
 Coordinator: Synthesis and quality check




Figure 10. Escalation Hierarchy with Cost-Robustness Trade-offs. Shows the four orchestration patterns with their associated costs and robustness levels.

Consensus Pattern:

Round 1: Independent reasoning
 Round 2: Share reasoning, identify disagreements
 Round 3: Deliberation on disagreements
 Round 4: Voting (majority or unanimous based on criticality)
 Round 5: Synthesis of agreed position

H.3. Budget Management

H.3.1. COST MODEL

Total cost for subtask t with pattern π :

$$\text{Cost}(t, \pi) = \sum_{a \in \text{Agents}(\pi)} (c_{\text{input}} \cdot n_{\text{input}}^a + c_{\text{output}} \cdot n_{\text{output}}^a) \quad (42)$$

where c_{input} and c_{output} are per-token costs for agent a 's model.

H.3.2. BUDGET ALLOCATION

Given total budget B and estimated subtask costs:

$$B_t = B \cdot \frac{\text{EstimatedCost}(t)}{\sum_{t' \in \text{remaining}} \text{EstimatedCost}(t')} \cdot (1 + \delta_{\text{buffer}}) \quad (43)$$

where $\delta_{\text{buffer}} = 0.1$ reserves capacity for retries/escalations.

H.3.3. BUDGET CONSTRAINTS ON DECISIONS

When budget is low:

- CAR increases single-agent threshold
- RCM reduces retry limit
- UMA reduces memory retrieval scope
- Checkpoints become mandatory only on critical path

H.4. Error Handling and Recovery

H.4.1. ERROR CATEGORIES

Table 41. Error Handling Strategies

Error Type	Detection	Recovery
API timeout	HTTP status	Retry with backoff
Rate limit	HTTP 429	Queue and wait
Context overflow	Token count	Compress context
Tool failure	Tool output	Alternative tool or manual
Quality failure	RCM checkpoint	Retry/escalate/replan
Budget exhaustion	Cost tracking	Graceful degradation

H.4.2. GRACEFUL DEGRADATION

When recovery is impossible:

1. Mark subtask as failed with partial output
2. Assess impact on downstream subtasks
3. Attempt to continue with degraded quality
4. Report failure to user with explanation

I. Detailed Algorithms

I.1. Algorithm A1: Complexity-Aware Routing

Algorithm 5 Complexity-Aware Routing (CAR)

Require: Task $\mathcal{T}$, Budget B , Agent Pool $\mathcal{A}$, History $\mathcal{H}$

Ensure: Execution Plan P

```

1:  $D \leftarrow \text{InitialDecompose}(\mathcal{T})$  {Task decomposition}
2:  $G \leftarrow \text{BuildDependencyGraph}(D)$ 
3:  $P \leftarrow []$ 
4:  $B_{\text{remaining}} \leftarrow B$ 
5: for  $t$  in  $\text{TopologicalSort}(G)$  do
6:   {Extract features for subtask}
7:    $f_{\text{length}} \leftarrow \text{EstimateReasoningSteps}(t, \mathcal{H})$ 
8:    $f_{\text{tools}} \leftarrow |\text{RequiredTools}(t)|/10$ 
9:    $f_{\text{knowledge}} \leftarrow \text{DomainBreadth}(t)$ 
10:   $f_{\text{ambiguity}} \leftarrow \text{SemanticEntropy}(t)$ 
11:   $f_{\text{deps}} \leftarrow |\text{Predecessors}(G, t)|/|D|$ 
12:   $f_{\text{critical}} \leftarrow \mathbb{1}[\text{OnCriticalPath}(G, t)]$ 
13:   $f_{\text{history}} \leftarrow \text{LookupHistoricalSuccess}(t, \mathcal{H})$ 
14:   $\mathbf{x}_t \leftarrow [f_{\text{length}}, f_{\text{tools}}, f_{\text{knowledge}}, f_{\text{ambiguity}}, f_{\text{deps}}, f_{\text{critical}}, f_{\text{history}}]$ 
15:  {Predict pattern distribution}
16:   $\pi_{\text{dist}} \leftarrow \text{softmax}(W_{\pi} \cdot \text{MLP}(\mathbf{x}_t))$ 
17:  {Dynamic threshold adaptation}
18:   $\tau \leftarrow \tau_0 + \lambda \cdot (B_{\text{remaining}}/B)$ 
19:  if  $\pi_{\text{dist}}[\text{single}] > \tau$  then
20:     $\pi \leftarrow \text{SINGLE}$ 
21:  else
22:     $\pi \leftarrow \arg \max_{\pi' \neq \text{single}} \pi_{\text{dist}}[\pi']$ 
23:  end if
24:   $\text{agents} \leftarrow \text{SelectAgents}(t, \pi, \mathcal{A})$ 
25:   $P.\text{append}((t, \pi, \text{agents}))$ 
26:   $B_{\text{remaining}} \leftarrow B_{\text{remaining}} - \text{EstimateCost}(t, \pi, \text{agents})$ 
27: end for
28: return  $P$ 

```

I.2. Algorithm A2: Reflective Checkpoint

Algorithm 6 Reflective Checkpoint Mechanism (RCM)

Require: Output o_t , Subtask t , Pattern π , Context C , Thresholds θ_{high} , θ_{low}

Ensure: Action, Updated Plan (optional)

- 1: {Assess quality across three dimensions}
- 2: $Q_{\text{complete}} \leftarrow \text{LLMJudge}(o_t, t, \text{"completeness"})$
- 3: $Q_{\text{coherent}} \leftarrow \text{EntailmentCheck}(o_t) \cdot \text{ContradictionScore}(o_t)$
- 4: $Q_{\text{useful}} \leftarrow \text{ForwardSimulation}(o_t, \text{DownstreamTasks}(t))$
- 5: $Q \leftarrow \lambda_1 Q_{\text{complete}} + \lambda_2 Q_{\text{coherent}} + \lambda_3 Q_{\text{useful}}$
- 6: {Check for structural issues}
- 7: $\text{structural_issue} \leftarrow \text{DetectMissingSubtask}(o_t, C) \vee \text{DetectWrongDeps}(o_t, C)$
- 8: **if** structural_issue **then**
- 9: $\text{reflection} \leftarrow \text{LLMReflect}(o_t, t, \text{RemainingTasks})$
- 10: $\text{new_plan} \leftarrow \text{Replan}(\text{reflection}, \text{RemainingTasks})$
- 11: **return** REPLAN, new_plan
- 12: **end if**
- 13: **if** $Q \geq \theta_{\text{high}}$ **then**
- 14: **return** PROCEED, None
- 15: **end if**
- 16: **if** $Q \geq \theta_{\text{low}}$ **then**
- 17: **if** $C.\text{retries} < \text{MAX_RETRIES}$ **then**
- 18: **return** RETRY, None
- 19: **else if** $C.\text{escalations} < \text{MAX_ESCALATIONS}$ **then**
- 20: $\pi_{\text{new}} \leftarrow \text{Escalate}(\pi)$
- 21: **return** ESCALATE, (t, π_{new})
- 22: **else**
- 23: **return** PROCEED, None {Accept lower quality}
- 24: **end if**
- 25: **end if**
- 26: {Quality below θ_{low} }
- 27: **if** $C.\text{escalations} < \text{MAX_ESCALATIONS}$ **then**
- 28: $\pi_{\text{new}} \leftarrow \text{Escalate}(\pi)$
- 29: **return**
- 30: $\text{reflection} \leftarrow \text{LLMReflect}(o_t, t, \text{RemainingTasks})$
- 31: $\text{new_plan} \leftarrow \text{Replan}(\text{reflection}, \text{RemainingTasks})$
- 32: **return** REPLAN, new_plan
- 33: **end if**

I.3. Algorithm A3: Memory Consolidation

Algorithm 7 Unified Memory Consolidation (UMA)

Require: Episodic Store M_E , Semantic Store M_S , Threshold τ
Ensure: Updated Semantic Store M_S

```

1: for episode in  $M_E$ .GetRecent() do
2:   {Extract key facts from episode}
3:   facts  $\leftarrow$  LLMExtractFacts(episode.trace)
4:   for fact in facts do
5:     {Compute importance based on success and frequency}
6:     importance  $\leftarrow \alpha \cdot \text{episode.success} + \beta \cdot \text{Frequency}(\text{fact}, M_E)$ 
7:     if importance  $> \tau$  then
8:       {Check for duplicates using embedding similarity}
9:       similar  $\leftarrow M_S$ .FindSimilar(fact, threshold = 0.9)
10:      if |similar| = 0 then
11:        {Add new fact to knowledge graph}
12:        entities  $\leftarrow$  ExtractEntities(fact)
13:        relations  $\leftarrow$  ExtractRelations(fact)
14:         $M_S$ .AddToGraph(entities, relations)
15:         $M_S$ .AddEmbedding(fact)
16:      else
17:        {Update existing fact with new evidence}
18:         $M_S$ .UpdateConfidence(similar[0], importance)
19:      end if
20:    end if
21:  end for
22:  {Archive processed episode}
23:   $M_E$ .Archive(episode)
24: end for
25: return  $M_S$ 

```

J. MARS Benchmark Specification

J.1. Task Schema

Each task in MARS follows this JSON schema:

```

{
  "id": "string",
  "category": "scientific_research|software_engineering|strategic_planning",
  "description": "string",
  "complexity_level": "low|medium_low|medium|medium_high|high",
  "subtasks": [
    {
      "id": "string",
      "description": "string",
      "complexity": 0.0-1.0,
      "required_tools": ["tool_names"],
      "dependencies": ["subtask_ids"],
      "optimal_pattern": "single|parallel|hierarchical|consensus",
      "estimated_steps": integer,
      "evaluation_criteria": {
        "completeness": "rubric",
        "correctness": "rubric",
        "efficiency": "rubric"
      }
    }
  ],
}

```

```

    "ground_truth": {
      "final_answer": "string_or_structured",
      "key_insights": ["strings"],
      "required_tool_calls": [{"tool": "name", "args": {...}}]
    },
    "metadata": {
      "domain": "string",
      "difficulty_rating": 1-5,
      "estimated_time_minutes": integer,
      "annotator_id": "string",
      "annotation_date": "ISO_date"
    }
  }
}

```

J.2. Task Category Details

J.2.1. SCIENTIFIC RESEARCH (847 TASKS)

Domains: Biology (284), Physics (198), Computer Science (212), Chemistry (153)

Task Types:

- Literature Review (35%): Synthesize research on a specific topic
- Methodology Analysis (25%): Evaluate experimental approaches
- Gap Identification (20%): Find unexplored research directions
- Hypothesis Generation (20%): Propose testable hypotheses

Example Task:

“Conduct a systematic review of transformer efficiency techniques published since 2022. Identify the top 5 approaches by computational savings, analyze their trade-offs, and propose a novel combination that could achieve better results.”

Optimal Pattern Distribution:

- Paper retrieval: Parallel (multiple databases)
- Individual paper analysis: Single agent
- Cross-paper synthesis: Hierarchical (coordinator + analysts)
- Novel proposal generation: Consensus (multiple perspectives)

J.2.2. SOFTWARE ENGINEERING (1,124 TASKS)

Domains: Web Development (412), Systems Programming (298), ML Engineering (248), DevOps (166)

Task Types:

- Bug Localization & Fix (30%): Identify and resolve issues
- Code Review (25%): Evaluate code quality and suggest improvements
- Feature Implementation (25%): Add new functionality
- Test Generation (20%): Create comprehensive test suites

Example Task:

“Debug the performance regression in the authentication module that appeared after commit abc123. Identify the root cause, implement a fix, write regression tests, and ensure backward compatibility.”

J.2.3. STRATEGIC PLANNING (876 TASKS)

Domains: Business Strategy (324), Policy Analysis (287), Product Design (265)

Task Types:

- Market Analysis (30%): Evaluate competitive landscape
- Decision Synthesis (25%): Integrate multiple factors for recommendations
- Risk Assessment (25%): Identify and evaluate potential risks
- Scenario Planning (20%): Develop contingency strategies

J.3. Annotation Process

Phase 1: Task Creation

1. Domain experts created initial tasks based on real-world scenarios
2. Tasks were reviewed for clarity and completeness
3. Subtask decomposition was performed by experienced annotators

Phase 2: Pattern Labeling

1. Three annotators independently labeled optimal patterns for each subtask
2. Inter-annotator agreement: Cohen's $\kappa = 0.78$
3. Disagreements resolved through discussion and empirical validation

Phase 3: Validation

1. Subset (10%) executed with different patterns to validate labels
2. Ground truth answers verified by domain experts
3. Quality control: 5% of tasks randomly re-annotated monthly


J.4. Evaluation Protocol

Task Success Criteria:

- Binary success: Task achieves $\geq 80\%$ subtask completion
- Subtask success: Output meets completeness, correctness, efficiency criteria
- Partial credit: Weighted by subtask importance

Orchestration Metrics:

- Orchestration Accuracy = $\frac{\text{Correct patterns}}{\text{Total subtasks}}$
- Adaptation Rate = $\frac{\text{Successful adaptations}}{\text{Adaptation attempts}}$
- Memory Utilization = $\frac{\text{Used retrieved memories}}{\text{Total retrieved}}$



Key Finding:
AMORE degrades more gracefully
as task complexity increases.
On 20+ subtask tasks:
• Single-agent: 5.8%
• AgentOrchestra: 24.8%
• **AMORE: 35.2%**

Figure 11. Performance vs. Task Complexity (Number of Subtasks)

K. Additional Experimental Results

K.1. Per-Domain MARS Results

Table 42. Detailed MARS Results by Domain

Domain	AMORE	AgentOrch.	Single	Δ Best
Scientific Research				
Biology	54.2	45.1	31.8	+9.1
Physics	48.7	41.3	28.4	+7.4
Computer Science	53.4	44.2	33.1	+9.2
Chemistry	50.1	42.8	29.7	+7.3
Software Engineering				
Web Development	61.2	52.4	41.2	+8.8
Systems Programming	55.8	48.7	36.5	+7.1
ML Engineering	58.2	49.1	38.8	+9.1
DevOps	56.4	50.2	37.4	+6.2
Strategic Planning				
Business Strategy	44.1	35.8	23.1	+8.3
Policy Analysis	42.8	34.2	21.5	+8.6
Product Design	48.7	40.5	29.8	+8.2

K.2. Scalability Analysis

Table 43. Performance vs. Number of Subtasks

Subtasks	AMORE	AgentOrch.	HALO	MetaGPT	Single
1-5	68.4	62.1	59.8	56.2	58.2
6-10	54.2	46.8	44.1	41.5	35.4
11-15	48.1	38.4	35.2	32.8	21.7
16-20	41.5	31.2	27.5	24.1	12.3
20+	35.2	24.8	20.1	17.2	5.8
Degradation	-49%	-60%	-66%	-69%	-90%

K.3. Pattern Selection Analysis

Table 44. CAR Pattern Prediction Accuracy by Complexity

Complexity	Overall	Single	Parallel	Hier.	Cons.
Low	92.4%	94.1%	88.2%	85.3%	78.1%
Medium-Low	89.7%	91.2%	89.5%	87.4%	82.3%
Medium	87.2%	85.4%	88.1%	88.7%	85.6%
Medium-High	85.8%	81.2%	84.3%	89.2%	88.1%
High	83.1%	75.4%	78.6%	86.8%	91.2%
Average	88.3%	85.5%	85.7%	87.5%	85.1%

K.4. Cost-Performance Detailed Analysis

Table 45. Detailed Cost Analysis (Cost per Task in \$)

Method	Sci.	SWE	Strat.	Avg	Succ/\$
GPT-4 + ReAct	2.45	1.87	2.10	2.14	15.0
AutoGen	4.21	3.52	3.82	3.85	9.4
MetaGPT	4.58	3.91	4.14	4.21	9.3
HALO	5.67	4.78	4.91	5.12	8.2
AgentOrchestra	7.12	5.98	6.31	6.47	6.8
AMORE	4.15	3.41	3.90	3.82	13.7

K.5. RCM Checkpoint Statistics

Table 46. Checkpoint Action Distribution

Action	Frequency	Success After	Avg Cost
Proceed	71.2%	N/A	1.0×
Retry	14.8%	68.4%	1.8×
Escalate	10.3%	72.1%	2.5×
Replan	3.7%	54.2%	3.2×




Figure 12. Unified Memory Architecture (UMA). Three-tier memory system showing how agents share and consolidate memory.

K.6. UMA Memory Statistics

Table 47. Memory Utilization Across Task Types

Category	Working	Episodic	Semantic	Consol. Rate
Scientific	45.2K	128.4K	23.1K	18.0%
Software	38.7K	95.2K	31.4K	33.0%
Strategic	52.1K	156.8K	28.7K	18.3%
Average	45.3K	126.8K	27.7K	21.8%

L. Implementation Details

L.1. Model Configuration

Table 48. Model Configuration

Component	Configuration
Coordinator LLM	GPT-4-Turbo (gpt-4-1106-preview)
Worker LLM	GPT-3.5-Turbo (gpt-3.5-turbo-1106)
Quality Judge	GPT-4 (for LLM-as-judge)
CAR MLP	3 layers, 256 hidden units, ReLU
Episodic Memory	ChromaDB with OpenAI embeddings
Semantic Memory	Neo4j + text-embedding-3-small

L.2. Hyperparameters

Table 49. Hyperparameter Settings

Parameter	Value	Tuning Range
θ_{high}	0.80	[0.70, 0.90]
θ_{low}	0.50	[0.40, 0.60]
τ_0 (base threshold)	0.60	[0.50, 0.70]
λ (threshold adaptation)	0.20	[0.10, 0.30]
Max retries	2	[1, 3]
Max escalations	2	[1, 3]
Consolidation threshold	0.70	[0.60, 0.80]
Quality weights ($\lambda_1, \lambda_2, \lambda_3$)	(0.4, 0.3, 0.3)	Grid search

L.3. Prompt Templates

L.3.1. QUALITY ASSESSMENT PROMPT

Given the subtask and its output, rate the quality on three dimensions.

Subtask: {subtask_description}

Output: {output}

Rate each dimension from 0.0 to 1.0:

1. **Completeness:** Does the output fully address all aspects of the subtask?
 - 1.0: All aspects fully addressed
 - 0.5: Most aspects addressed, some gaps
 - 0.0: Major aspects missing
2. **Coherence:** Is the output internally consistent and well-structured?
 - 1.0: Perfectly coherent and logical
 - 0.5: Mostly coherent with minor issues
 - 0.0: Contradictory or confusing
3. **Usefulness:** Will this output support downstream tasks effectively?
 - 1.0: Highly useful, provides clear foundation
 - 0.5: Somewhat useful, needs supplementation
 - 0.0: Not useful for downstream tasks

Provide ratings in JSON format:

```
{"completeness": X.X, "coherence": X.X, "usefulness": X.X}
```

L.3.2. REFLECTION PROMPT FOR REPLANNING

The following subtask produced unsatisfactory results despite multiple attempts.

Subtask: {subtask_description}

Last Output: {output}

Quality Score: {quality_score}

Remaining subtasks:

```
{remaining_subtasks}
```

Analyze what went wrong and propose a revised plan:

1. Identify the root cause of the failure
2. Suggest how to restructure the remaining work
3. Recommend which subtasks need different approaches

Provide your analysis and revised plan.

L.4. Compute Resources

- Training CAR: 4 hours on 1x NVIDIA A100 (40GB)
- Evaluation: 200 hours total API time across all experiments
- Memory systems: 32GB RAM for vector DB, 16GB for graph DB
- Total cost: Approximately \$8,500 in API calls

L.5. Cost Accounting Methodology

We provide detailed cost breakdowns for reproducibility.

API Pricing (as of experiments, November 2024):

- GPT-4-Turbo: \$0.01/1K input tokens, \$0.03/1K output tokens
- GPT-3.5-Turbo: \$0.0005/1K input tokens, \$0.0015/1K output tokens
- text-embedding-3-small: \$0.00002/1K tokens

Per-Subtask Cost Breakdown:

Table 50. Average Cost per Subtask by Pattern (\$)

Component	Single	Parallel	Hier.	Cons.
Coordinator (GPT-4)	0.15	0.28	0.45	0.62
Workers (GPT-3.5)	–	0.08	0.24	0.35
CAR Routing	0.02	0.02	0.02	0.02
Quality Assessment	0.03	0.03	0.03	0.03
Memory Retrieval	0.01	0.02	0.03	0.03
Total	0.21	0.43	0.77	1.05

Average Prompt Lengths:

- System prompt: 800-1200 tokens (varies by agent role)
- Task context: 500-2000 tokens (depends on subtask complexity)
- Tool outputs: 200-1500 tokens per tool call
- Memory retrieval: 300-800 tokens (top-5 episodic + top-3 semantic)
- Average output: 400-1200 tokens per agent turn

Dollar Conversion: All costs reported in the paper are direct API costs. We do not include infrastructure costs (compute for running experiments, database hosting) as these vary by deployment. The \$3.82 average cost per task on MARS corresponds to approximately 380K input tokens and 95K output tokens across all agents and patterns used.

M. Extended Related Work

M.1. Test-Time Compute and Reasoning

Recent work has explored scaling computation at inference time rather than training time. OpenAI’s o1 (OpenAI, 2024) demonstrates that learned chain-of-thought via reinforcement learning can dramatically improve reasoning capabilities. Snell et al. (2024) show that moving computation from training to test time can allow smaller models to outperform 14× larger models.

This relates to AMORE’s adaptive approach: rather than always using expensive multi-agent orchestration, we allocate more “test-time compute” (via more agents or more sophisticated patterns) only when task complexity warrants it.

M.2. Cognitive Architectures for AI

AMORE’s UMA draws inspiration from cognitive architectures like ACT-R and SOAR (Laird, 2022), which distinguish between working, procedural, and declarative memory. Recent work has begun exploring these ideas in LLM contexts:

- Sumers et al. (2023) propose a cognitive architecture framework for LLM agents

- Park et al. (2023) use memory streams for generative agents
- Zhong et al. (2024) develop memory banks for long-horizon tasks

M.3. Quality Estimation and Self-Correction

RCM’s quality assessment relates to work on LLM self-evaluation:

- Self-Refine (Madaan et al., 2023) iteratively improves outputs (NeurIPS 2023)
- Reflexion (Shinn et al., 2024) uses verbal reinforcement learning
- Constitutional AI (Bai et al., 2022) employs critique-revision cycles
- SelfCheck (Miao et al., 2023) zero-shot verifies reasoning steps (ICLR 2024)

AMORE extends these ideas to multi-agent settings with escalation as a novel intervention mechanism.

M.4. Adaptive Systems in Machine Learning

The concept of adaptive computation has precedents in ML:

- Adaptive computation time (Kalinin et al., 2025) for RNNs
- Mixture of Experts (Pavlitska et al., 2025) for conditional computation
- Early exit networks (Simioni et al., 2025) for efficiency

AMORE applies similar principles to multi-agent orchestration, selecting computation “depth” (number and type of agents) based on task complexity.

N. Detailed Feature Extraction Latency Analysis

Reviewer concern: “How was the 0.5s feature-extraction overhead per subtask measured given multiple LLM calls (step estimation, paraphrases for ambiguity)? Please provide concrete timing distributions and amortization details.”

We provide comprehensive latency measurements from production API calls with detailed breakdowns.

N.1. Per-Feature Timing Distributions

Table 51 presents timing statistics for each CAR feature extraction operation:

Table 51. Feature extraction latency breakdown (ms). Measured over 1,000 subtasks using GPT-4-turbo API with batch optimization.

Feature	P50	P95	P99	Method
f_{length} (steps)	45	82	127	Single LLM call
f_{tools} (tool count)	38	71	98	Single LLM call
$f_{\text{knowledge}}$ (domains)	41	78	112	Single LLM call
$f_{\text{ambiguity}}$ (paraphrase)	156	243	312	3 LLM calls
f_{deps} (dependencies)	12	24	38	DAG lookup
f_{critical} (critical path)	8	15	22	DAG analysis
f_{history} (success rate)	3	7	12	DB lookup
Sequential total	303	520	721	—
Batched total	187	312	445	Parallel calls
With caching	98	178	267	Similar subtasks

N.2. Amortization Strategy

The reported 0.5s overhead is achieved through several optimizations:

(1) Batched LLM calls: Features f_{length} , f_{tools} , and $f_{\text{knowledge}}$ are extracted via a single structured prompt, reducing 3 sequential calls to 1 parallel batch. This accounts for $\sim 120\text{ms}$ savings.

(2) Cached embeddings: For $f_{\text{ambiguity}}$, we cache paraphrase embeddings across similar subtask descriptions (Jaccard similarity > 0.7), achieving 67% cache hit rate on MARS.

(3) Pre-computed DAG features: f_{deps} and f_{critical} are computed once during task decomposition and reused, adding negligible overhead.

(4) Async pipeline: Feature extraction runs in parallel with prior subtask execution, hiding latency for 78% of subtasks.

N.3. Latency vs. Accuracy Trade-off

Table 52. CAR routing accuracy vs. feature extraction budget.

Configuration	Latency (ms)	Accuracy (%)	Cost (\$)
Minimal (3 features)	52	71.3	0.002
Standard (7 features)	187	83.7	0.008
Extended (7 + ensemble)	423	85.2	0.019
Full (+ cross-validation)	891	86.1	0.041

The standard 7-feature configuration provides the optimal trade-off, achieving 83.7% routing accuracy at 187ms median latency and \$0.008 cost per subtask.

O. CAR Label Stability and λ Sensitivity Analysis

Reviewer concern: “For CAR’s ‘optimal pattern’ labels, how sensitive are results to the cost weight λ and to the small training set (423 subtasks)? Could you report label stability analyses and router sensitivity to λ ?”

O.1. Label Stability Under Different λ Values

We analyze how the “optimal pattern” labels change as the cost-quality trade-off weight λ varies in the scoring function:

$$\text{Score}(p) = \text{Quality}(p) - \lambda \cdot \text{NormalizedCost}(p)$$

Table 53. Label distribution and stability across λ values.

λ	Single	Parallel	Hier.	Cons.	Flip Rate
0.0 (quality only)	18%	22%	31%	29%	—
0.1	24%	28%	28%	20%	12.3%
0.2 (default)	31%	32%	24%	13%	8.7%
0.3	38%	33%	20%	9%	6.2%
0.5	47%	31%	16%	6%	4.1%
1.0 (cost only)	71%	19%	8%	2%	—

Flip Rate measures the percentage of subtasks whose optimal label changes when λ varies by ± 0.05 . The default $\lambda = 0.2$ shows reasonable stability (8.7% flip rate).

O.2. Router Performance Across λ Values

Table 54. AMORE performance when CAR is trained with different λ values.

Training λ	AgentBench	WebArena	MARS	Cost
0.1 (quality-biased)	60.2%	30.1%	53.1%	\$4.21
0.2 (balanced)	59.7%	29.7%	52.3%	\$3.47
0.3 (cost-aware)	58.4%	28.9%	51.2%	\$2.89
0.5 (cost-focused)	55.8%	27.2%	48.7%	\$2.31

Insight: $\lambda = 0.2$ provides the best accuracy-cost trade-off. Lower λ (0.1) yields +0.5-0.8% accuracy at +21% cost; higher λ (0.3+) reduces cost but degrades accuracy.

Table 55. CAR routing accuracy vs. training set size (bootstrap analysis).

Training Size	Accuracy (%)	Std. Dev.	95% CI
100 subtasks	72.1	4.3	[68.2, 76.0]
200 subtasks	78.4	2.8	[75.9, 80.9]
300 subtasks	81.9	2.1	[80.0, 83.8]
423 subtasks (full)	83.7	1.4	[82.4, 85.0]
600 subtasks [†]	85.1	1.1	[84.1, 86.1]
800 subtasks [†]	85.8	0.9	[85.0, 86.6]

[†]Projected via data augmentation (paraphrasing).

O.3. Training Set Size Sensitivity

We analyze CAR stability with different training set sizes via bootstrap sampling:

Diminishing returns: Performance plateaus around 500 samples, consistent with VC-dimension analysis. The 423-subtask dataset captures 97% of achievable routing accuracy.

O.4. Cross-Seed and Cross-Model Label Stability

Reviewer concern: “Can you report label stability across cost settings and different seeds/models used for the exhaustive pattern runs?”

We analyze label stability across (1) random seeds for pattern execution, (2) different backbone models, and (3) budget normalization schemes.

Table 56. CAR label stability analysis across seeds, models, and cost settings.

Variation	Label Agreement	CAR Acc. Impact
<i>Cross-Seed Stability (5 seeds, same model)</i>		
Seed 42 vs. Seed 123	96.8%	±0.3%
Seed 42 vs. Seed 456	95.9%	±0.5%
Seed 42 vs. Seed 789	96.2%	±0.4%
Seed 42 vs. Seed 101	97.1%	±0.2%
Average cross-seed	96.5%	±0.35%
<i>Cross-Model Variance (same seed)</i>		
GPT-4 vs. Claude-3.5	93.4%	±1.2%
GPT-4 vs. Llama-3-70B	89.7%	±2.1%
Claude-3.5 vs. Llama-3-70B	88.2%	±2.4%
Average cross-model	90.4%	±1.9%
<i>Budget Normalization Schemes</i>		
Per-pattern fixed (\$5/\$10/\$15/\$20)	94.2%	±0.8%
Token-based (50K/100K/150K/200K)	92.8%	±1.1%
Time-based (30s/60s/90s/120s)	91.5%	±1.4%

Key findings:

- **Cross-seed stability:** 96.5% label agreement across 5 random seeds, indicating low stochastic variance in pattern execution outcomes.
- **Cross-model variance:** 90.4% agreement between different backbone models. Disagreements primarily occur on medium-complexity tasks where multiple patterns achieve similar quality.
- **Budget normalization:** Dollar-based normalization shows highest stability (94.2%); token-based and time-based show slightly more variance due to model-specific efficiency differences.

Pareto analysis: $\lambda \in [0.10, 0.25]$ forms the Pareto frontier. The default $\lambda = 0.20$ achieves the best stability (94.7%) while maintaining competitive success rates. Users can adjust λ based on cost constraints without significant accuracy degradation.

Table 57. Detailed λ sensitivity with confidence intervals (5 seeds each).

λ	MARS Success	Avg Cost	Label Stability	Pareto Optimal?
0.05	53.8 $\pm$ 0.9%	\$4.82	91.2%	No (dominated by 0.10)
0.10	53.1 $\pm$ 0.7%	\$4.21	93.7%	Yes
0.15	52.7 $\pm$ 0.6%	\$3.89	94.5%	Yes
0.20	52.3 $\pm$ 0.5%	\$3.47	94.7%	Yes (default)
0.25	51.4 $\pm$ 0.6%	\$3.12	93.8%	Yes
0.30	50.2 $\pm$ 0.8%	\$2.89	92.4%	No (suboptimal quality)

P. Native Decomposition Baseline Comparison

Reviewer concern: “The shared decomposition DAG improves parity but may undercut baselines whose strength lies in native planning. Can you include a setting where each baseline uses its own decomposition and report both settings?”

We evaluate all methods using both (1) shared AMORE decomposition and (2) each method’s native planning/decomposition.

P.1. Native Decomposition Results

Table 58. Comparison: Shared vs. Native decomposition on MARS benchmark.

Method	Shared Decomp.		Native Decomp.	
	Acc.	Cost	Acc.	Cost
Single-Agent	38.9%	\$1.24	37.2%	\$1.31
AutoGen	43.2%	\$3.89	41.8%	\$4.21
MetaGPT	44.1%	\$4.12	46.3%	\$4.67
CAMEL	41.7%	\$3.67	40.2%	\$3.94
HALO	45.8%	\$5.21	48.2%	\$5.89
AgentOrchestra	44.1%	\$5.87	43.7%	\$6.12
AMORE	52.3%	\$3.47	52.3%	\$3.47

Key findings:

1. **MetaGPT** benefits most from native decomposition (+2.2%), as its SOP-based planning is tightly coupled with role assignment.
2. **HALO** gains +2.4% with native MCTS-based workflow search, confirming its planning strength.
3. **AMORE** maintains the same performance as it uses its own CAR-guided decomposition by design.
4. Most other baselines perform comparably or slightly worse with native decomposition, suggesting the shared DAG is not systematically disadvantaging them.

P.2. Cross-Benchmark Native Decomposition

Table 59. Native decomposition results across all benchmarks.

Method	AgentBench	WebArena	MARS
MetaGPT (shared)	51.2%	23.1%	44.1%
MetaGPT (native)	52.8%	24.7%	46.3%
Δ	+1.6%	+1.6%	+2.2%
HALO (shared)	53.9%	24.4%	45.8%
HALO (native)	55.1%	25.8%	48.2%
Δ	+1.2%	+1.4%	+2.4%
AMORE	59.7%	29.7%	52.3%
Δ vs best native	+4.6%	+3.9%	+4.1%

Conclusion: Even against baselines using their native decomposition (which favors them), AMORE maintains 4-5% absolute improvement, confirming that adaptive orchestration provides value beyond decomposition quality.

Q. Non-Parametric kNN Router Baseline

Reviewer concern: “Can you add a non-parametric kNN router baseline for CAR (using the same features or a standard embedding), in light of recent routing studies showing strong kNN performance?”

We implement and evaluate kNN-based routing alternatives to our learned CAR model.

Q.1. kNN Router Variants

We evaluate three kNN configurations:

1. **kNN-Feature:** Uses the same 7-dimensional CAR feature vector with Euclidean distance
2. **kNN-Embed:** Uses text-embedding-3-large (3072-dim) embeddings of subtask descriptions
3. **kNN-Hybrid:** Concatenates normalized CAR features with compressed (128-dim) embeddings

Table 60. Routing accuracy comparison: CAR vs. kNN variants.

Router	Routing Acc.	Latency (ms)	Memory
CAR (learned)	83.7%	187	2.1 MB
kNN-Feature (k=5)	76.2%	23	0.3 MB
kNN-Feature (k=15)	78.4%	31	0.3 MB
kNN-Embed (k=5)	79.8%	156	12.4 MB
kNN-Embed (k=15)	81.3%	178	12.4 MB
kNN-Hybrid (k=10)	82.1%	198	6.8 MB

Q.2. End-to-End Performance

Table 61. MARS benchmark performance with different routers.

Router	Success (%)	Quality	Cost	Latency
CAR (learned)	52.3%	0.71	\$3.47	1.00×
kNN-Feature (k=5)	47.8%	0.64	\$3.21	0.89×
kNN-Feature (k=15)	49.2%	0.66	\$3.34	0.91×
kNN-Embed (k=15)	50.1%	0.68	\$3.89	1.02×
kNN-Hybrid (k=10)	51.2%	0.69	\$3.67	1.05×

Analysis:

- **kNN-Feature** is fastest but least accurate—the 7-dim feature space lacks discriminative power for fine-grained routing.
- **kNN-Embed** improves significantly (+3.6% routing accuracy) but at higher memory cost and similar latency.
- **kNN-Hybrid** approaches CAR performance (82.1% vs. 83.7%) and represents a viable alternative for scenarios requiring interpretability.
- **CAR’s advantage** comes from learning nonlinear decision boundaries that capture complex feature interactions, which kNN cannot model efficiently in 7-dim space.

Recommendation: For production deployments, learned CAR provides the best accuracy. For interpretability or debugging, kNN-Hybrid offers a competitive 82% routing accuracy with transparent nearest-neighbor explanations.

R. Per-Pattern Escalation Overhead Analysis

Reviewer concern: “What is the escalation overhead (tokens, wall-clock) per additional pattern level, and how often do rollbacks occur in the wild beyond the reported 7%?”

R.1. Escalation Cost Breakdown

Notes: Token counts include context re-encoding for new agents. Wall-clock includes agent initialization and coordination setup.

Table 62. Per-escalation overhead by pattern transition.

Transition	Tokens	Wall-clock	API Cost	Freq.
Single → Parallel	2,340	3.2s	\$0.047	8.7%
Single → Hierarchical	4,120	6.1s	\$0.082	3.2%
Parallel → Hierarchical	3,890	5.4s	\$0.078	5.1%
Parallel → Consensus	6,210	8.7s	\$0.124	2.8%
Hierarchical → Consensus	4,780	7.2s	\$0.096	4.3%
Average escalation	4,268	6.1s	\$0.085	—

R.2. Rollback Analysis by Benchmark

Table 63. Rollback frequency and outcomes by benchmark.

Benchmark	Rollback %	Quality Δ	Wasted Cost
AgentBench	5.8%	+0.12	\$0.31/rollback
WebArena	8.9%	+0.08	\$0.42/rollback
MARS - Scientific	6.2%	+0.15	\$0.28/rollback
MARS - Software	4.1%	+0.09	\$0.24/rollback
MARS - Strategic	9.7%	+0.11	\$0.47/rollback
Overall	7.0%	+0.11	\$0.34/rollback

Key insights:

- Strategic planning tasks have highest rollback rate (9.7%) due to subjective quality criteria
- Software engineering has lowest (4.1%) due to verifiable outputs
- Average quality improvement from rollback (+0.11) justifies the \$0.34 wasted cost
- Only 2.1% of rollbacks result in final quality *lower* than pre-escalation (double-rollback cases)

R.3. Cumulative Escalation Statistics

Table 64. Escalation depth distribution across 2,847 MARS tasks.

Escalation Depth	Tasks	Avg. Quality	Avg. Cost
0 (no escalation)	71.2%	0.68	\$2.41
1 escalation	19.3%	0.74	\$4.12
2 escalations	7.8%	0.78	\$6.34
3+ escalations	1.7%	0.81	\$9.87

Conclusion: Most tasks (71%) complete without escalation. Each escalation level adds ~\$1.5-2.5 cost but improves quality by 0.03-0.06. The escalation budget cap (3 levels) prevents runaway costs.

S. UMA Contamination Audit and Error Analysis

Reviewer concern: “Could you release quantitative ‘memory audit’ outcomes and error cases, and measure how often incorrect episodic entries get consolidated and later corrected?”

S.1. Contamination Detection Results

We audit 10,000 memory write operations across all MARS tasks:

Definitions:

- **False positive:** Valid memory blocked incorrectly
- **False negative:** Contaminated memory accepted

Table 65. UMA contamination audit results.

Metric	Episodic	Semantic	Working
Total write attempts	8,234	2,891	14,567
Blocked (cross-task)	312 (3.8%)	89 (3.1%)	0 (0%)
Blocked (confidence $< \tau$)	421 (5.1%)	156 (5.4%)	0 (0%)
Accepted writes	7,501	2,646	14,567
False positives	23 (0.28%)	8 (0.28%)	—
False negatives	67 (0.81%)	31 (1.07%)	—

S.2. Incorrect Consolidation Analysis

Table 66. Episodic $\rightarrow$ Semantic consolidation error analysis.

Error Type	Occurrences	Rate
Incorrect facts consolidated	31	1.17%
Incomplete consolidation	47	1.78%
Redundant consolidation	89	3.36%
Total problematic	167	6.31%
Later corrected	142	85.0% of errors
Persisted errors	25	0.95% of all

Self-correction mechanism: 85% of consolidation errors are detected and corrected within 3 subsequent task executions through:

1. Consistency checking against new episodic memories
2. Confidence decay for unverified entries
3. Explicit contradiction detection

S.3. Error Case Examples

Example 1: False positive (blocked valid memory)

Entry: “The API rate limit is 100 requests/minute”

Reason blocked: Cross-task reference detected (mentioned task-id from different session)

Actual issue: Generic rate limit applies across tasks—incorrectly flagged

Example 2: False negative (accepted contaminated memory)

Entry: “User prefers Python over JavaScript”

Issue: Preference from Task A incorrectly written without task scope

Correction: Flagged when Task B user expressed JavaScript preference

S.4. Impact on Downstream Performance

Table 67. Performance impact of UMA errors on subsequent tasks.

Scenario	Success Rate	Quality Score
No memory errors	53.1%	0.72
With false positives	51.8%	0.70
With false negatives	49.2%	0.67
With consolidation errors	50.7%	0.69

Conclusion: UMA’s safeguards reduce contamination to 0.95% persistent error rate, with downstream performance impact of -2-4% when errors occur.

S.5. MemoryBench-Style Evaluation Metrics

Reviewer concern: “Could you provide a more detailed breakdown of UMA’s contribution using memory-centric metrics (e.g., retrieval precision/recall, contamination rate, provenance utilization), possibly aligned with MemoryBench-style evaluations?”

We evaluate UMA against baseline memory systems using standardized memory-centric metrics:

Table 68. UMA memory-centric metrics comparison (MemoryBench-style evaluation).

Metric	UMA	RAG	MemGPT	AriGraph
<i>Retrieval Quality</i>				
Precision@5	0.84	0.72	0.76	0.81
Recall@10	0.89	0.78	0.82	0.86
MRR	0.82	0.72	0.76	0.84
<i>Memory Integrity</i>				
Contamination Rate	3.1%	8.4%	6.2%	4.8%
Cross-task Leakage	0.3%	2.1%	1.4%	0.8%
Conflict Resolution Acc.	87.2%	N/A	72.4%	81.5%
<i>Provenance & Consolidation</i>				
Provenance Utilization	78.4%	41.2%	58.7%	68.9%
Consolidation Yield	91.5%	N/A	84.2%	88.1%
Temporal Coherence	94.2%	82.1%	87.5%	91.8%
<i>Efficiency</i>				
Retrieval Latency (ms)	68	45	125	185
Storage Overhead (MB)	2.4	1.2	2.1	4.5
Update Throughput (ops/s)	842	1,245	523	312

Metric Definitions:

- **Provenance Utilization:** Percentage of retrieved memories where source attribution is correctly used by downstream agents
- **Consolidation Yield:** Percentage of episodic memories successfully abstracted to semantic memory without information loss
- **Temporal Coherence:** Accuracy of temporal ordering in retrieved memory sequences
- **Cross-task Leakage:** Rate of task-specific information incorrectly accessible by unrelated tasks

Key Findings:

1. UMA achieves highest retrieval quality (Precision@5: 0.84) while maintaining lowest contamination (3.1%)
2. Provenance utilization (+37% over RAG) enables agents to assess memory reliability
3. The three-tier architecture trades modest latency increase (+23ms vs RAG) for substantially better integrity guarantees
4. AriGraph’s knowledge graph structure achieves competitive MRR (0.84) but at $2.7\times$ higher latency

T. Open-Weight Model Evaluation

Reviewer concern: “How do results change with open-weight models (e.g., Qwen/DeepSeek) to improve reproducibility?”

T.1. Full Open-Weight Results

Key findings:

- **DeepSeek-V3** achieves 94% of GPT-4 performance at 19% of the cost
- **Qwen2.5-72B** provides best cost-efficiency ratio for self-hosted deployments
- Open-weight models show consistent 4-8% accuracy gap vs. closed-source
- AMORE’s adaptive benefits transfer across all backbones (14-17% vs. single-agent)

Table 69. AMORE performance with open-weight backbone models.

Backbone	AgentBench	WebArena	MARS	Cost [†]
<i>Closed-source (reference)</i>				
GPT-4-turbo	59.7%	29.7%	52.3%	\$3.47
Claude-3-Opus	58.2%	28.4%	50.8%	\$4.12
<i>Open-weight</i>				
Qwen2.5-72B-Instruct	54.8%	26.2%	47.9%	\$0.89
DeepSeek-V3	56.1%	27.8%	49.4%	\$0.67
Llama-3.1-70B-Instruct	52.3%	24.9%	45.2%	\$0.78
Mixtral-8x22B	49.7%	22.1%	42.8%	\$0.52
Yi-1.5-34B	47.2%	20.8%	40.1%	\$0.41

[†]Self-hosted cost estimate (A100 GPU @ \$2/hr).

Table 70. AMORE component effectiveness with open-weight backbones (MARS).

Configuration	DeepSeek-V3	Qwen2.5-72B	Llama-3.1-70B
Full AMORE	49.4%	47.9%	45.2%
– CAR (random routing)	42.1%	40.8%	38.9%
– RCM (no checkpoints)	45.7%	44.2%	41.8%
– UMA (isolated memory)	46.8%	45.3%	43.1%
CAR contribution	+7.3%	+7.1%	+6.3%
RCM contribution	+3.7%	+3.7%	+3.4%
UMA contribution	+2.6%	+2.6%	+2.1%

T.2. Component-Level Analysis

Observation: AMORE component contributions are consistent across backbones, demonstrating architecture-agnostic benefits.

T.3. Reproducibility Configuration

For full reproducibility, we recommend:

```
{
  "backbone": "deepseek-v3",
  "api_base": "https://api.deepseek.com",
  "temperature": 0.0,
  "car_model": "car_weights_deepseek.pt",
  "rcm_threshold_high": 0.78,
  "rcm_threshold_low": 0.48
}
```

Adjusted RCM thresholds account for DeepSeek’s slightly lower calibration than GPT-4.

U. Proxy Implementation Details for Reproducibility

Reviewer concern: “Given the limited availability of some orchestrator baselines, can you publish your proxy implementations and unify reward formulations so the community can reproduce/extend the comparisons?”

U.1. Proxy Implementation Specifications

We provide detailed specifications for our proxy implementations:

xRouter-Proxy:

- PPO-based router with 3-layer MLP policy (256, 128, 64 units)
- State: concatenation of subtask embedding + historical success rates
- Action space: 4 patterns (same as AMORE)
- Reward: $r = 0.7 \cdot \text{quality} + 0.3 \cdot (1 - \text{norm_cost})$

- Training: 50K episodes on MARS training split

DAAO-Proxy:

- VAE-based difficulty estimator (encoder: 512-256-64, decoder: 64-256-512)
- Difficulty signal: latent z mapped to $[0,1]$ via sigmoid
- Operator selection: difficulty thresholds at $[0.3, 0.6, 0.8]$ for pattern assignment
- Layered execution: sequential processing of difficulty-sorted subtasks

MoMA-Proxy:

- Model selection router (not orchestration pattern router)
- Adapted to select patterns instead of models
- Gating network: 2-layer MLP with softmax output
- Training: supervised on AMORE’s ground-truth labels for fair comparison

U.2. Unified Reward Formulation

All proxy baselines use the same reward formulation for fair comparison:

$$R(s_i, p_i, o_i) = \alpha \cdot Q(o_i) + \beta \cdot (1 - C(p_i)/C_{\max}) + \gamma \cdot \mathbb{I}[\text{success}] \quad (44)$$

where $\alpha = 0.5$, $\beta = 0.2$, $\gamma = 0.3$, $Q(\cdot)$ is the quality score, $C(\cdot)$ is cost, and $C_{\max}$ is the consensus pattern cost.

U.3. Code Availability

All proxy implementations are available at:

<https://anonymous.4open.science/r/AMORE-ICML2026-FB66/proxies/>

Including:

- `xrouter_proxy.py`: PPO implementation with training script
- `daao_proxy.py`: VAE difficulty estimator and operator allocation
- `moma_proxy.py`: Adapted gating network
- `unified_reward.py`: Shared reward computation
- `reproduce_baselines.sh`: Full reproduction script

V. Per-Environment Breakdown with Cost and Latency

Reviewer concern: “Can you release per-environment breakdown tables (AgentBench/WebArena), token costs, and wall-clock latencies per pattern to help reproduce/validate cost savings and to guide deployment decisions?”

V.1. AgentBench Detailed Breakdown

Table 71. AgentBench per-environment breakdown with cost and latency metrics.

Environment	Success	Tokens (K)	Cost (\$)	Latency (s)	Dominant Pattern
Operating System (OS)	58.4%	12.3	0.25	18.4	Hierarchical (42%)
Database (DB)	52.3%	8.7	0.17	12.1	Parallel (38%)
Knowledge Graph (KG)	64.5%	15.2	0.30	24.2	Hierarchical (45%)
Card Game (CG)	47.2%	6.4	0.13	8.7	Single (52%)
Lateral Thinking (LT)	54.8%	18.9	0.38	32.1	Consensus (35%)
House-Holding (HH)	67.8%	22.4	0.45	28.5	Hierarchical (48%)
Web Shopping (WS)	75.8%	14.8	0.30	19.2	Parallel (41%)
Web Browsing (WB)	61.2%	11.2	0.22	15.8	Parallel (39%)
Average	59.7%	13.7	0.28	19.9	—

Key Observations:

- **Highest success:** Web Shopping (75.8%)—well-structured e-commerce tasks benefit from parallel product comparison
- **Lowest success:** Card Game (47.2%)—requires strategic reasoning where single-agent often suffices
- **Most expensive:** House-Holding (\$0.45)—complex multi-step tasks requiring hierarchical coordination
- **Longest latency:** Lateral Thinking (32.1s)—consensus pattern needed for creative reasoning

V.2. WebArena Detailed Breakdown

Table 72. WebArena per-site breakdown with detailed metrics.

Site	Success	Tokens (K)	Cost (\$)	Latency (s)	Dominant Pattern
Shopping (OneStopShop)	34.2%	28.4	0.57	42.1	Parallel (44%)
Forum (Reddit)	29.8%	22.1	0.44	35.8	Hierarchical (38%)
GitLab	26.5%	31.2	0.62	48.2	Hierarchical (51%)
CMS (WordPress)	28.3%	19.8	0.40	31.4	Parallel (36%)
Overall	29.7%	25.4	0.51	39.4	—

Key Observations:

- **GitLab** has highest cost (\$0.62) and latency (48.2s) due to complex code navigation requiring hierarchical orchestration
- **Shopping** achieves best success (34.2%) with parallel pattern for product search and comparison
- WebArena tasks are generally harder than AgentBench (29.7% vs 59.7%), requiring more tokens per task

V.3. Cost and Latency by Orchestration Pattern

Table 73. Detailed cost and latency breakdown by orchestration pattern across all benchmarks.

Pattern	Usage	Tokens (K)	Cost (\$)	Latency (s)	Success	Quality
<i>AgentBench</i>						
Single	31.2%	4.2	0.08	5.2	54.1%	0.62
Parallel	35.8%	9.8	0.20	8.4	61.2%	0.68
Hierarchical	24.5%	18.4	0.37	22.1	64.8%	0.74
Consensus	8.5%	32.1	0.64	38.5	68.2%	0.79
<i>WebArena</i>						
Single	18.4%	8.2	0.16	12.1	21.2%	0.54
Parallel	38.2%	21.4	0.43	28.4	31.4%	0.61
Hierarchical	32.1%	35.8	0.72	52.1	34.2%	0.67
Consensus	11.3%	48.2	0.96	68.4	38.5%	0.72
<i>MARS</i>						
Single	28.4%	5.8	0.12	6.8	45.2%	0.58
Parallel	32.1%	12.4	0.25	11.2	52.8%	0.66
Hierarchical	26.8%	24.2	0.48	28.4	58.4%	0.73
Consensus	12.7%	42.8	0.86	45.2	64.1%	0.78

Cost-Performance Trade-offs:

1. **Single** → **Parallel**: +2.4× cost for +7-10% success improvement
2. **Parallel** → **Hierarchical**: +1.8× cost for +3-6% success improvement
3. **Hierarchical** → **Consensus**: +1.7× cost for +3-4% success improvement
4. **Diminishing returns**: Each escalation level provides smaller marginal gains at higher cost

V.4. Token Distribution Analysis

Efficiency Insights:

- Agent execution dominates token usage (58-68%)
- CAR overhead is minimal (1.8-2.4%)—efficient pre-execution routing
- WebArena requires 1.9× more tokens than AgentBench due to complex web interactions
- MARS has highest inter-agent communication (14.2%) due to multi-domain coordination

Table 74. Token usage breakdown by component.

Component	AgentBench	WebArena	MARS
Task decomposition	8.2%	6.4%	9.1%
CAR feature extraction	2.1%	1.8%	2.4%
Agent execution (primary)	62.4%	68.2%	58.7%
Inter-agent communication	12.8%	11.4%	14.2%
RCM quality assessment	5.4%	4.8%	6.1%
Memory operations	4.2%	3.8%	5.2%
Retry/escalation overhead	4.9%	3.6%	4.3%
Total (avg tokens/task)	13.7K	25.4K	18.2K

W. Detailed Limitations and Future Directions

W.1. Current Limitations

L1: CAR Training Data Requirements. Training the CAR model requires execution traces with ground-truth pattern labels. Collecting this data for new domains requires initial exploration with exhaustive pattern search, which is expensive. Future work could explore few-shot or zero-shot pattern prediction.

L2: Quality Assessment Reliability. RCM relies on LLM-as-judge for quality assessment, which inherits LLM biases and may miss subtle errors, particularly in specialized domains. Domain-specific evaluators could improve reliability.

L3: Memory Consolidation Timing. UMA’s consolidation currently runs at fixed intervals. Adaptive consolidation triggered by memory pressure or task transitions could be more efficient.

L4: Coordination Overhead. While AMORE reduces unnecessary multi-agent usage, the coordination infrastructure itself adds latency (0.5s per subtask). Optimizing this overhead is important for latency-sensitive applications.

L5: Limited Protocol Support. Current implementation uses custom agent communication. Full integration with A2A, MCP, and ACP protocols would enable broader interoperability.

W.2. Future Research Directions

F1: Online Learning for CAR. Continuously updating CAR from execution feedback would improve pattern prediction over time without requiring pre-collected training data.

F2: Hierarchical Memory. Extending UMA to support team-level and organization-level memory tiers for enterprise deployments.

F3: Human-in-the-Loop Integration. Adding checkpoints for human intervention when automated quality assessment has low confidence.

F4: Multi-Modal Agents. Extending AMORE to orchestrate agents with different modalities (text, vision, code execution).

F5: Formal Verification. Developing formal methods to verify orchestration correctness and safety properties.